

Systems I – CS 6013

Computer Architecture and Operating Systems

Lecture 4: Function Calls in ASM

MASTER OF SOFTWARE DEVELOPMENT (MSD) PROGRAM

J. DAVISON DE ST. GERMAIN

SPRING 2024

Miscellaneous

- Questions?
- `gcc -o doit doit.c -Wall`
 - `-g`
 - Debug mode
 - `-O`
 - Optimized mode
 - `-S`
 - Output assembly code (.s file)
- `doit & #` What is the ampersand for?
[1] 7353
 - Runs in the *background*. What is 7353? What is 1?
 - PID – process identification number
 - Background process number.
 - `fg %1, c-z` (control z)
- `./doit` # What is the `./` and why is it there?
 - `echo $PATH`
- Make sure to look at Week 2 reading list
- New HW assignment has been posted
 - [Has old Apple MacBook instructions that do not work with new MacBooks (Apple Silicon / ARM)]
 - Not due for 2 weeks.
- Diego, please see footnote 7 in the reading (Introduction (Chapter 2)). ☺

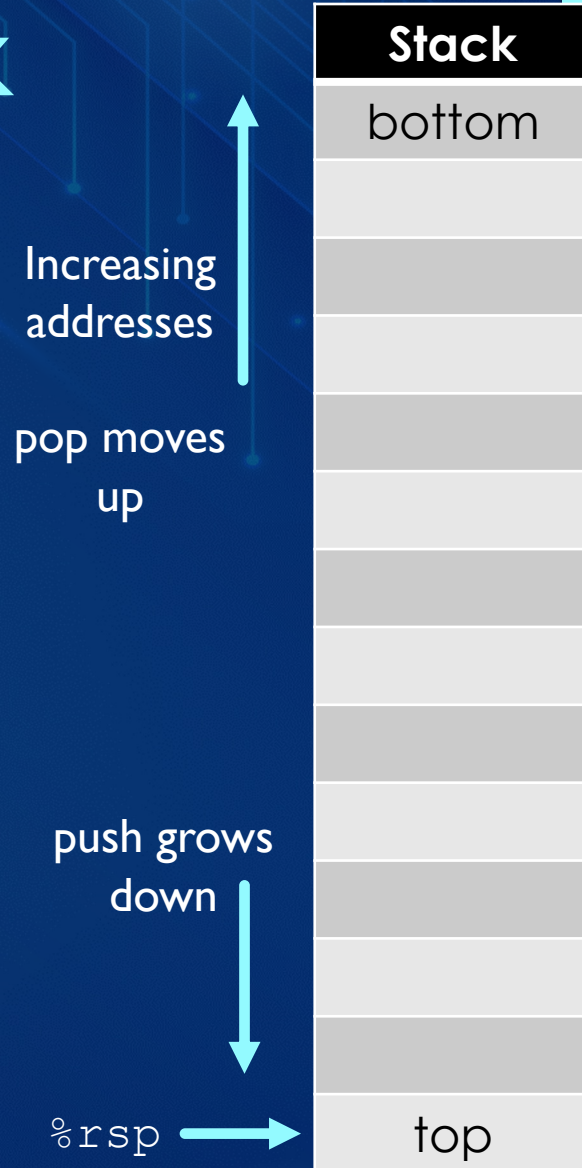
Lecture 4 – Topics

3

- Call Stack
- Function Parameters

Function Calls and The Stack

- The call stack (remember this?)
 - While in previous lectures, to match a “physical stack” I’ve displayed the stack as growing upward... in reality, with respect to memory addresses, it grows downward.
- When a function is called, information about it is pushed onto the stack. Space for what information?
 - Parameters
 - Local Variable
 - Return Value(s)
 - Location of the previous stack frame
 - etc.
- What registers can I use?
- Answers to these questions are described by the *Calling Conventions* or *Application Binary Interface (ABI)*



Example: System V ABI for x86_64

- Basically the ABI for Unix – Specific to language / operation system. Note, not the same as the API (Application Programming Interface).
- First 6 int or pointer parameters/arguments to a function must be placed in rdi, rsi, rdx, rcx, r8, r9
- Floating point parameters in xmm0, ... xmm7
- Int return values must be placed in rax (and you look for them there when trying to capture return value).
- rbx, rbp, r12-r15 are “callee-saved” registers. If the function uses them, it must restore them before it returns
- Other registers may be “clobbered” by the function
- Other parameters are passed on the stack

Passing / Returning Values in Assembly

6

- Passing parameters:

```
int main() {  
    int x = doit( 9, 12, 45 );  
    // mov rdi, 9  
    // mov rsi, 12  
    // mov rdx, 45  
    // call doit  
}
```

- Compute and return a value

```
int doit( int a, int b, int c ) {  
    int result = a + b + c;  
    return result;  
    // doit:  
    //     mov rbx, 0 ; rbx is result  
    //     add rbx, rdi  
    //     add rbx, rsi  
    //     add rbx, rdx  
    //     mov rax, rbx  
}
```

- Can do the above without first and last mov.

Passing / Returning Values in Assembly

7

- Passing parameters:

```
int main() {  
    int x = doit( 9, 12, 45 );  
    // mov rdi, 9  
    // mov rsi, 12  
    // mov rdx, 45  
    // call doit  
}
```

- Compute and return a value

```
int doit( int a, int b, int c ) {  
    int result = a + b + c;  
    return result;  
    // doit:  
    //   mov rbx, rdi ; rax is result  
    //   add rbx, rsi  
    //   add rbx, rdx  
}
```


Register Protocols

- Some registers are temporaries
 - Call a function => register value may have changed on return
 - A.k.a. Caller-saved
 - eg: %r10, %rsi
- Some registers are preserved
 - Call a function => register value the same on return
 - a.k.a. callee-saved
 - eg: %rbx, %rsp
- Classification of registers is part of an:
 - Application Binary Interface (ABI)

x86-64 Linux Register Usage

9

| | register | usage |
|---------------------|----------|---------------|
| Caller-saved | %rax | return value |
| | %rdi | 1st argument |
| | %rsi | 2nd argument |
| | %rdx | 3rd argument |
| | %rcx | 4th argument |
| | %r8 | 5th argument |
| | %r9 | 6th argument |
| | %r10 | temporary |
| | %r11 | temporary |
| Callee-saved | %rbx | preserved |
| | %r12 | preserved |
| | %r13 | preserved |
| | %r14 | preserved |
| | %rbp | stack frame |
| | %rsp | stack pointer |

Application Binary Interface

An OS-specific ABI defines:

- How arguments are passed to functions
 - So far, only integer and address args
- How results are returned from functions
 - So far, only integer and address results
- Which registers are preserved (and not)
 - There are more registers...
- Other constraints, such as stack alignment
 - x86-64 Linux: stack aligned on call at 8 mod 16
- Optional debugging protocols

Managing the Stack

II

- Function **must** set up stack before running its body
- Must **clean up** before returning
- Setup is called “prologue”
- Tear down is called “epilogue”

Prologue

- Save the base pointer (bp) to the stack, then overwrite it with sp
- Make room on stack for other function calls by *subtracting* from sp (stack grows downward)
- Copy arguments into locations relative to bp

```
int g( int y ) {  
}
```

```
int f( int x ) {  
    int z = 44;  
    return g( z )  
}
```

```
main() {  
    int a = 2;  
    f( a );  
}
```

Note, “main()” and “a”, and “Address” don’t actually exist on the Stack. They are implicitly there, but we write them down like this just for our own information.

| Stack | | | |
|--------|---|---|---------|
| | | | Address |
| main() | a | 2 | 1004 |
| | | | 1000 |
| | | | 996 |
| | | | 992 |
| | | | 988 |
| | | | 984 |
| | | | 980 |
| | | | ... |
| | | | |
| | | | |
| | | | |
| | | | |
| | | | |
| | | | |
| | | | |
| | | | |

Prologue

- Save the base pointer (bp) to the stack, then overwrite it with sp
- Make room on stack for other function calls by *subtracting* from sp (stack grows downward)
- Copy arguments into locations relative to bp

```
int g( int y ) {  
}  
  
int f( int x ) {  
    int z = 44;  
    return g( z )  
}  
  
main() {  
    int a = 2;  
    f( a );  
}
```

| Registers | |
|-----------|--------|
| bp | 1004 |
| sp | 996 |
| ip | <addr> |
| | |
| | |
| | |
| | |
| | |
| | |

| Stack | | | |
|--------|---|----|---------|
| | | | Address |
| main() | a | 2 | 1004 |
| | | | 1000 |
| f() | x | 2 | 996 |
| | z | 44 | 992 |
| | | | 988 |
| | | | 984 |
| | | | 980 |
| | | | ... |
| | | | |
| | | | |
| | | | |
| | | | |
| | | | |
| | | | |
| | | | |

13

- Old / simplistic view... But what actually happens?
- Need to handle the prologue...

Prologue

- Save the base pointer (bp) to the stack, then overwrite it with sp
- Make room on stack for other function calls by *subtracting* from sp (stack grows downward)
- Copy arguments into locations relative to bp

```
int g( int y ) {  
}  
  
int f( int x ) {  
    int z = 44;  
    return g( z )  
}  
  
main() {  
    int a = 2;  
    f( a );  
}
```

| Registers | |
|-----------|--------|
| bp | 1004 |
| sp | 996 |
| ip | <addr> |
| | |
| | |
| | |
| | |
| | |

| Stack | | | |
|--------|---|---|---------|
| | | | Address |
| main() | a | 2 | 1004 |
| | | | 1000 |
| | | | 996 |
| | | | 992 |
| | | | 988 |
| | | | 984 |
| | | | 980 |
| | | | ... |
| | | | |
| | | | |
| | | | |
| | | | |
| | | | |
| | | | |
| | | | |
| | | | |

14

- What does the prologue require?

Prologue

- Save the base pointer (bp) to the stack, then overwrite it with sp
- Make room on stack for other function calls by *subtracting* from sp (stack grows downward)
- Copy arguments into locations relative to bp

```
int g( int y ) {  
}  
  
int f( int x ) {  
    int z = 44;  
    return g( z )  
}  
  
main() {  
    int a = 2;  
    f( a );  
}
```

| Registers | |
|-----------|--------|
| bp | 996 |
| sp | 984 |
| ip | <addr> |
| | |
| | |
| | |
| | |
| | |
| | |

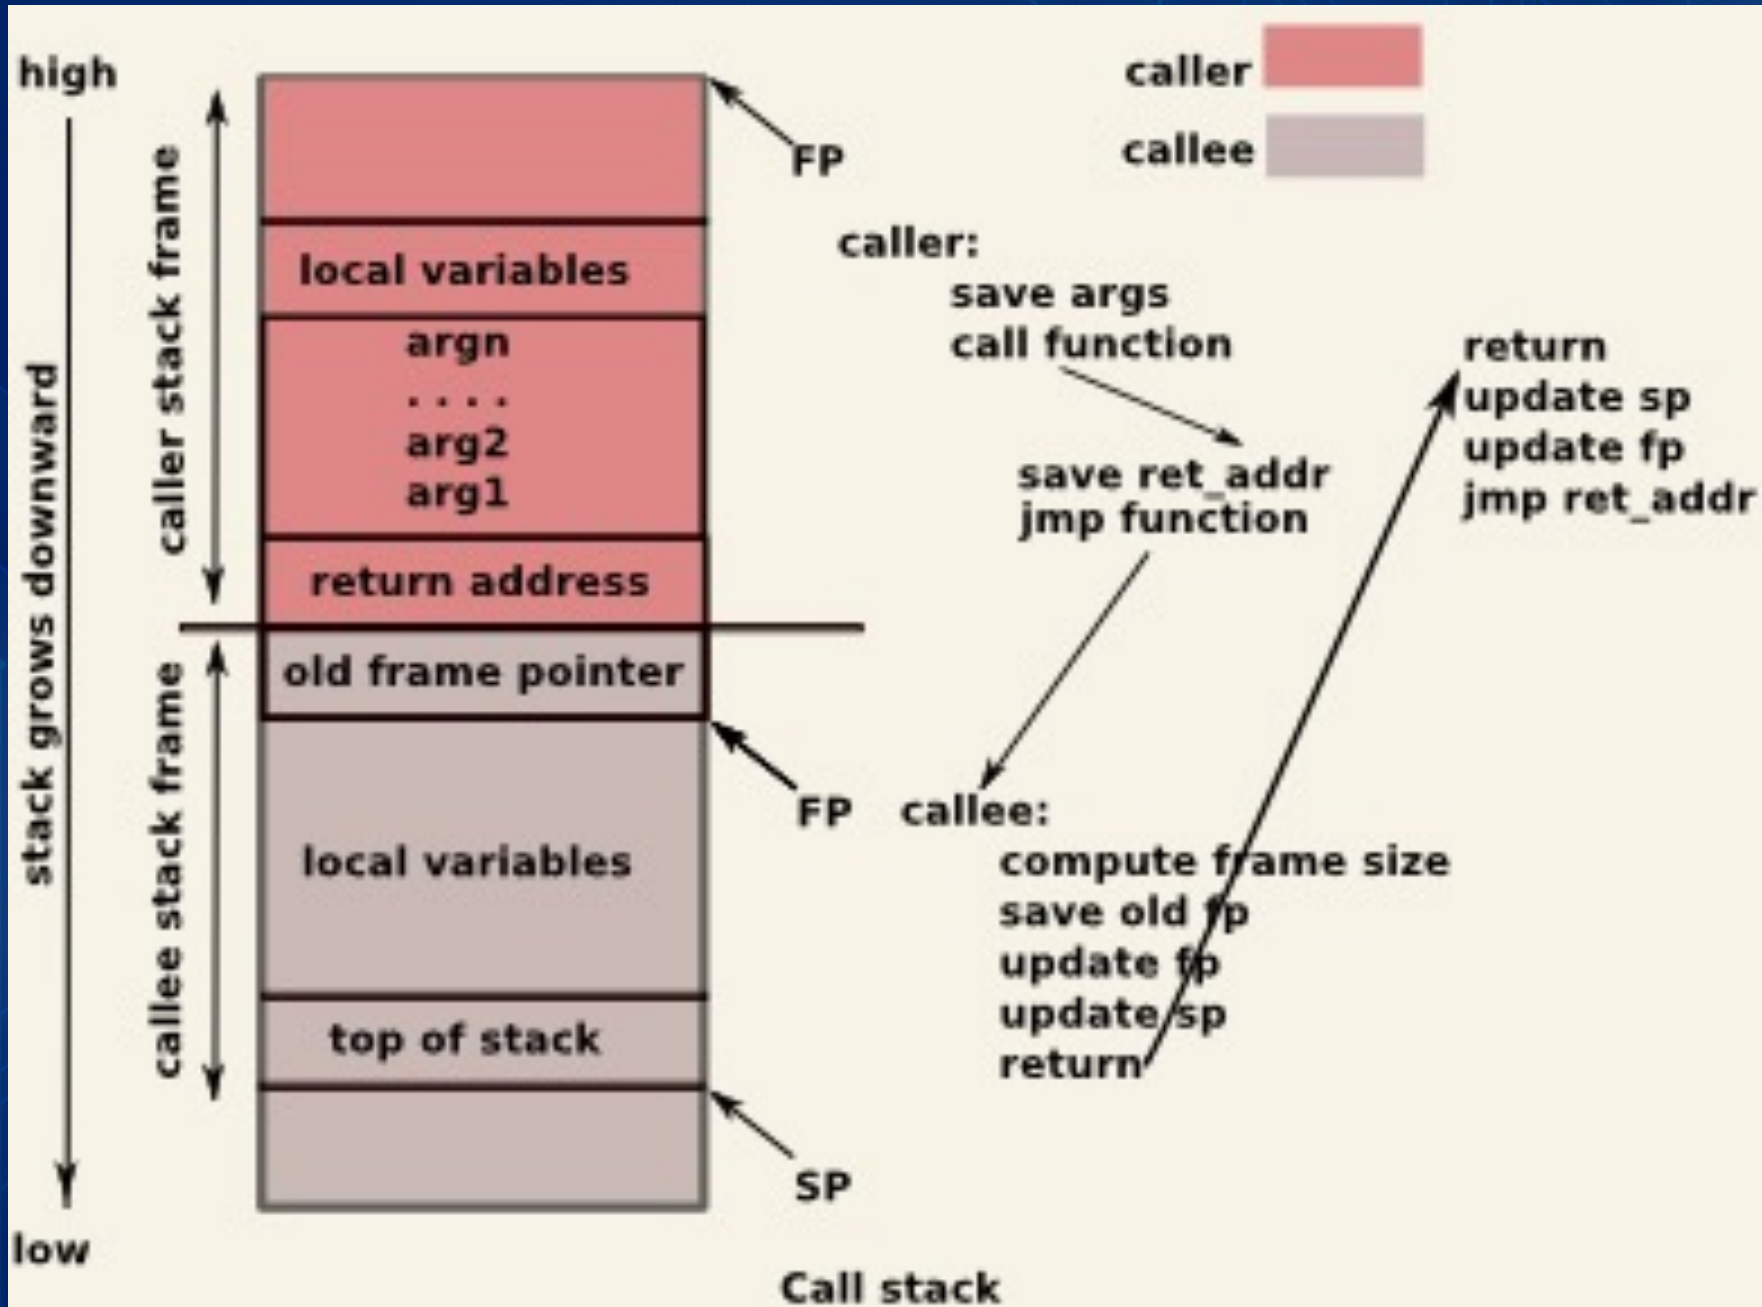
| Stack | | | |
|--------|------|------|---------|
| | | | Address |
| main() | a | 2 | 1004 |
| | [ip] | addr | 1000 |
| f() | [bp] | 1004 | 996 |
| | x | 2 | 992 |
| | z | 44 | 988 |
| | | | 984 |
| | | | 980 |
| | | | ... |
| | | | |
| | | | |
| | | | |
| | | | |
| | | | |
| | | | |
| | | | |
| | | | |

What (Why) is ip?
mov [bp-4], rdi
What does this do?
Copies rdi into the address of base pointer - 4.
What is rdi?
The 1st parameter - in this case x (with the value of 2)
What would it look like after calling g()?

Epilogue

16

- Store return int into rax. Why?
- Then, undo what we just did in the prologue...
- Restore the base pointer so that when we return, the caller's stack pointer is in the right place.
- `pop rbp`
 - What is rbp vs bp from previous slide?
 - r == 64 bit version of register.



- Note 1: In this diagram the BP (Base Pointer) is referred to as the FP (Frame Pointer).
- Note 2: SP (Stack Pointer) always points to the top of the stack (remember, stack is growing downward)

Special Instructions

- `call functionLabel ; save the program counter ; (the special/secret IP ; register) to the stack, ; then jump to functionLabel`
- `leave ; restore rbp from the stack`
- `ret ; jump to the return address saved by call ; instruction and pop it off the stack`

Examples

- Write some code and see the assembly...
 - Godbolt.org
 - **Assembly is close, but not exact to our syntax...**

```
int main() { return 0; }
```

- push, pop, why eax,0

```
int main() { int i = 3; return 0; }
```

- `mov [rbp-4], 3 ; what is this?`
 - Save the value of 3 on the stack. (4 bytes)

```
int doit() { return 10; }
```

```
int main() { int g = doit(); return g; }
```

- Look at assembly output with `-O` option...

~ Fin ~