

Systems I – CS 6013

Computer Architecture and Operating Systems

Lecture 7: Processes I

MASTER OF SOFTWARE DEVELOPMENT (MSD) PROGRAM

J. DAVISON DE ST. GERMAIN

SPRING 2024

*Adapted from Ryan Stutsman's slides

CS 6013 – Spring 2024

Miscellaneous

- Dereferencing a Pointer
 - Address vs Value
- Setting up and taking down new Stack Frame – epilogue / prologue

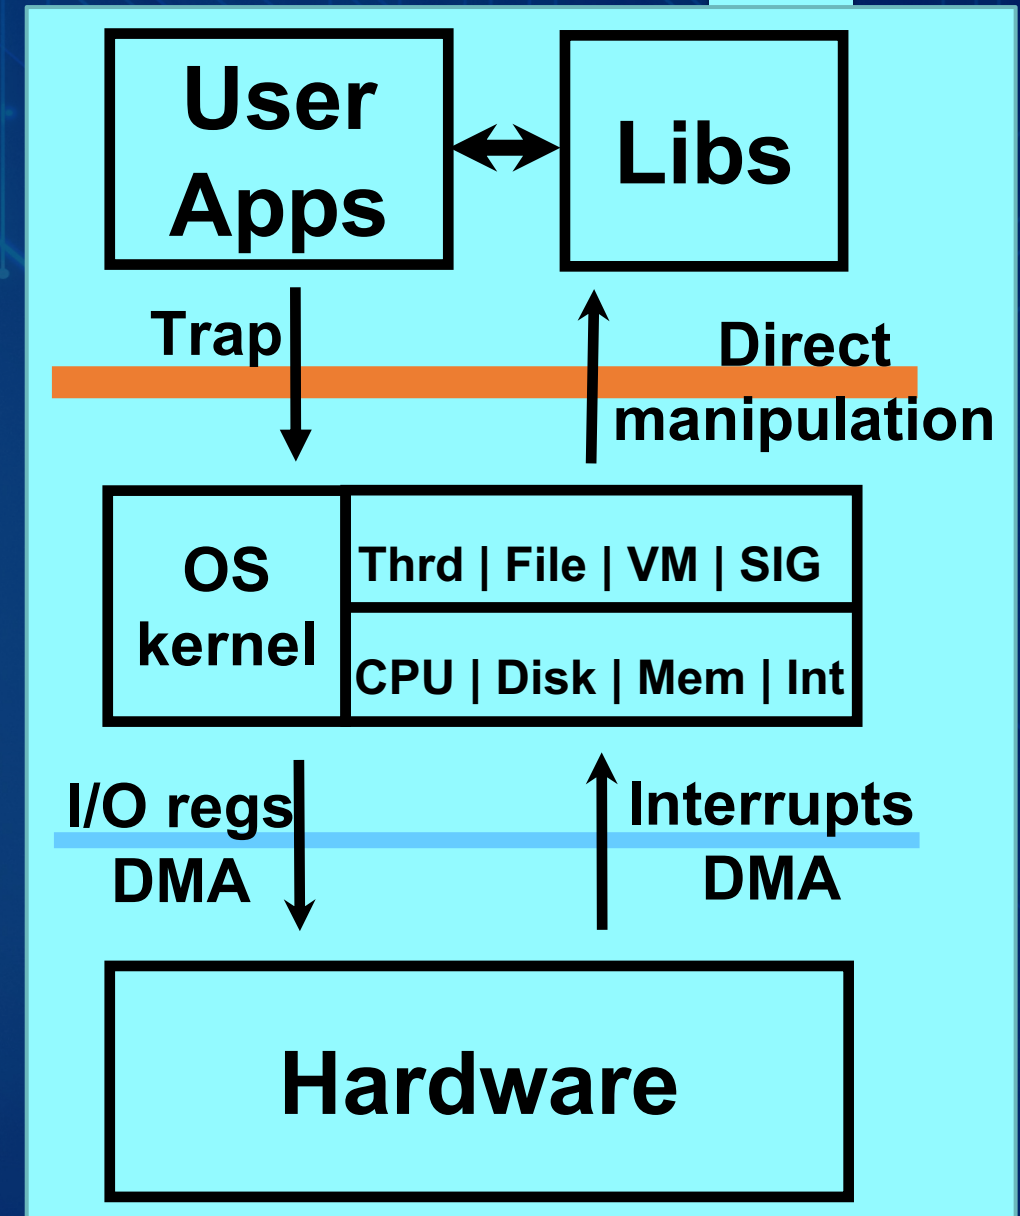
Lecture 6 – Topics

3

- Processes
 - Memory Layout
 - PCB – Process Control Block
 - Creating Processes – Fork()

Standard OS Structure

- User-level
 - Applications
 - Libraries: many common “OS” functions
 - Example: malloc() vs sbrk()
- Kernel-level
 - Top-level: machine independent
 - Bottom-level: machine dependent
 - Runs in protected mode
 - Need a way to switch (user \leftrightarrow kernel)
- Hardware-level
 - Device maker exports known interface
 - Device driver initiates operations by writing to device registers
 - Devices use interrupts to signal completion
 - DMA – Direct Memory Access (offloads work, but has restrictions)



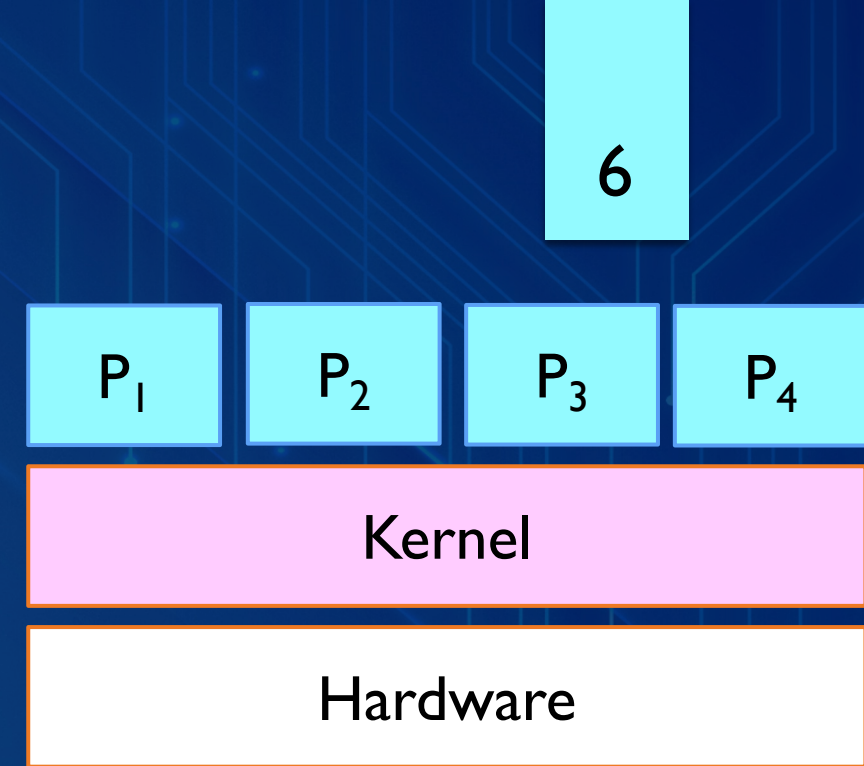
What is a Process?

5

- **Process**: execution context of running program
- A **process** does not equal a **program**!
 - Process is an *instance* of a program. Where have we talked about instances before?
 - Many copies of same program can be running at same time
- OS manages a variety of activities
 - User programs
 - Batch jobs and scripts (which are just other programs)
 - System programs – print spool, file servers, net daemons
 - BTW, what is a daemon? (Pronounced Demon)
 - A process that runs in the background... Usually on a long term basis. (eg: sshd, ntpd, httpd, ...)
- Each of these activities is encapsulated in a process
- Everything happens either in kernel or a process
 - The OS is not a process

Isolating Processes

- Lots of running processes
- Each with own code, data
- Each need to interact with devices, memory, CPU
- How do we **multiplex** the hardware among them?
- How do we make this safe? Efficient?



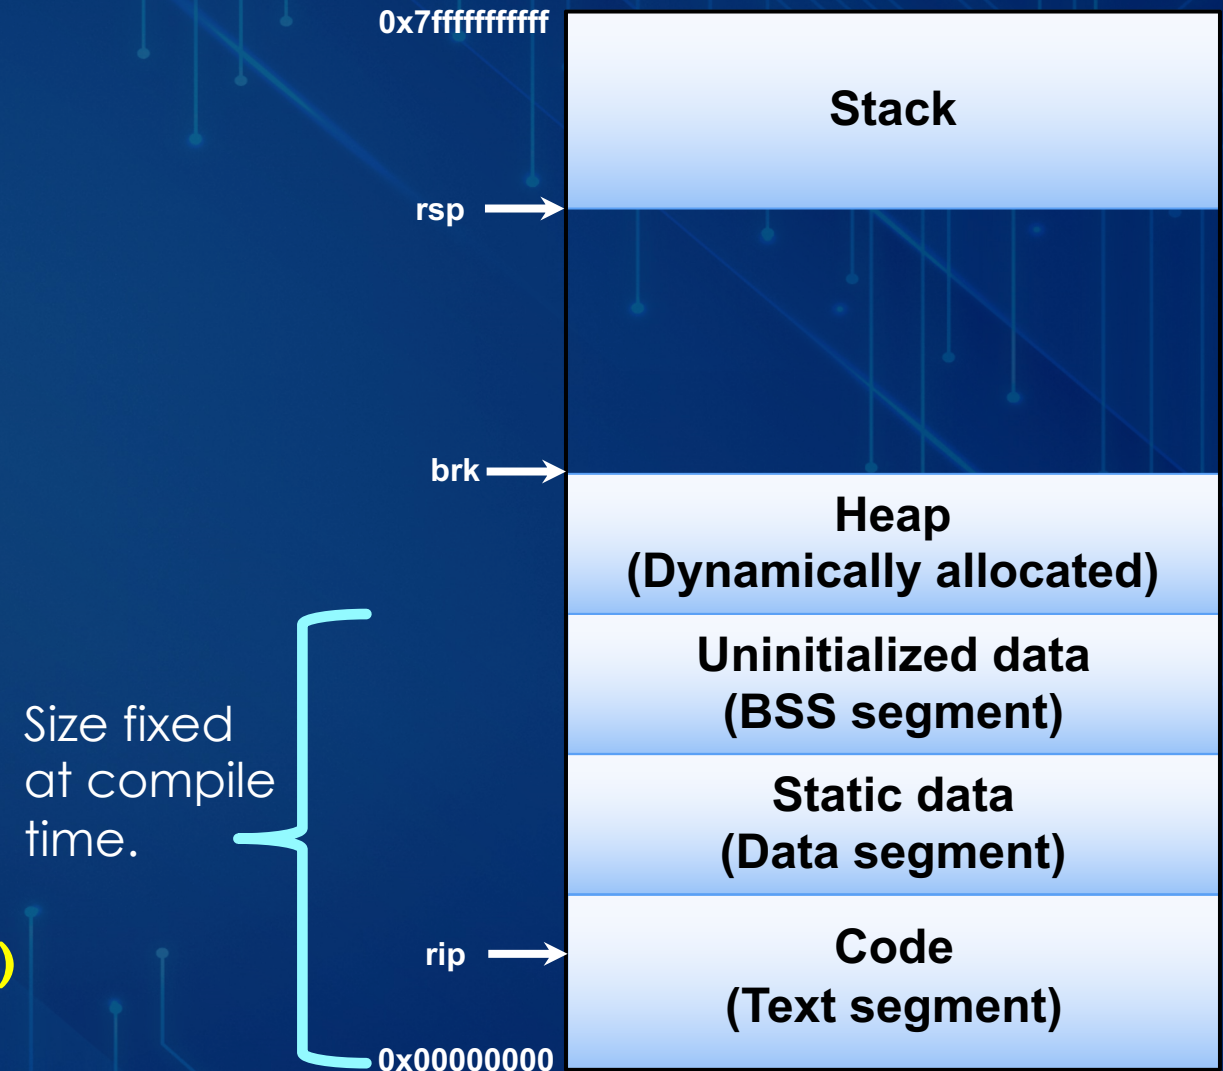
What is in a Process?

- Process state consists of:
 - Memory state: code, data, heap, stack
 - Processor state: IP, registers, etc.
 - Kernel state:
 - Process state: ready, running, etc.
 - Resources: open files/sockets, etc.
 - Scheduling: priority, CPU time, etc.
- **Address space** consists of:
 - Code
 - Static data (data and BSS)
 - Dynamic data (heap and stack)
 - See: Unix `size` command
- Special pointers:
 - IP: current instruction being executed
 - brk: top of heap (explicitly moved)
 - sp: bottom of stack (implicitly moved)

All tracked in a **Process Control Block (PCB)**

Process Memory Layout


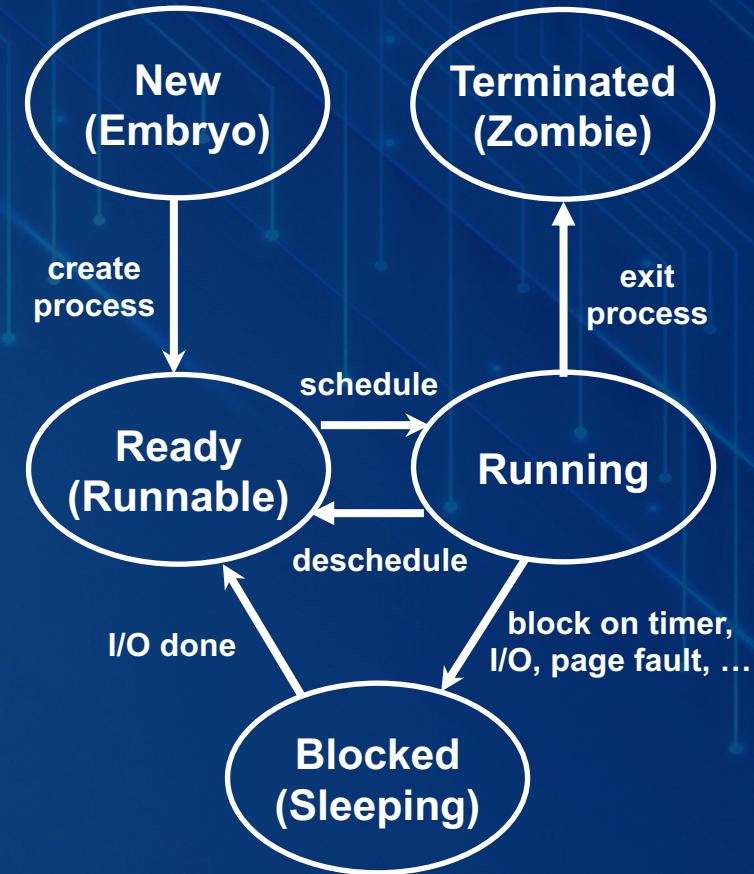
7



Process State Machine

8

- Each process has a state:
 - new: OS is setting up process
 - ready: runnable, but not running
 - running: executing instructions on CPU
 - blocked: stalled for some event (e.g., IO)
 - terminated: process is dead or dying
- Process moves from state to state as a result of its actions or external actions
 - Program: sleep(), IO request, ...
 - OS action: scheduling
 - External: interrupts, IO completion
- Key to efficiency and high utilization
 - Blocked processes yield CPU
 - They are just data structures



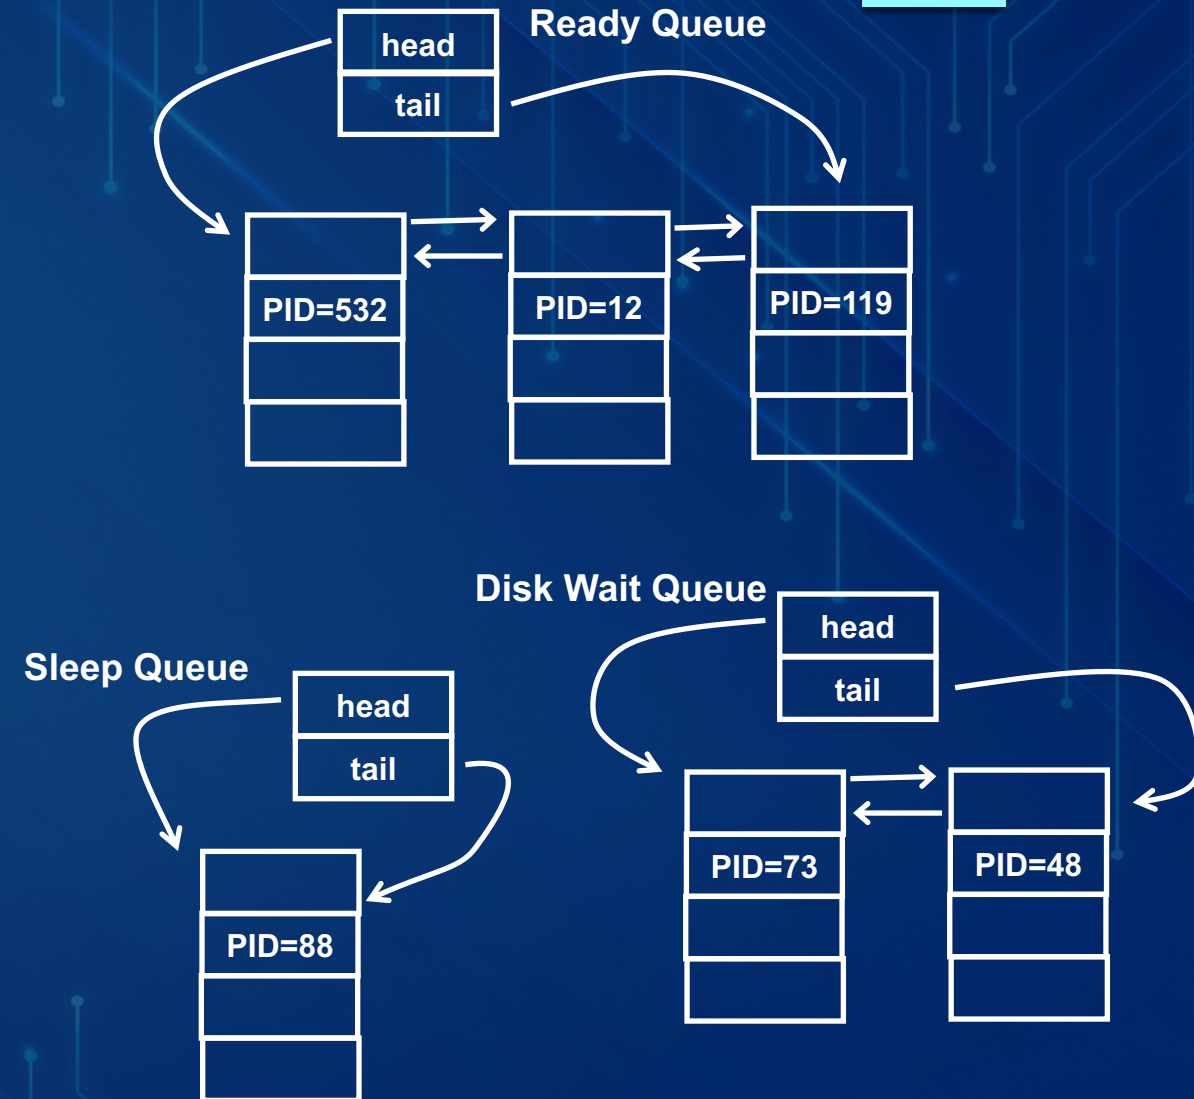
Process Control Block (PCB)

- the PCB is a data structure - One per process, allocated in kernel memory
- Stores state of a process, typically including:
 - Process state (running, waiting, ...)
 - PID (process identifier, often a 16-bit integer)
 - Machine state: IP, SP, registers (if not currently running)
 - Memory management info
 - Open file table (open socket table)
 - Queue pointers (waiting queue, I/O, sibling list, parent, ...)
 - Scheduling info (e.g., priority, time used so far, ...)
- On process creation, the kernel will: allocate PCB, initialize, put on ready queue (queue of runnable processes)
- On exit: clean up all process state (close files, release memory, page tables, etc)

Process State Queues

10

- OS tracks PCBs using queues
- Ready processes on the “ready queue”
- Each I/O device has a wait queue
 - Queue traversed when I/O interrupt handled
- OS invariant: A process is always running. Other processes are either on the ready queue, or on a single wait queue
 - Implications of this?
- Processes linked to parents and siblings
 - Needed to support wait()



xv6 PCB

```
struct context {  
    uint edi;  
    uint esi;  
    uint ebx;  
    uint ebp;  
    uint eip;  
};
```

- What does “e” (in registers listed) tell you about this system?
 - 32-bit computer

```
enum proc_state { UNUSED, EMBRYO, SLEEPING,  
                  RUNNABLE, RUNNING, ZOMBIE };  
  
struct proc {  
    uint sz; // Proc mem size (bytes)  
    pde_t* pgdir; // Page table  
    char *kstack; // Bottom of kstack  
    enum procstate state; // Process state  
    int pid; // Process ID  
    struct proc *parent; // Parent process  
    struct trapframe *tf; // Trap frm for syscall  
    struct context *context; // swtch() here to run  
    void *chan; // If !0, sleeping on chan  
    int killed; // If !0, have been killed  
    struct file *ofile[NOFILE]; // Open files  
    struct inode *cwd; // Current directory  
    char name[16]; // Process name  
};
```


Process Management

12

- OS manages processes:
 - Creates, deletes, suspends, and resumes processes
 - Schedules processes to manage CPU allocation
 - Manages inter-process communication and synchronization
 - Allocates resources to processes (and takes them away)
- Processes use OS functionality to cooperate
 - Signals, sockets, pipes, files to communicate

Practical Process Management

13

- On a Unix machine, try:
 - `ps -Af` # lots of info on all running processes
 - `kill -9 547` # terminates process with PID 547 (Don't actually try this one...!?)
 - `top` # displays dynamic info on top running jobs
 - Write a program that calls:
 - `getpid()`: returns current process's PID (process id)
 - `fork()`: create a new process
 - `wait()`: wait for exit of a child process
 - `exec()`: load a new program into the current process
 - `sleep()`: puts current process to sleep for specified time
- Commands work on macOS / linux
- On Windows → Task manager (CTL-ALT-DEL)

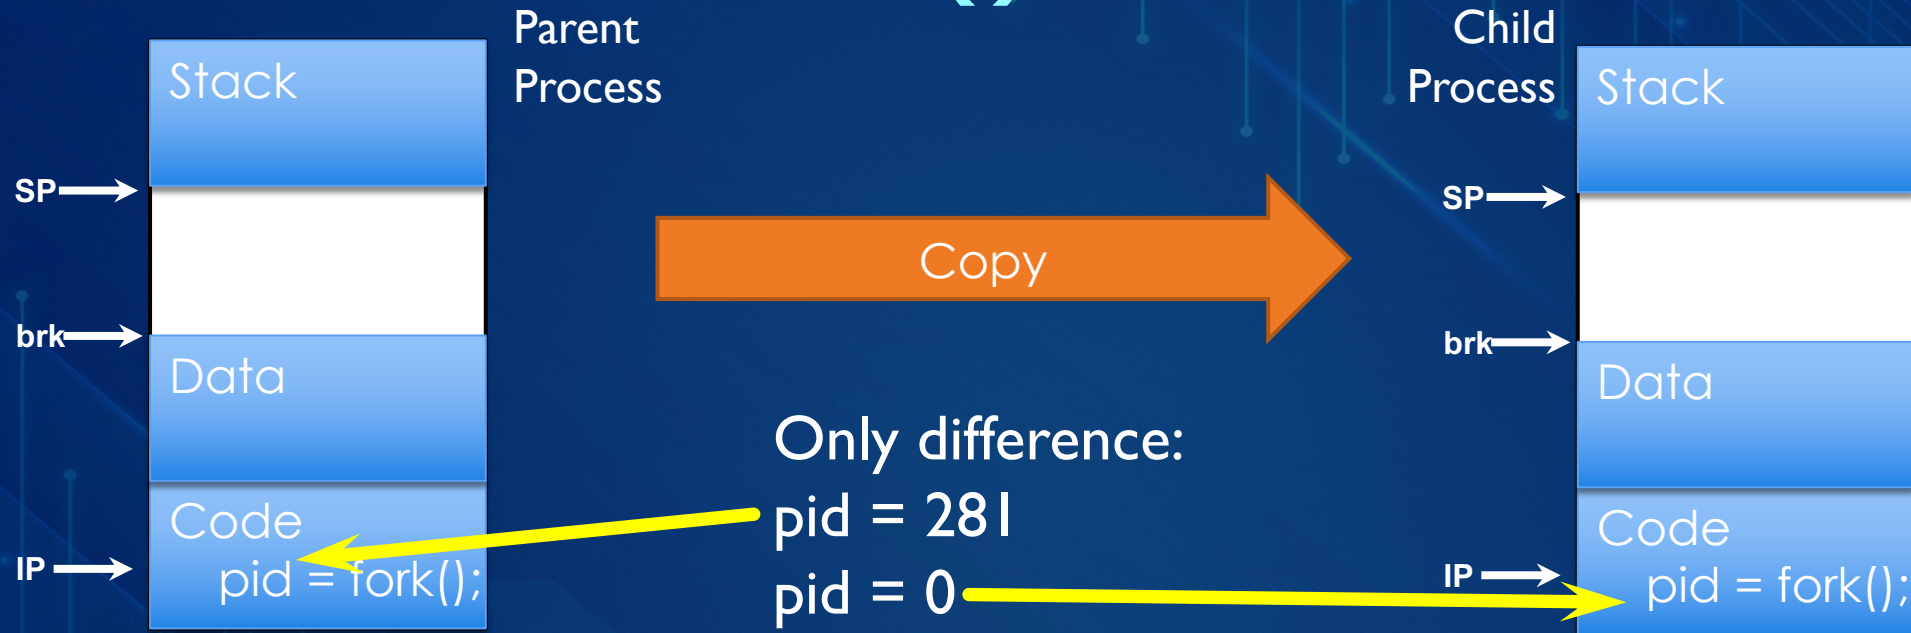
Creating Processes: `fork()`

14

- Creates process that is near-clone of forking parent
 - Address space and running state is cloned
- Return of `fork()` differs:
 - 0 in child
 - child PID in parent
 - 3rd Option?
 - -1 – ERROR! Check error codes.
- Many kernel resources are shared between the parent and child...
 - open files and sockets
 - have to be careful
- `wait()` lets a process wait for the exit of a child
- To spawn a new program, use some form of `exec()`

Semantics of fork()

15



- fork(), exit(), and exec() are weird!
 - fork() returns twice – once in each process
 - exit() does not return at all
 - exec() usually “does not return”: replaces process’ program

fork() and wait() demo

16

- Note: the `fork()` call can return first in child or parent process...

Fork() / Wait ()

17

```
int main(){
    pid_t pid = fork();
    if( pid < 0 ) {
        perror( "Fork failed..." );
        return -1;
    }
    else if( pid == 0 ) {
        printf( "HC: hello from child\n" );
        exit(17);
    }
    else {
        int child_status;
        printf( "HP: hello from parent\n" );
        waitpid( pid, &child_status, 0 ); // Waits for child to end
        printf( "CT: child result %d\n", WEXITSTATUS( child_status ) );
    }
    printf( "Bye\n" );
    return 0;
}
```


~ Fin ~