

Systems I – CS 6013

Computer Architecture and Operating Systems

Lecture I2: Scheduling (Part 2) and Inter-process Communication (IPC)

MASTER OF SOFTWARE DEVELOPMENT (MSD) PROGRAM

J. DAVISON DE ST. GERMAIN

SPRING 2024

Lecture 12 – Topics

2

- Scheduling
 - Note: Review book chapters 8, 9
- IPC – Inter-Process Communication

Miscellaneous

- Reminders
 - Week 4 Vocabulary Assignment - Due Tonight
 - Week 5 Reading Assignment (Chapters 12-15)
- Lab after lecture today
 - It will help prepare you for the Shell assignment
- Emacs is cool... Emacs vs PowerPoint
 - What does a PowerPoint slide look like?

Scheduling Basics

Workloads:

- Arrival Time
- Run Time

Schedulers:

- FIFO
- SJF
- STCF
- RR

Metrics:

- turnaround time
- response time


* First In, First Out; Shortest Job First; Shortest Time to Completion First; Round Robin

Workload Assumptions

- ~~1. Each job runs for the same amount of time~~
- ~~2. All jobs arrive at the same time~~
- ~~3. All jobs only use the CPU (no I/O)~~
4. The run-time of each job is known

What Happens if Scheduler is Not I/O Aware


6



- If the Scheduler is not aware of a Process doing I/O...
- Don't let Job A hold on to CPU while blocked waiting for disk

I/O Aware (Overlap)

7



Treat Job A as 3 separate **CPU bursts** (sub-jobs)
When Job A completes I/O, “another” Job A is ready

Each CPU burst is shorter than Job B, so with STCF,
Job A preempts Job B

Workload Assumptions

- ~~1. Each job runs for the same amount of time~~
 - ~~2. All jobs arrive at the same time~~
 - ~~3. All jobs only use the CPU (no I/O)~~
 - ~~4. The run-time of each job is known~~
- We've used these assumptions to exam the basics of scheduling... But with modern systems, these assumptions do not apply, so...
 - We need a smarter, fancier scheduler!

MLFQ (Multi-Level Feedback Queue)

9

- Design Goal: general-purpose scheduling
- Must support two job types with distinct goals:
 - interactive programs care about **response time**
 - batch programs care about **turnaround time**
- Approach: multiple queues, each for a different priority level.
- Higher queue has a higher priority.
- If more than one job in one queue, use Round Robin (RR) on that queue.

Priorities

10

Q3 has the highest priority

Rule 1: If $\text{Priority}(A) > \text{Priority}(B)$, A runs

Rule 2: If $\text{Priority}(A) == \text{Priority}(B)$, A & B run in RR

More rules to come...



Q1



Can think of this as a “Multi-level Priority Queue”

How do we set priorities?

History of Process' Past Behavior


II

- Use past behavior of process to predict future behavior (**Heuristic**)
 - Common technique in systems
- Assumption: Most processes alternate between I/O and CPU work
- Guess how CPU burst (job) will behave based on past CPU bursts (jobs) of this process

More MLFQ Rules

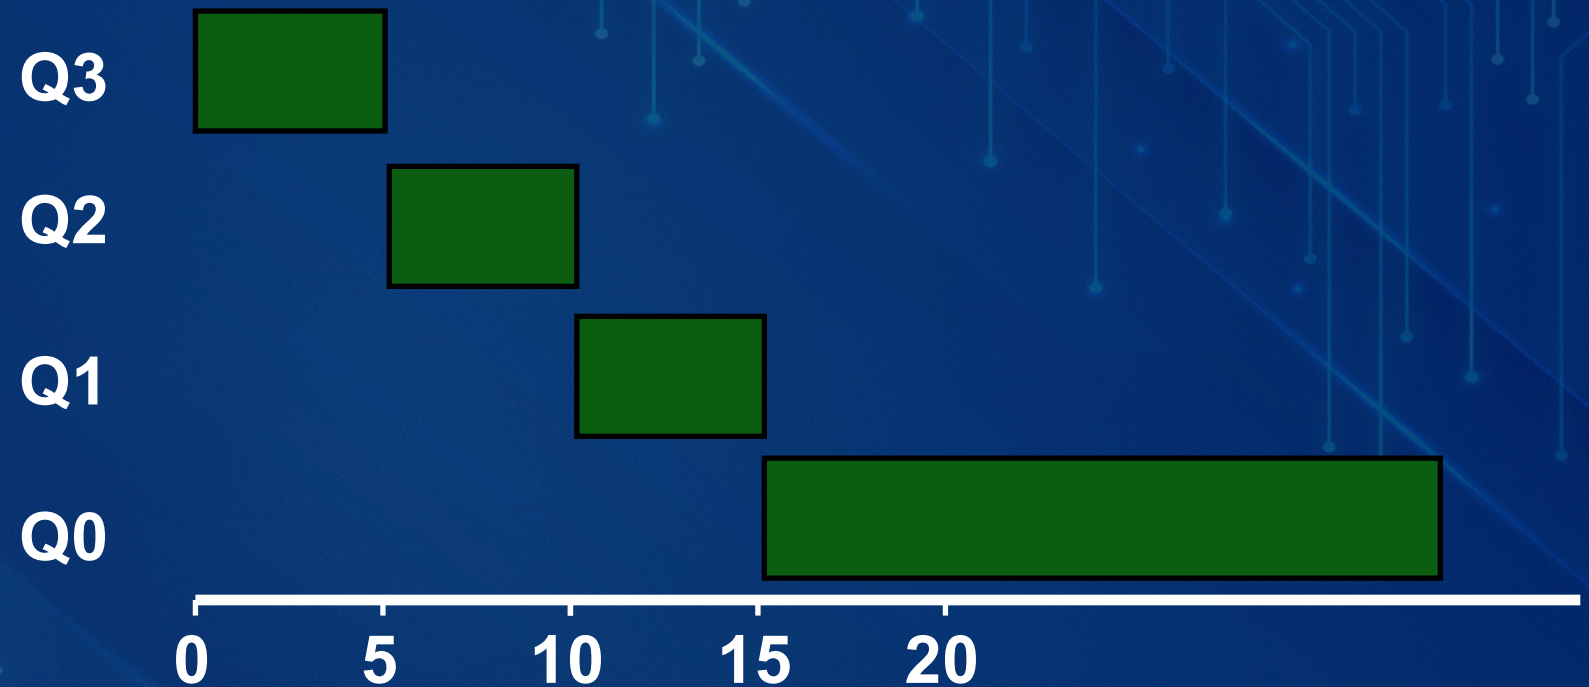
12

- Rule 1: If $\text{priority}(A) > \text{Priority}(B)$,
A runs
- Rule 2: If $\text{priority}(A) == \text{Priority}(B)$,
A & B run in RR
- Rule 3: Processes start at top priority (highest queue)
- Rule 4: If job uses its whole slice, demote process,
 - else leave at same level.
 - What does it mean if this case occurs?
 - Process is **not** doing I/O



One Long Job – MLFQ Example

I3




- Process uses entire time slice. What happens?
 - Drops in priority...
 - Ends up at the bottom of the priority queue...

Example With Two Processes - MLFQ

14

Notes:

- *10 ms time slicing
- Q3 == top priority




Red (Interactive) and Green (Batch) Process.

- Based on this diagram, what can you determine about the red / green processes?
 - Red is an interactive process (using I/O) - thus never uses its entire CPU time slice.
 - Interactive processes never use entire time slice, so are not demoted
 - Why isn't the red process allowed to use its entire time slice on the CPU?
 - Because it blocks on I/O and the Scheduler puts another process onto the CPU (while the red process waits for the I/O to finish).

Problems with MLFQ?

15




- Problems
 - Long running (ie: batch) jobs will never receive CPU time (**starvation**)
 - This is not apparent in the 2-process case, but imagine having 10s of “Red” processes...
 - Smart users could **game** the scheduler.
 - Just yield the CPU right before the time slice is about to end.

Prevent Starvation

16

Interactive Processes:
Blue, Gold

CPU Intensive Process:
Green



Problem: Low priority job may never get scheduled
Periodically boost priority of all jobs
(or all jobs that haven't been scheduled; aging)

Prevent Gaming

17

- Job can yield just before time slice ends to retain CPU
- Fix: Account for job's total run time at priority level
 - Instead of just this time slice;
 - downgrade when threshold exceeded

Updated MLFQ Rules

- Rule 1: If $\text{priority}(A) > \text{Priority}(B)$, A runs
- Rule 2: If $\text{priority}(A) == \text{Priority}(B)$, A & B run in RR
- Rule 3: Processes start at top priority (highest queue)
- Rule 4: Once a job uses allotment at given level, it moves down a queue (prevents gaming).
- Rule 5: After some time period, move (boost) all jobs to the topmost queue and repeat (prevents starvation).

Lottery Scheduling

- Goal: proportional (fair) share, but allow for weights
- Approach:
 - Give processes lottery tickets
 - Whoever wins runs
 - More tickets → higher priority/share
- Amazingly simple to implement

Lottery Scheduler Example

20

```
int counter = 0;
int winner = getrandom( 0, totaltickets );

node_t *current = head;
while( current ) { // Walk list to find winner
    counter += current->tickets;
    if (counter > winner) { break; }
    current = current->next;
}

// current gets to run
```



Scheduling Summary

- Understand goals (metrics) and workload, then design scheduler around that
- General purpose schedulers need to support processes with different goals
- Past behavior is good predictor of future behavior
- Random algorithms (lottery scheduling) can be simple to implement and can avoid corner cases.

Interprocess Communication (IPC)

*adapted from slides by Scott Brandt at UC Santa Cruz and other general sources

IPC Introduction

- Inter-Process Communication (IPC) enables processes to communicate with each other, to share information
 - **Pipes (half duplex)**
 - **FIFOs (named pipes)**
 - Stream pipes (full duplex)
 - Named stream pipes
 - Message queues
 - Semaphores
 - Shared Memory
 - Sockets (Can be on the same machine)
 - Signals

Pipes

24

- What is a (Linux/Shell) Pipe?
 - A communication channel (between processes*)
 - A connection between two file descriptors**
- **File Descriptors
 - An integer that refers to the location in the array of open files.
 - `stdin == 1, stdout == 2... What is #3?`
 - `stderr`
 - Everything in Unix (Linux) is a file...
 - Keyboard, Disk, Monitor, etc
- `dup2, pipe`
 - Unix commands to setup and manipulate file descriptors
 - `read()` / `write()` / `cin` / `cout`

Pipes

- Oldest (and perhaps simplest) form of UNIX IPC
- Half duplex
 - Data flows in only one direction
- Only usable between processes with a common ancestor
 - Usually parent-child
 - Also child-child

Pipes on the command line

- Chain output of one command into the input of another.
- **Syntax:** `command1 | command2`
 - `ls | head`
 - `w | grep days`
- In Unix, all processes in a pipe are started “simultaneously”.
- Pipes implement buffering under the hood if the read/write speeds of the two processes are different.

Pipes (cont.)

- `#include <unistd.h>`
- `int pipe(int fildes[2]);`
 - How is the param `fildes` passed?
 - Reference... What?
 - It is actually `int * fildes`
 - `fildes[0]` is open for **reading** (remember – half duplex)
 - `fildes[1]` is open for **writing**
- The output of `fildes[1]` is the input for `fildes[0]`

Understanding Pipes

- Within a single process
 - Data written to `filides[1]` can be read on `filides[0]`
 - Not very useful
- Typically used between processes.

Pipes and fork()

- Create a pipe.
- Then fork().
- Parent and Child processes each have a copy of the pipe... however, it can only be used to send data in one direction.
 - Parent and Child must decide which one will write to the pipe, and which one will read from it.
- Pipes are half-duplex.

Pipes and fork()

- To send messages from parent to child:
 - Have parent close the read end of its pipe:
 - `close(filedes[0]);`
 - Have child close the write end of its pipe:
 - `close(filedes[1]);`
 - Use `write()` in the parent to write to the pipe. Write to `filedes[1]`.
 - Use `read()` in the child to read from the pipe. Read from `filedes[0]`.

Summary

- IPC allows processes to communicate.
- One mechanism is via pipes.
- Pipes can be used programmatically via `pipe()` and `fork()` to allow IPC between parent and child.
- Pipes can be used on the command line to chain commands together without using intermediate files.

~ Fin ~