

Systems I – CS 60I3

Computer Architecture and Operating Systems

Lecture I6: Virtual Memory, Part 2

MASTER OF SOFTWARE DEVELOPMENT (MSD) PROGRAM

J. DAVISON DE ST. GERMAIN

SPRING 2024

*(adapted from slides by Ryan Stutsman, Andrea Arpaci-Dusseau, and Sarah Diesburg, and MSD presentations)

Lecture 16 – Topics

2

- Address Translation
 - Base / Bound
 - Segmentation

Announcements / Questions


3

- Unix Shell?
 - Check return codes on system calls
 - Everyone on track to have it done by Friday?

Unix Shell

4

- The user has typed “ls -l | nl | tail”
 - Draw a process bubble picture that includes processes that exist, and every piece of information they are currently storing.
- Now, at the beginning (1st or 2nd line) of `getCommands()`, draw a new bubble diagram that shows all data that now exists...



Unix Shell

```
ls -l | nl | tail
```

5

- Beginning of `getCommands()`
 - Note, 0 == Std In FD
 - 1 == Std Out FD
- Now draw the bubble diagram detailing the way things look at the end of `getCommands()`

Shell

`getCommands():`

```
vector<struct Command> cmds = [  
    ["", [], 0, 0, False],  
    ["", [], 0, 0, False],  
    ["", [], 0, 0, False]  
]
```

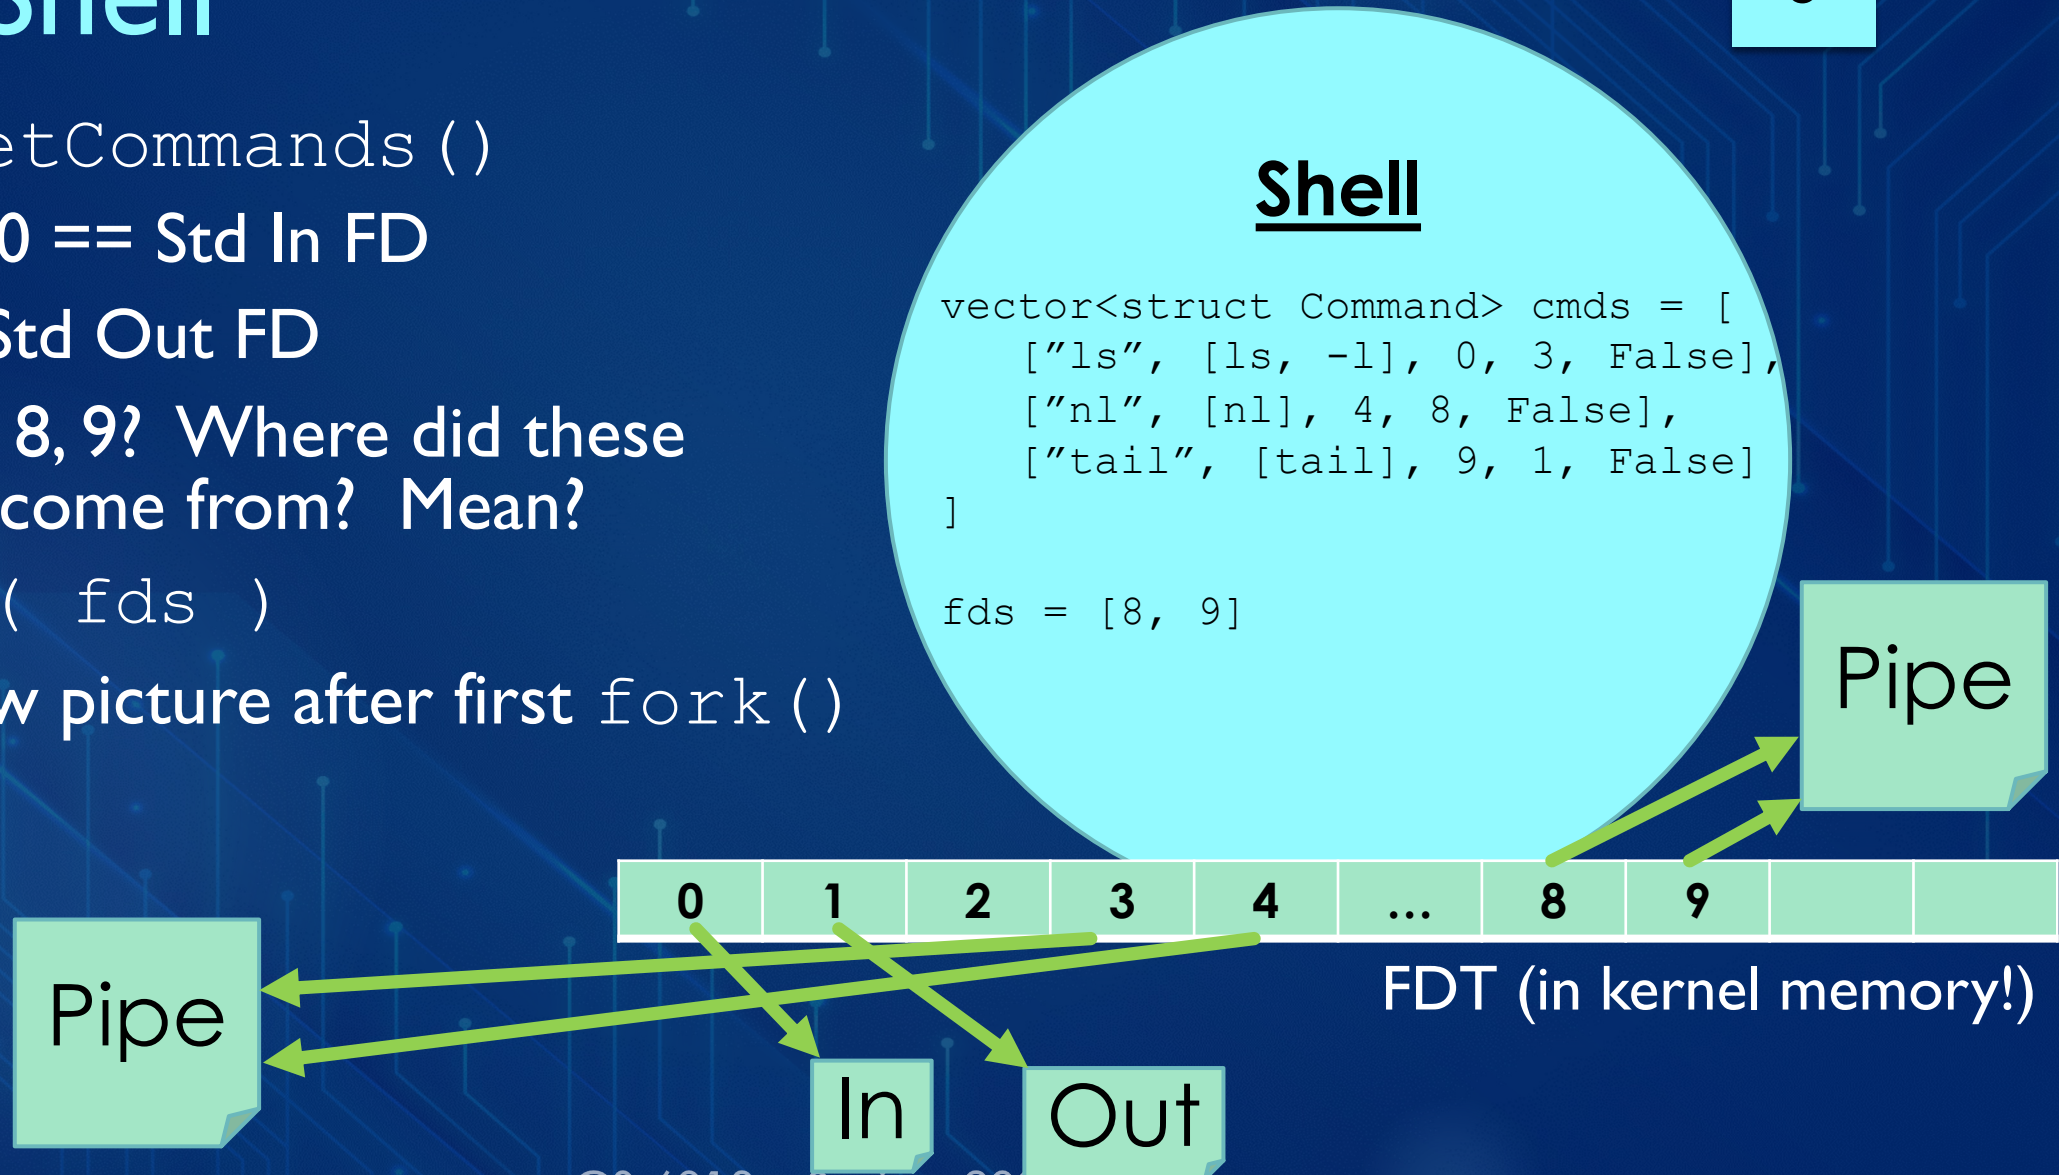


Unix Shell

```
ls -l | nl | tail
```

6

- End of `getCommands()`
 - Note, 0 == Std In FD
 - 1 == Std Out FD
- Why 3, 4, 8, 9? Where did these numbers come from? Mean?
 - `pipe(fds)`
- Now draw picture after first `fork()`



Unix Shell

7

```
ls -l | nl | tail
```

- After first fork():
- Where is “ls” process?
- Not here yet...
- Now draw picture before exec...

shell

```
vector<struct Command>  
cmds = [  
  ["ls", "ls, -l", 0, 3, False],  
  ["nl", "nl", 4, 8, False],  
  ["tail", "tail", 9, 1, False] ]
```

fds = [8, 9]

...

exec()

In Out

Shell

```
vector<struct Command>  
cmds = [  
  ["ls", "ls, -l", 0, 3, False],  
  ["nl", "nl", 4, 8, False],  
  ["tail", "tail", 9, 1, False] ]
```

fds = [8, 9]

0 1 2 3 4 ... 8 9

FDT

0 1 2 3 4 ... 8 9

FDT

Pipe

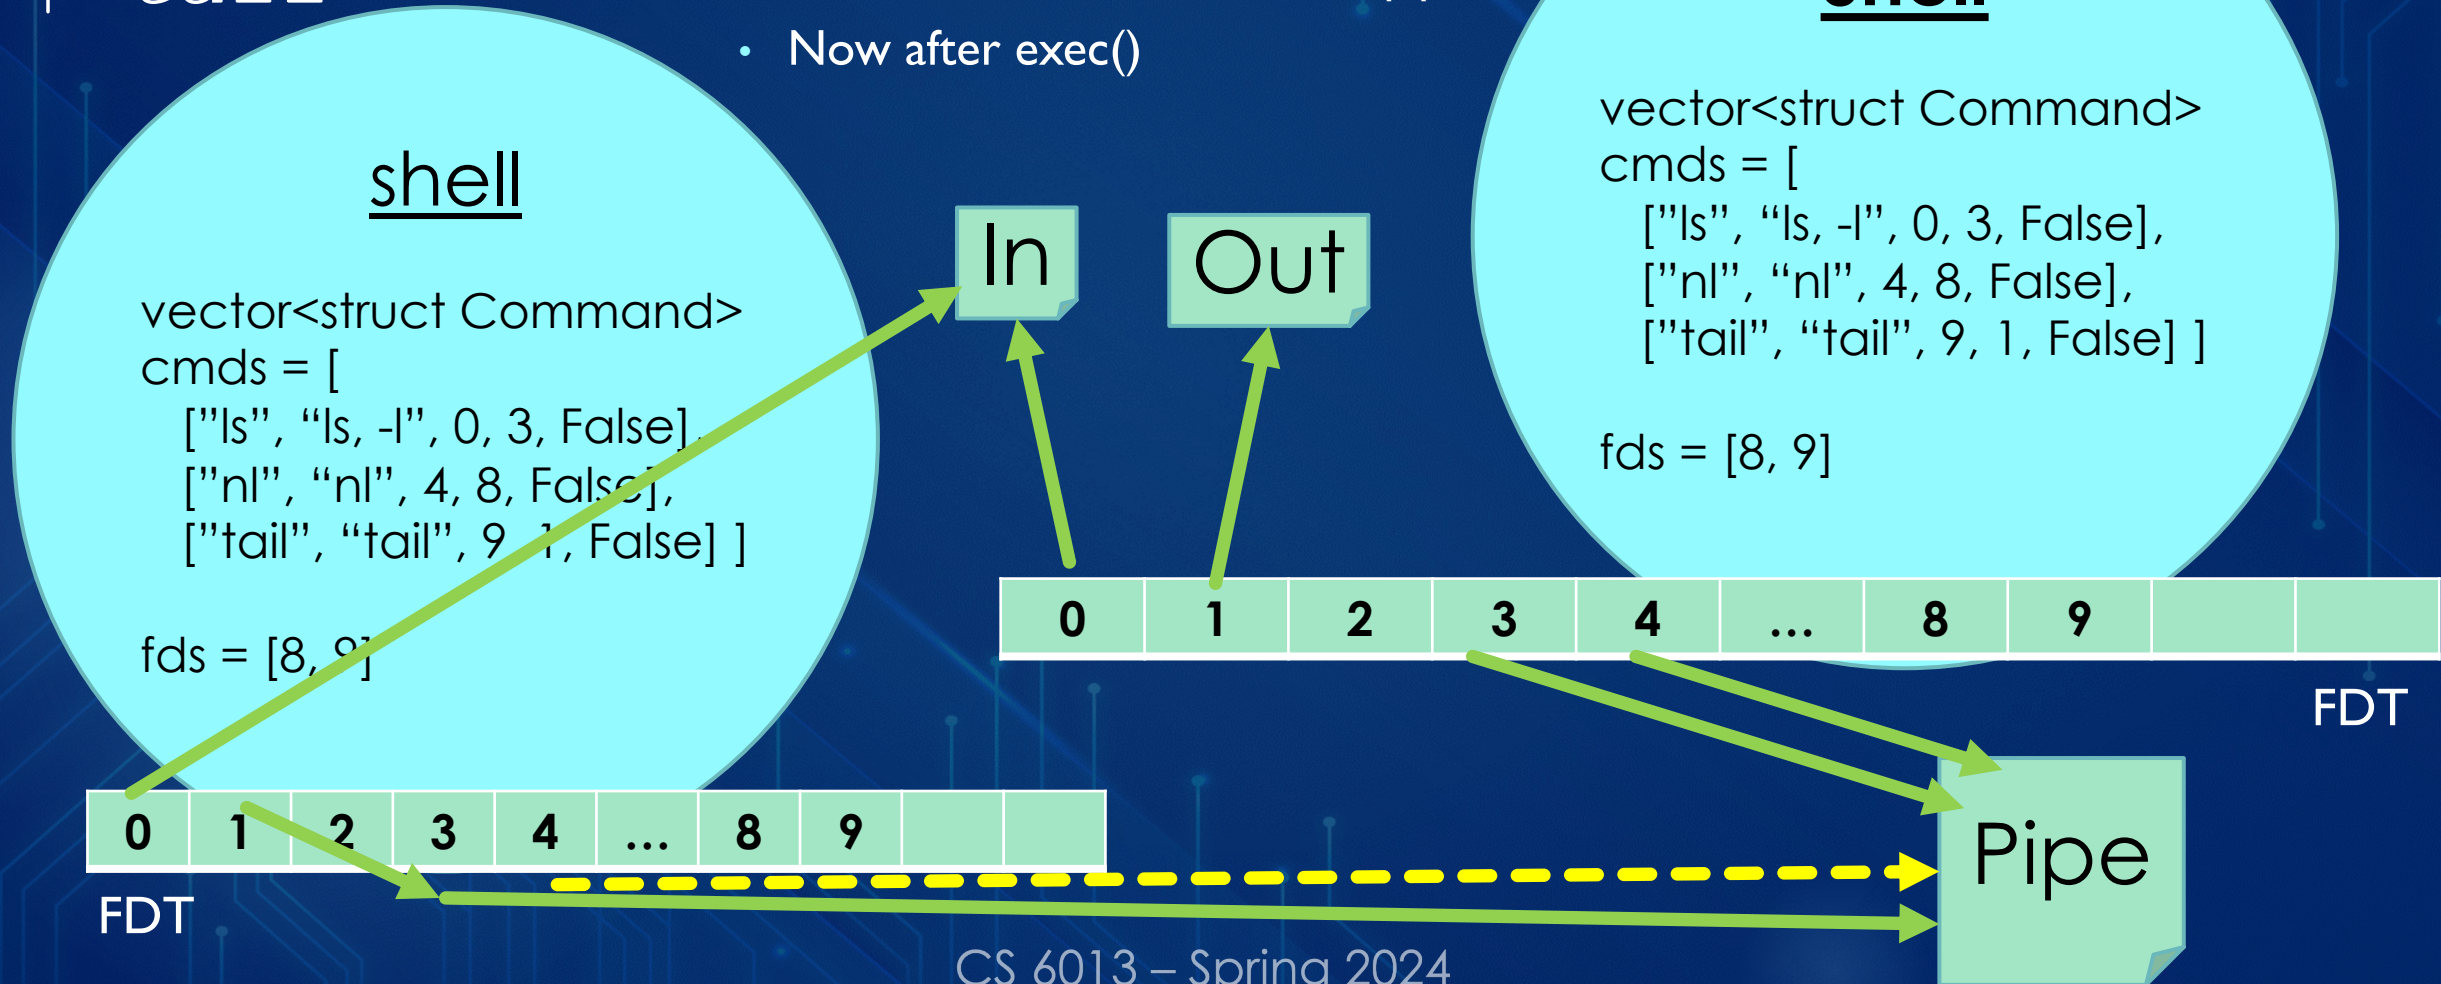
Pipe

Unix Shell

8

```
ls -l | nl  
| tail
```

- Before exec()
 - Why / how is 1 pointing to 3?
 - 1 is "redirected" to 3 via dup2
 - FD #4 (in new shell) should be closed because (soon to be) `ls` doesn't need the read side of the pipe...
- Now after exec()

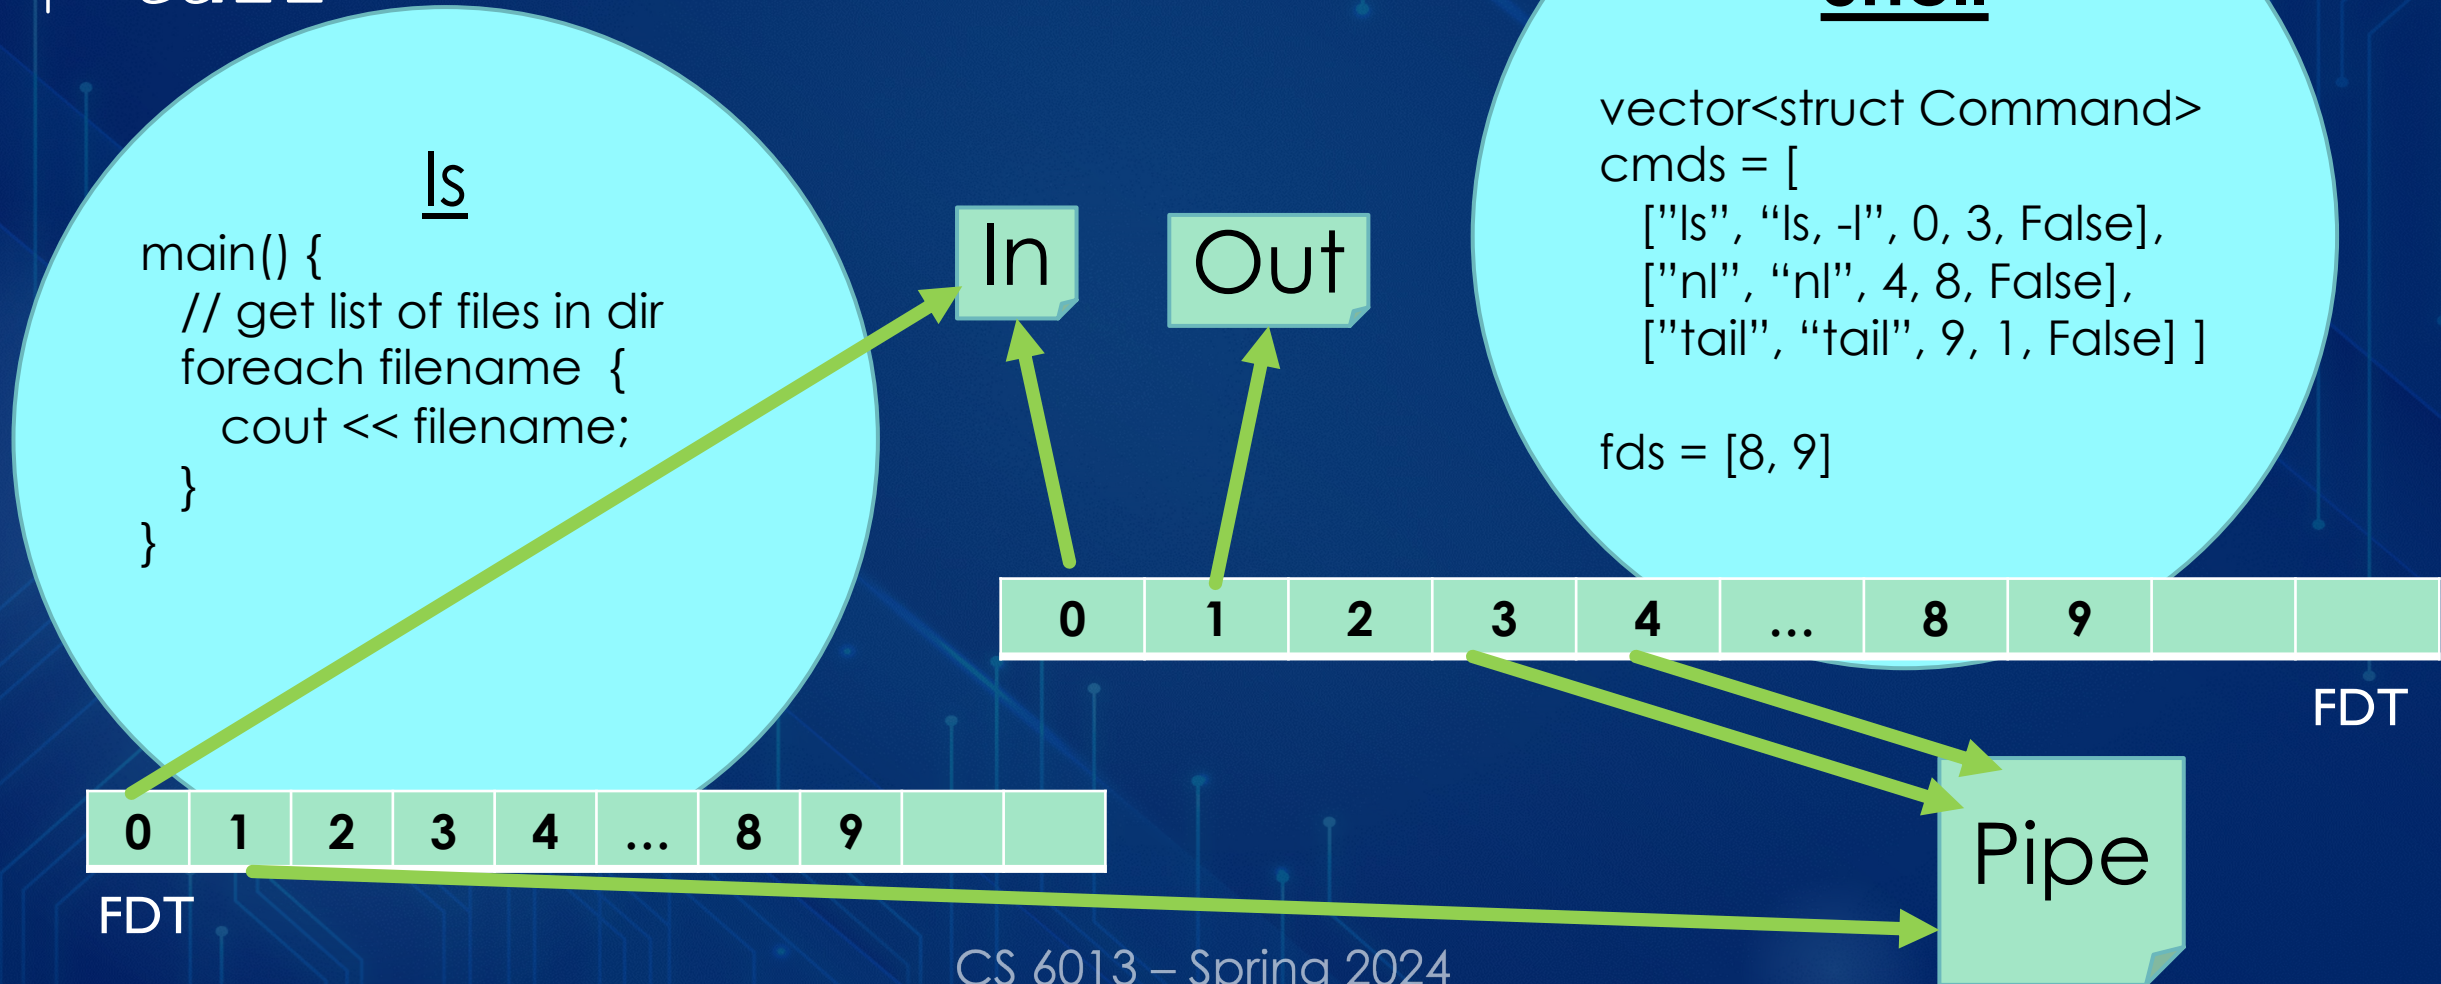


Unix Shell

9

- After exec()
- Now what happens?

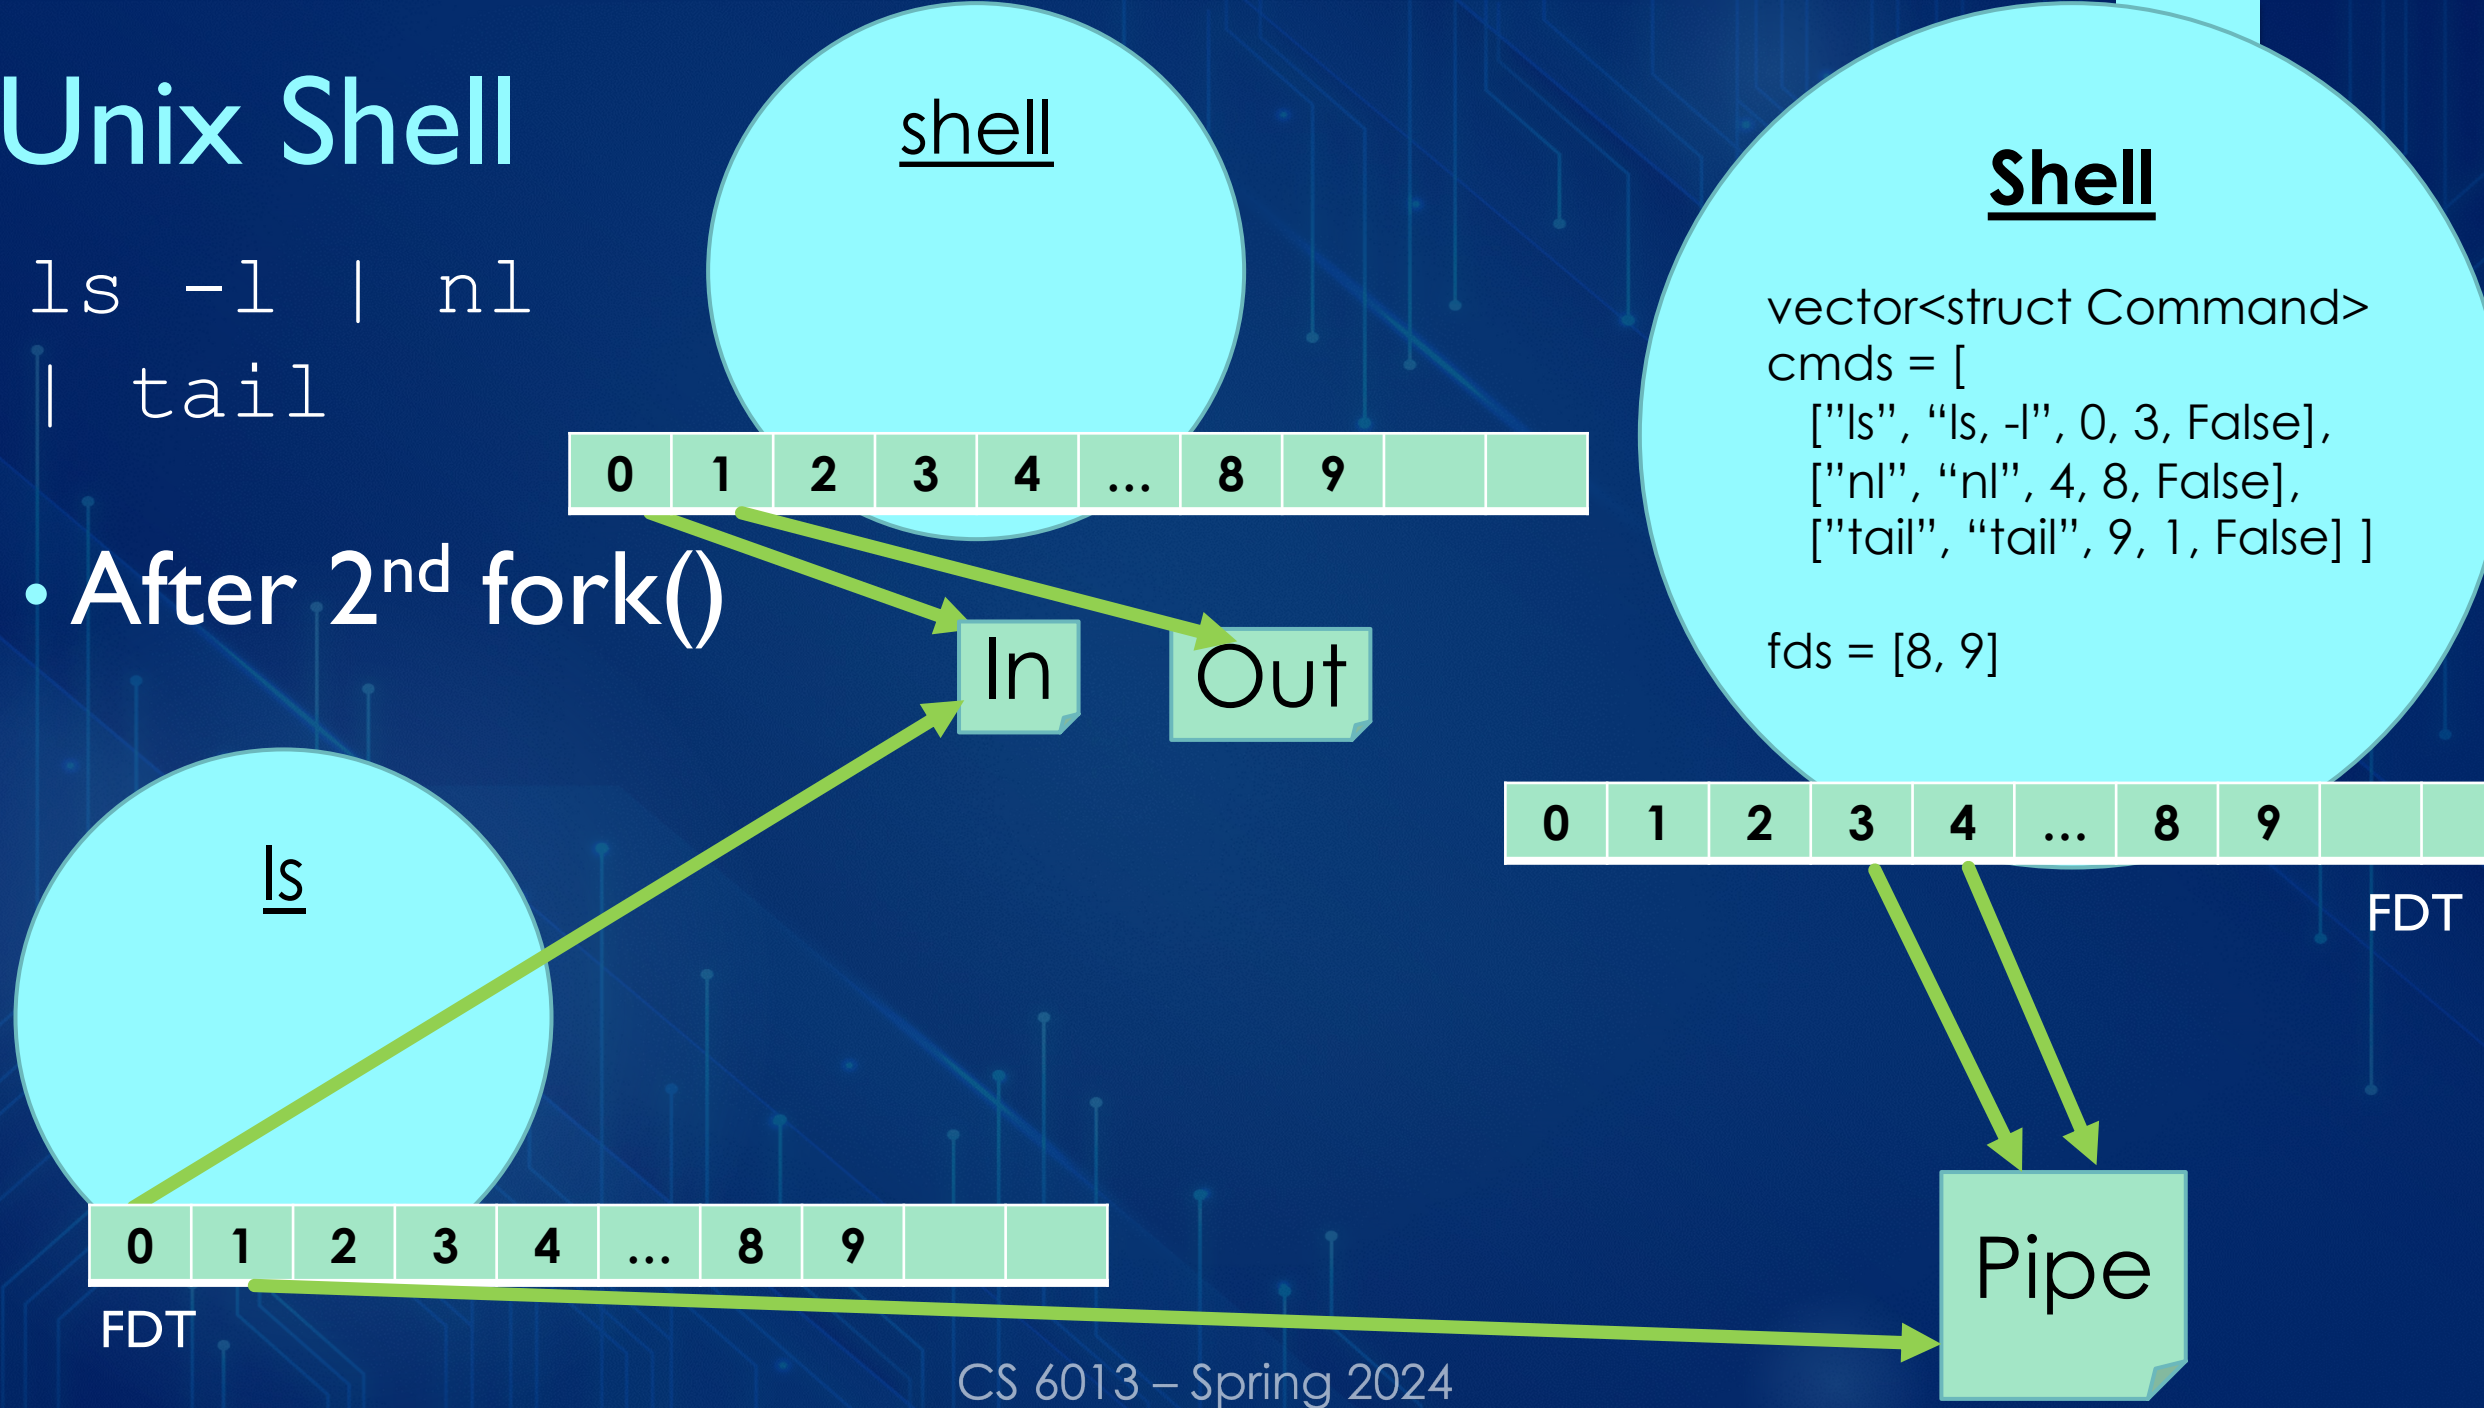
```
ls -l | nl  
| tail
```



Unix Shell

```
ls -l | nl  
| tail
```


- After 2nd fork()



Unix Shell

```
ls -l | nl  
| tail
```


- Before 2nd exec()



Unix Shell

```
ls -l | nl  
| tail
```

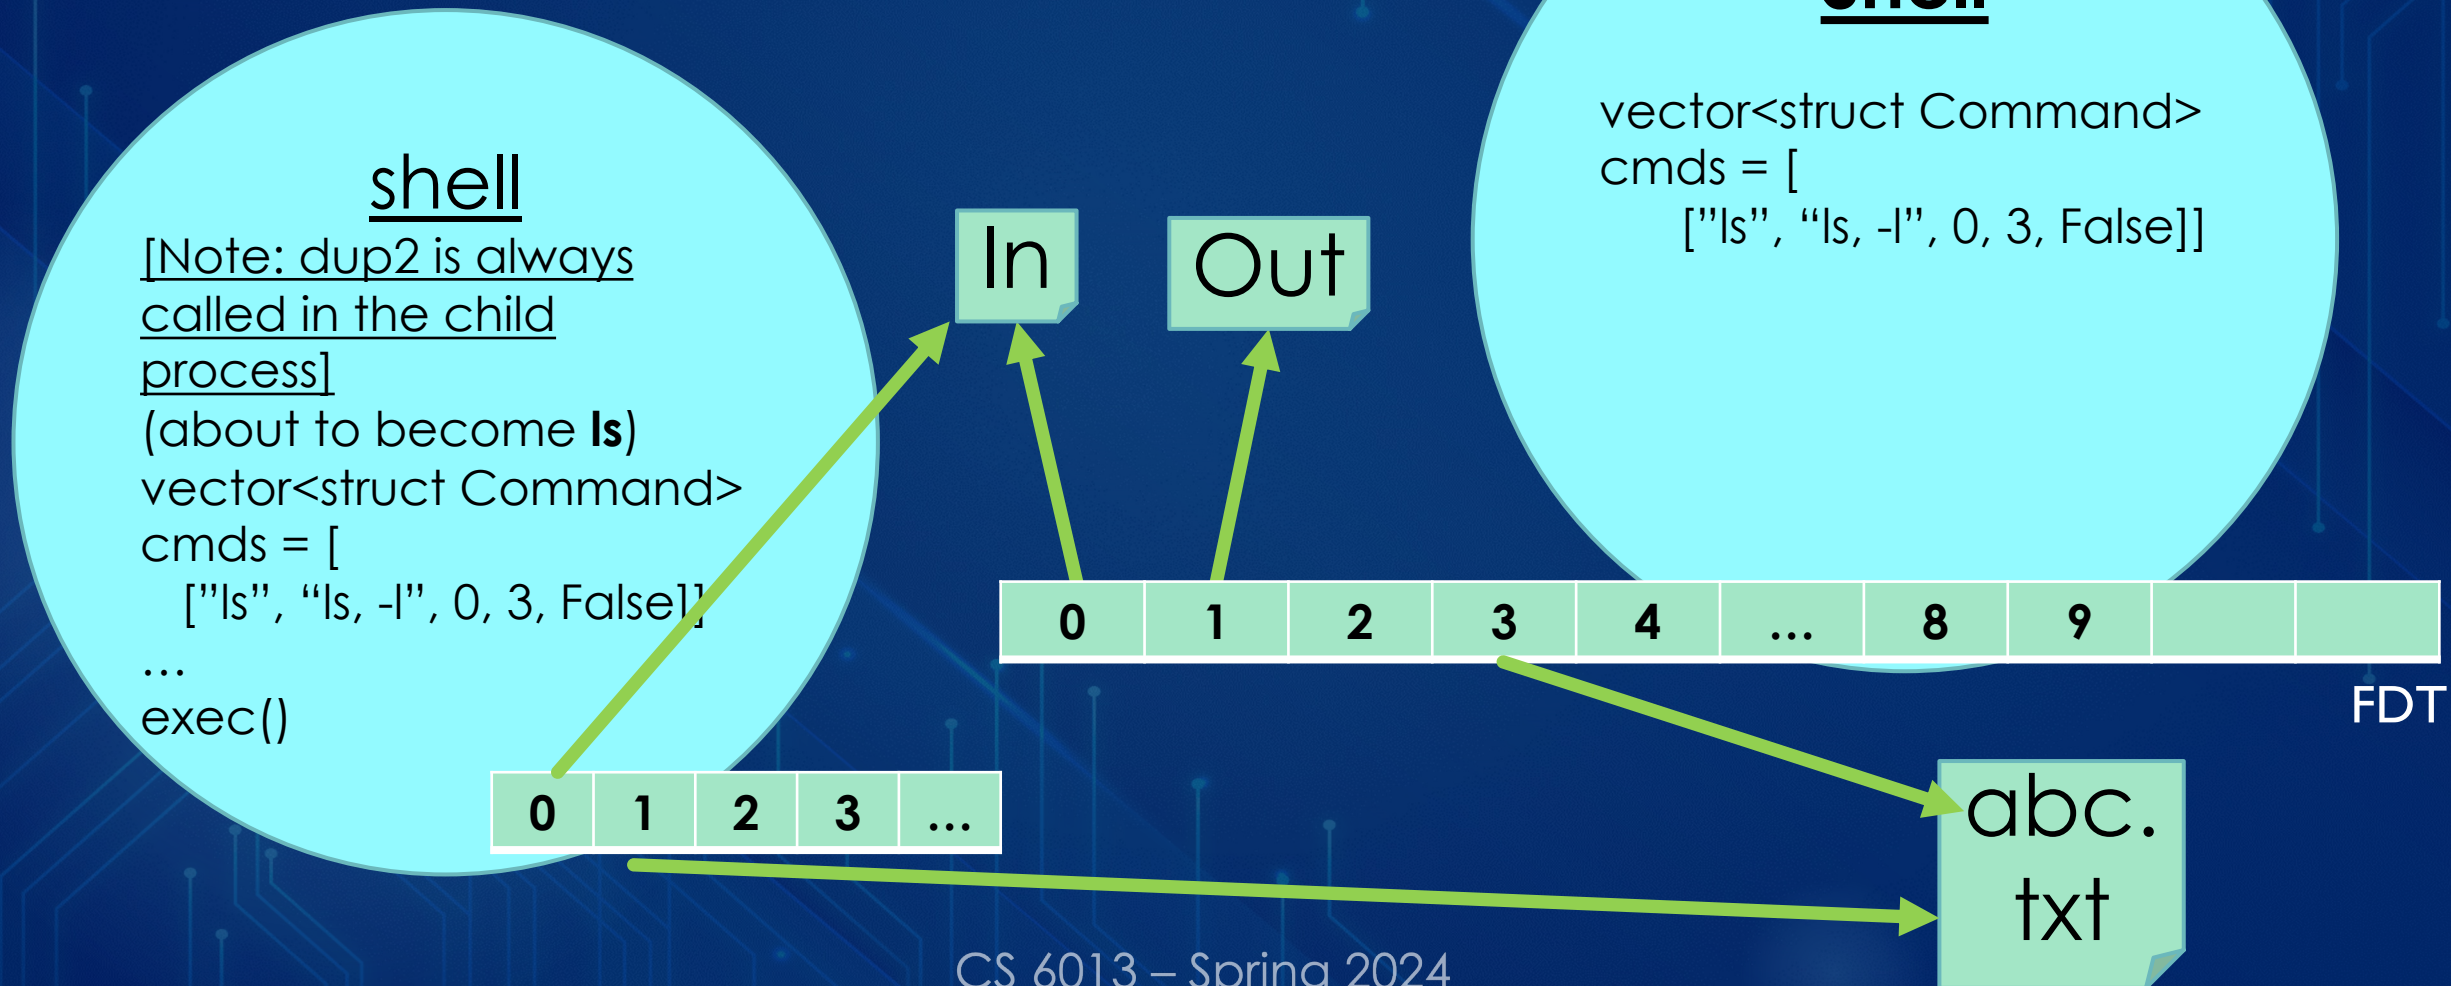
- After 2nd exec()



- What would “ls -l > abc.txt” look like after `getCommands()` and `fork()` and just before `exec()`?

Unix Shell

13



Address Translation (v1)

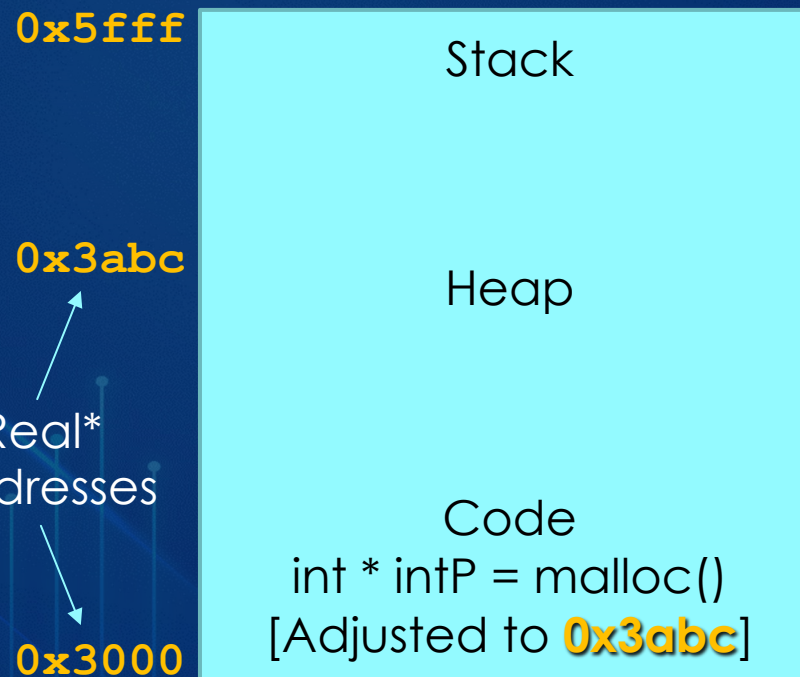

14

- Static Relocation

- At **load** (run) time, OS updates all addresses based on a base location.
- Say **+0x3000** (Max Process Memory Size)

Program [as written / compiled]

Process (P4) [when running]



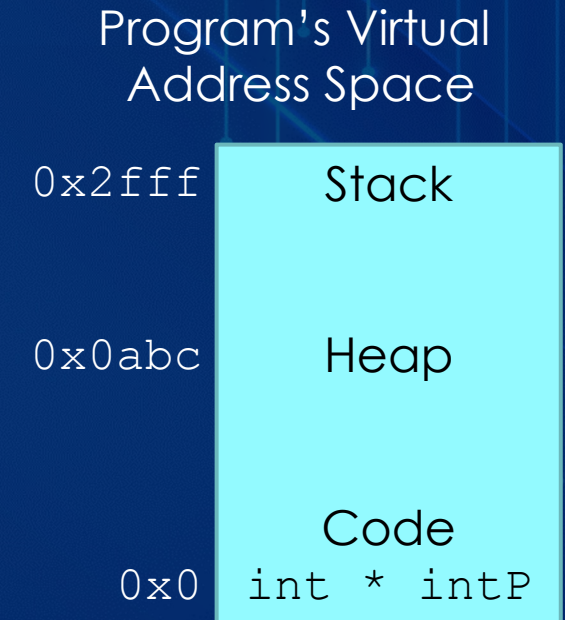
Physical Memory	
0x18000	P2
0x15000	P3
0x12000	
0x9000	
0x6000	P1
0x3000	P4
0x0	OS

- How is real memory arranged?
 - What's at 0x0?
 - OS
 - All other processes are slotted in.
- And P4?

Base and Bound Approach (Translation v2)

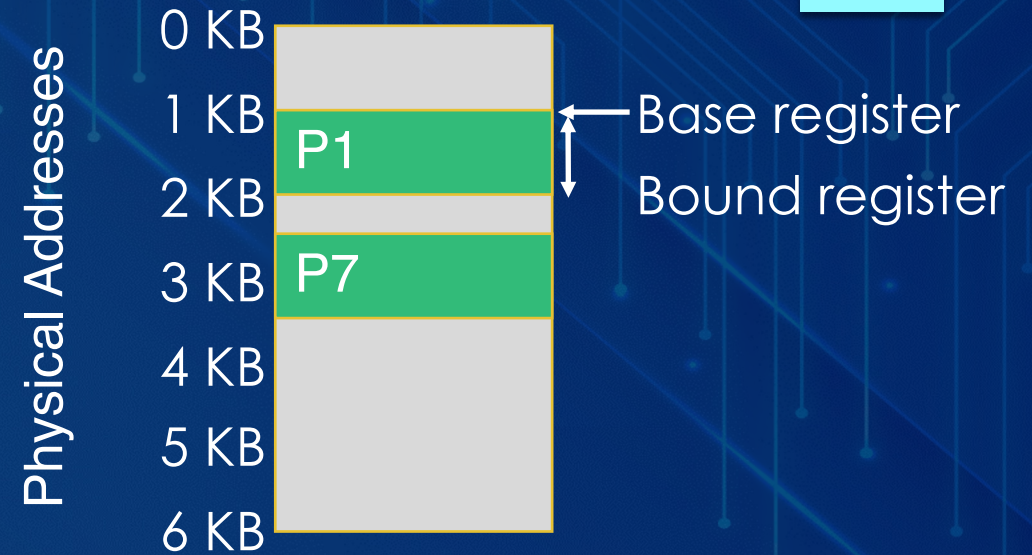
15

- Fairly straightforward and easy to implement!
 - Similar to the Static Relocation – but done “on the fly”
- OS sets up address translation with privileged registers (cr3 and others) in CPU
 - Indicates where in physical memory the address space of the process starts (**base**)
 - And to what physical memory it extends (**bound**)
- All addresses are automatically translated by MMU
 - Hardware does the translation (ie: adds **base** (cr3) to every address)
- Of course, a user-mode process cannot change these registers
 - The OS changes the values of the base/bound register at context switch



Base & Bound Example

I6




P1 is running

- Where are the base and bound values stored?
 - In registers (in the MMU – which is part of the CPU), which are only accessible by the OS (ie, when in kernel mode)

Base & Bound Implementation

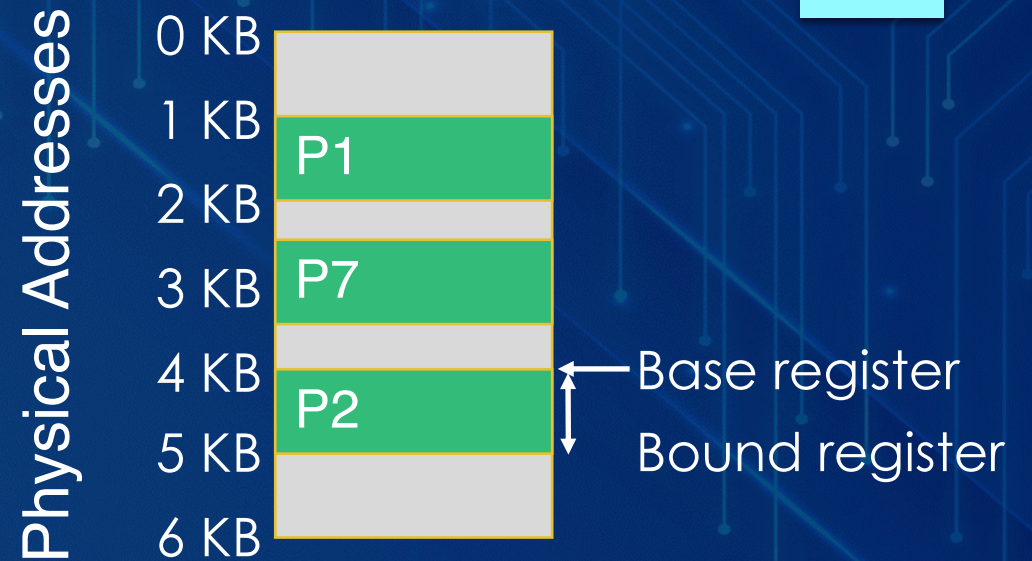
- On every memory access of process
 - MMU compares logical address to bounds register
 - if logical address is greater (or *), then generate error
 - MMU adds base register to logical address to form physical address

* or less than base...



Base & Bound Example

18




P2 starts running

- How (in general terms) does P2 to start running?
 - Context switch – which registers (in this situation) must be updated?
 - Base / Bound Registers (and of course, all the others)
- FYI, what does KB mean?
 - Kilobyte (1024 bytes)

Base & Bound Example

19

Can P1 access P2 memory?




	Virtual	Physical
P1:	load R1, [100]	load R1, [1124]
P2:	load R1, [100]	load R1, [4196]
P2:	load R1, [1000]	load R1, [5096]
P1:	load R1, [1000]	load R1, [2024]
P1:	load R1, [3072]	Illegal bound error

- What happens to this virtual instruction?
 - Why did we add 1124?
 - This is the **base** register's value.
- What makes this happen?
 - MMU and base/bounds registers

Base and Bounds Advantages

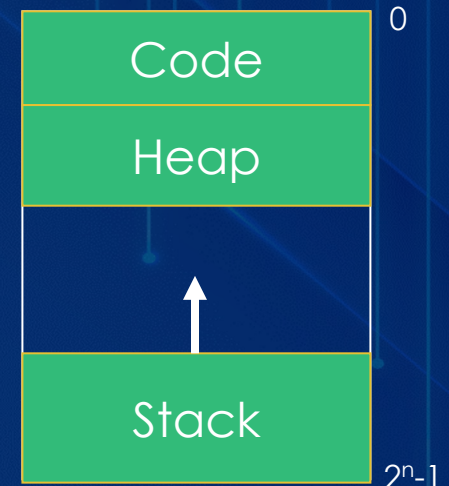
- Advantages
 - Protection (access/no access) across address spaces
 - Supports dynamic relocation* (see disadvantages)
 - Simple, inexpensive implementation
 - Only need a few registers and a little logic in MMU
 - Fast
 - Add and compare in parallel

Physical
Addresses




Base and Bounds Problems

- Disadvantages
 - Each process must be allocated contiguously in physical memory
 - Must allocate memory that may not be used by process (free space between stack and heap in Figure below)
 - I.e., “empty space” is mapped to physical memory – even though it isn’t being used.
 - No partial sharing: cannot share limited parts of address space
 - Not used because of its inflexibility... but is a good starting point for building a reasonable memory management system.



Segmentation


- Divide address space into logical **segments**
 - Each corresponds to logical entity in address space
 - Code, Stack, Heap
- Each segment can independently:
 - Be placed (separately) in physical memory
 - Grow and shrink
 - Be protected (separate read/write/execute protection bits)



Segmented Addressing

23


- Process specifies segment and offset within segment
- How does process designate a particular segment?
 - Explicit: use part of virtual address
 - Top bits select segment (physical piece of real memory)
 - Low bits indicate offset within segment
- Each segment gets its own base and bound!
 - Base gives starting address for a segment.
 - Bound tells us the limit of that segment in memory.
- How many bits to use as the segment number?
 - Number of segments allowed.
 - Total size of virtual address space, etc



Segmented Address

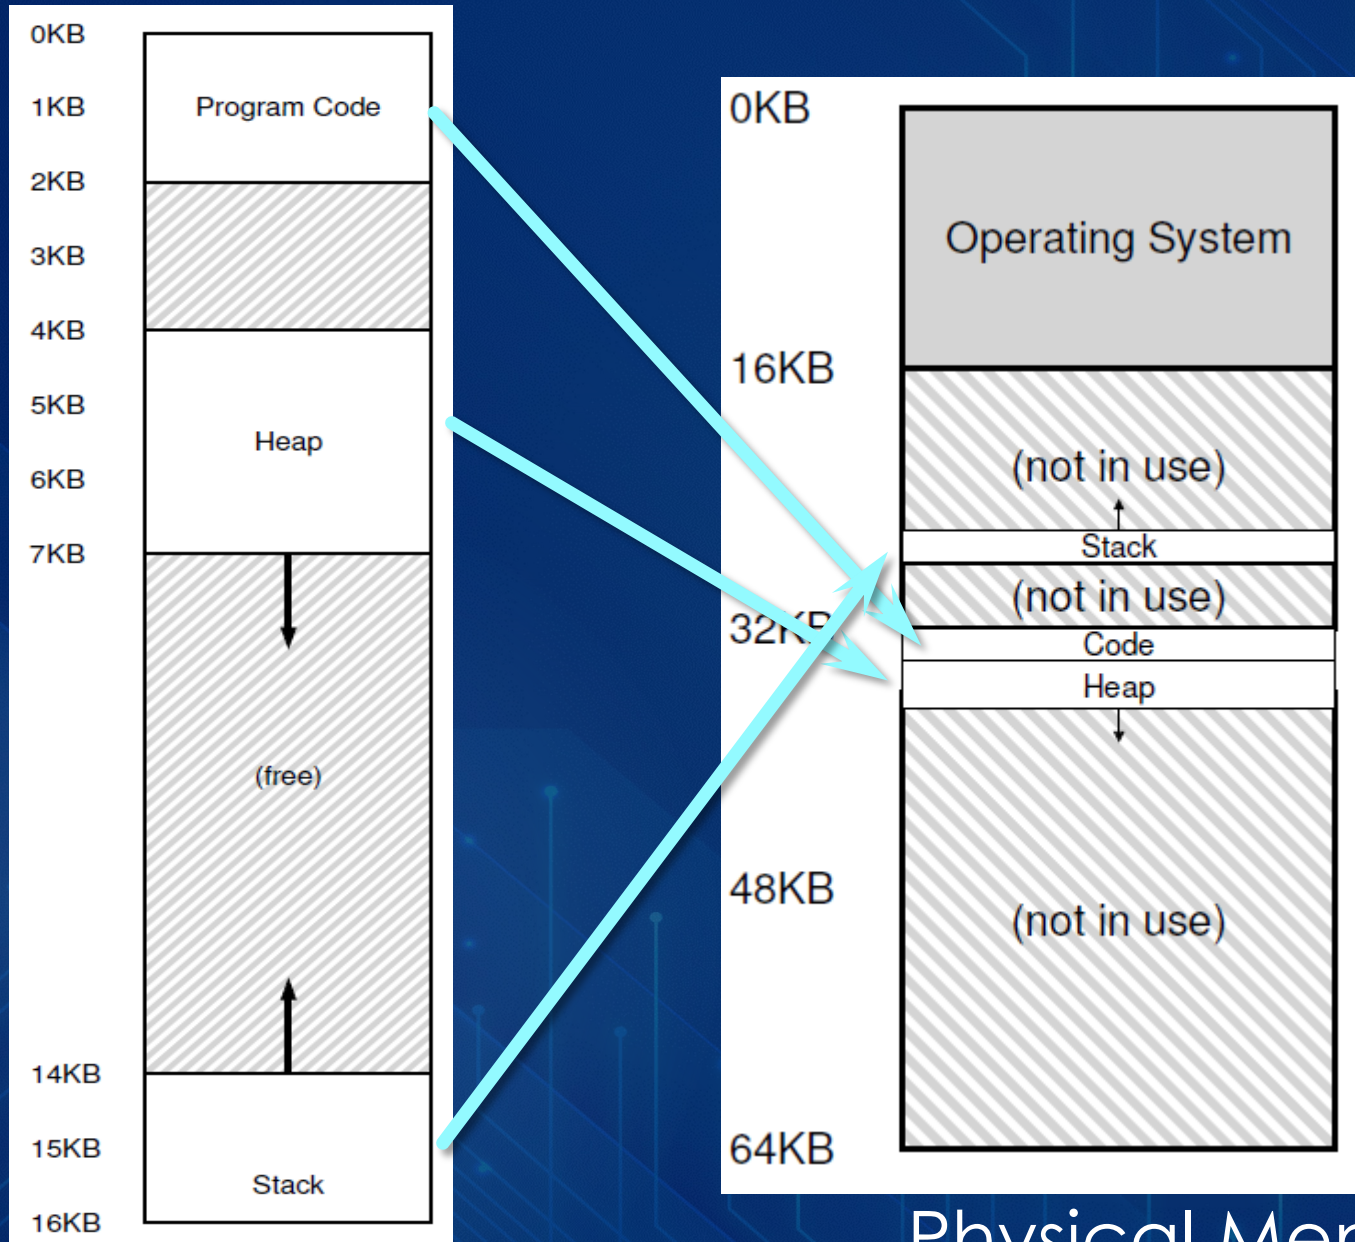
- Assuming we are using 14-bit addresses...
- Amount of memory a program could use?
 - $16\text{ K} - 2^{14}$
- Now assume using Segmentation, and physical memory is divided into (only) 4 segments...
 - How many bits to represent the Seg #?
 - 2
 - How big can a segment of memory be?
 - $4\text{ K} - 2^{12}$

Segmented
Virtual Address



Examples

25



Virtual address space (Process 1).

Physical Memory

CS 6013 – Spring 2024

- Each segment on the left is mapped onto the real memory on the right.
- Notice OS is taking up some physical memory
- How much real memory does Process 1's Code segment use?
 - 2 KB of real memory – same as the virtual memory is uses

Example 1

- Say code segment begins at 32KB in physical memory and is allotted 2KB.
- Consider the virtual address 100.
- What is the corresponding physical address?

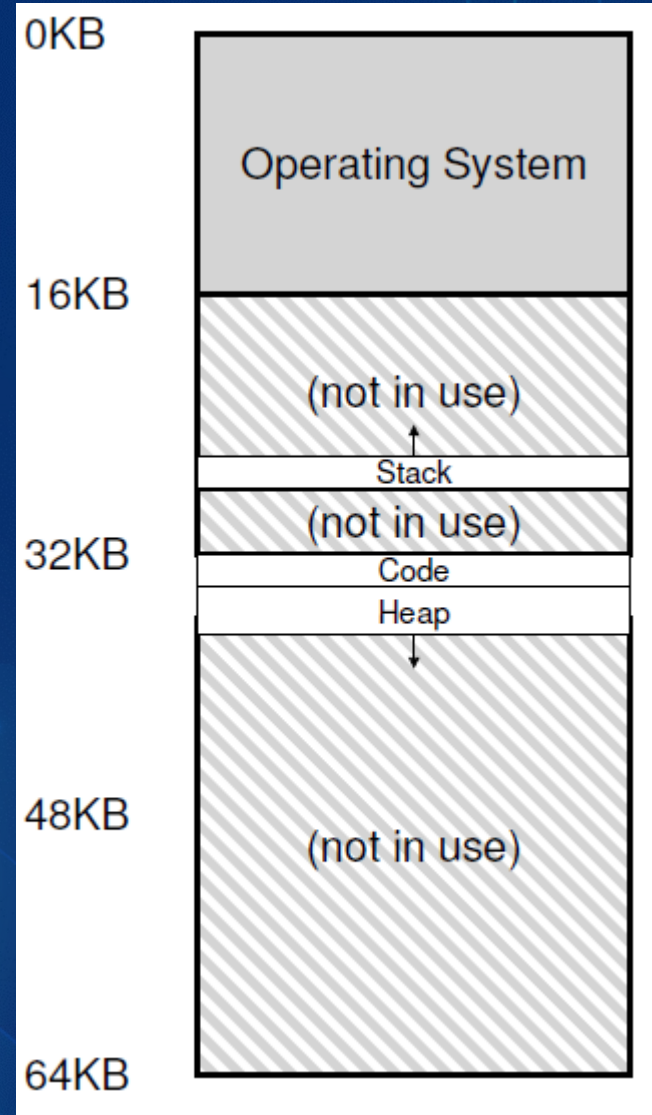

Example 1 - Answer

- Say code segment begins at 32KB in physical memory and is allotted 2KB.
- What is its physical address space?
 - 32KB to 34KB – 32768 to 34816
 - Note: $32 * 1 \text{ KB} == 32 * 1024 == 32768$
- What is physical address of virtual address 100?
 - Physical address = $32\text{KB} + 100 = 32868$.
- Note: this is still within bounds of $32\text{KB} + 2\text{KB}$.

Example 2

28

- Say heap segment starts at 34KB in physical memory, has size 3K.
- What is the **virtual** starting address of the heap segment?
 - Just looking at the figure on the left, we see the virtual starting address is 4096.
- What is the ending **physical** address of the heap?





Physical Memory

Virtual address space (Process 1).

Example 3

29

- What is the physical address of virtual address 4200?
 - This one is a little trickier.
- First find the offset into the heap.
 - This is $4200 - (\text{starting virtual address of heap})$.
 - That is, $4200 - 4096 = 104$.
- Then add this offset to the base physical address of heap (34KB).
 - Result = $34\text{KB} + 104\text{B} = 34816 + 104 = 34920$.



Physical Memory

Virtual address space (Process 1).

Segmentation Implementation

30

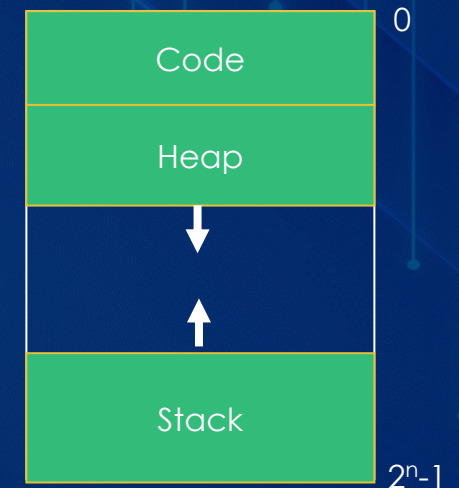
MMU has base/bound/permissions register per segment

- Example: 14-bit logical address, 4 segments
- How many bits for segment number?
- How many bits for offset?

Segment	Base	Bounds	R W
0	0x2000	0x6ff	1 0
1	0x0000	0x4ff	1 1
2	0x3000	0xfff	1 1
3	0x0000	0x000	0 0

Advantages of Segmentation

- Enables sparse allocation of address space
 - Stack and heap can grow independently
 - Heap: if no free memory, then `sbrk()`
- Protections for different segments
 - Read-only status for code
- Enables sharing of selected segments
- Supports dynamic relocation of each segment



Disadvantages of Segmentation

- External Fragmentation
 - Memory is full of little chunks of free space, but none of them are large enough for a single new segment.
 - Memory allocation requests may be denied.
- May still be really wasteful.
 - What if the heap is barely used, but it gets its own large segment?
 - Wasted space, contributes to fragmentation, and is not really fine-grained.

Conclusion

- HW+OS work together to virtualize memory
 - Give illusion of private address space to each process
- Add MMU registers for base+bounds so translation is fast
 - OS not involved with every address translation, only on context switch or errors
 - CPU has built-in hardware to do address mapping (translation)
- Dynamic relocation with segments is a good building block
 - Next lecture: Solve fragmentation with paging

~ Fin ~