

Systems I – CS 60I3

Computer Architecture and Operating Systems

Lecture 22: Malloc Assignment

MASTER OF SOFTWARE DEVELOPMENT (MSD) PROGRAM

J. DAVISON DE ST. GERMAIN

SPRING 2024

Lecture 22 – Topics

2

- Malloc Assignment
 - Structure Sizes
 - Malloc / Free
 - mmap
 - Hash Table

Announcements / Questions

3

- Week 9 Readings – Concurrency: OSTEP

Chapters:

- 25: Dialogue
- 26: Concurrency and Threads
- 27: Thread API
- 28: Locks

What is Malloc?

- `malloc()` allocates memory for us...
 - But it is only a function call...
 - Although a fairly advanced piece of code...
- But wait, if we are grabbing memory from the system, don't we need a system call?
 - Yes, to get raw blocks of memory, we use the `mmap()` system call.
 - However, `malloc()` does this behind the scenes for us so that we don't have to worry about it.

Storing an Address / Size in HT

```
int * i = malloc( sizeof( int ) )
```

```
double * d = malloc( 8 )
```

```
ht.insert( i, 4 );
```

```
ht.insert( d, 8 );
```

- **Declaration** of insert()?
- What is the *type* of a generic pointer?
 - void *
- What is the size of (all) pointers?
 - 64 (or 32) bits.
- How can we tell?
 - sizeof(void *)
 - g++ compilation flag: -m32

- So, insert() is?

```
void insert( void *, size )
```

- What is the *type* of size?

- int? long?
- How can we handle it regardless of 32 vs 64 bit system / compilation?
- Built-in type:
 - size_t

```
void insert( void * address,  
            size_t size );
```


man malloc

- **malloc**

- Defined in header `<stdlib.h>`
- `void* malloc(size_t size);`
- Allocates size bytes of **uninitialized** storage.
- If allocation succeeds, returns a pointer that is suitably **aligned** for any object type with fundamental alignment.
- If size is zero, the behavior of malloc is implementation-defined. For example, a null pointer may be returned. Alternatively, a non-null pointer may be returned; but such a pointer should not be dereferenced, and should be passed to free to avoid memory leaks.
- malloc is **thread-safe**: it behaves as though only accessing the memory locations visible through its argument, and not any static storage. (synchronized)

How big is a structure?

7

- How many bytes does struct Data use in memory?
- How can we ask the computer?
 - `sizeof(struct Data)`
- Let's figure it out ourself...

```
struct Data {  
    bool    b1;  
    int     i1;  
    bool    b2;  
    bool    b3;  
    void *  p;  
    bool    b4;  
    bool    b5;  
    bool    b6;  
    bool    b7;  
    bool    b8;  
    int     i2;  
};
```


How big is a structure?

- Size of each piece?
 - So the size is?
 - 24!
 - Well, no, it's 40...
 - What?
 - Let's continue by printing out the addresses...

```
struct Data {  
    bool    b1;    // 1      209808  
              // 3 bytes of padding  
    int     i1;    // 4      209812 <- Aligned!  
    bool    b2;    // 1      209816  
    bool    b3;    // 1      209817  
              // 6 bytes of padding  
    void *  p;     // 8      209824 <- Aligned!  
    bool    b4;    // 1      209832  
    bool    b5;    // 1      209833  
    bool    b6;    // 1      209834  
    bool    b7;    // 1      209835  
    bool    b8;    // 1      209836  
              // 3 bytes of padding  
    int     i2;    // 4      209840 <- Aligned!  
};
```


How big is a structure?

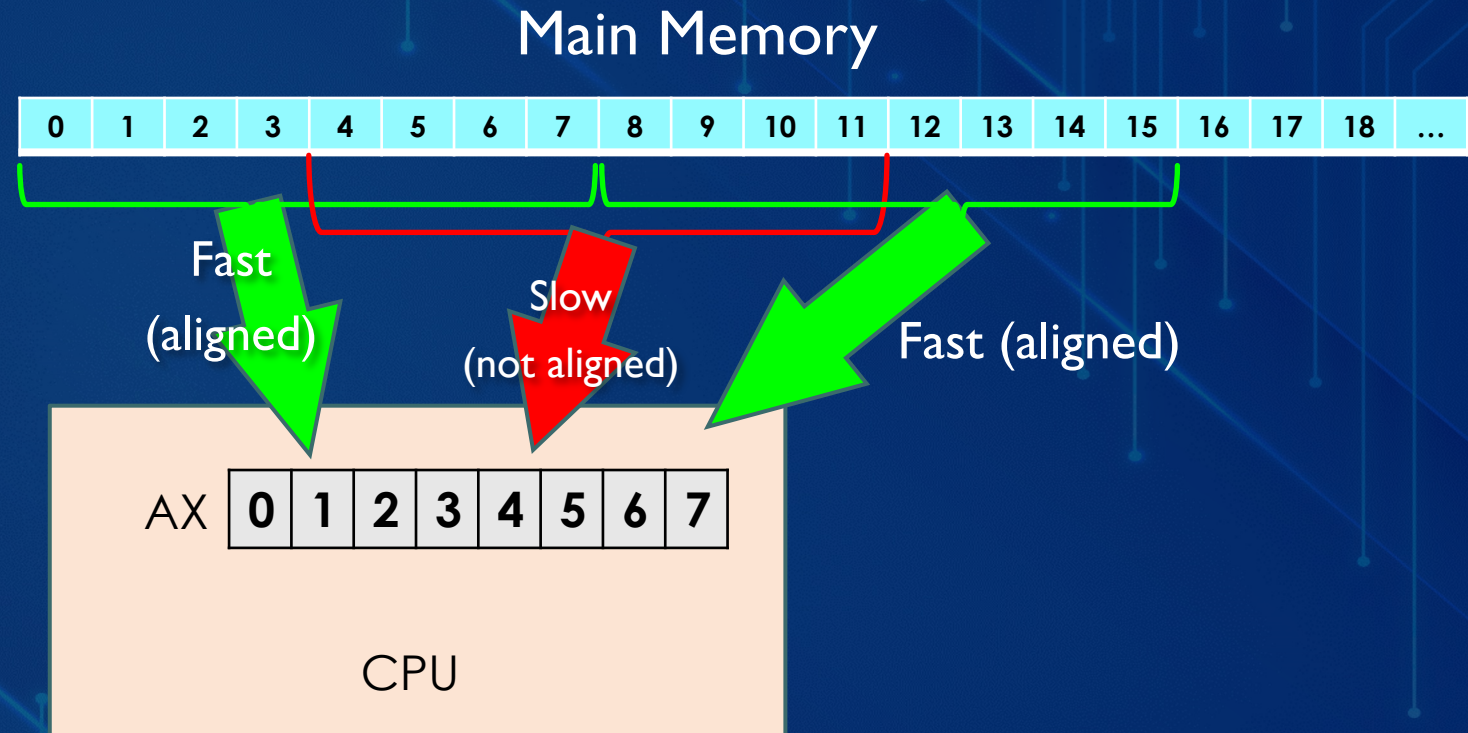
- Ah ha! The size is:
 - 36!
- Well, no...
 - What!?
 - Structs must be aligned based on the size of pointers...
 - 8, 16, 24, 32, 40, etc
- So...
 - Need 4 more bytes of padding... giving the final answer as:
- 40!

```
struct Data {  
    bool    b1;    // 1  
                // 3 bytes of padding  
    int     i1;    // 4    209812  
    bool    b2;    // 1    209816  
    bool    b3;    // 1    209817  
                // 6 bytes of padding  
    void *  p;    // 8    209824  
    bool    b4;    // 1    209832  
    bool    b5;    // 1    209833  
    bool    b6;    // 1    209834  
    bool    b7;    // 1    209835  
    bool    b8;    // 1    209836  
                // 3 bytes of padding  
    int     i2;    // 4    209840  
};                // 4 bytes of padding
```


Copying Data From Memory to CPU

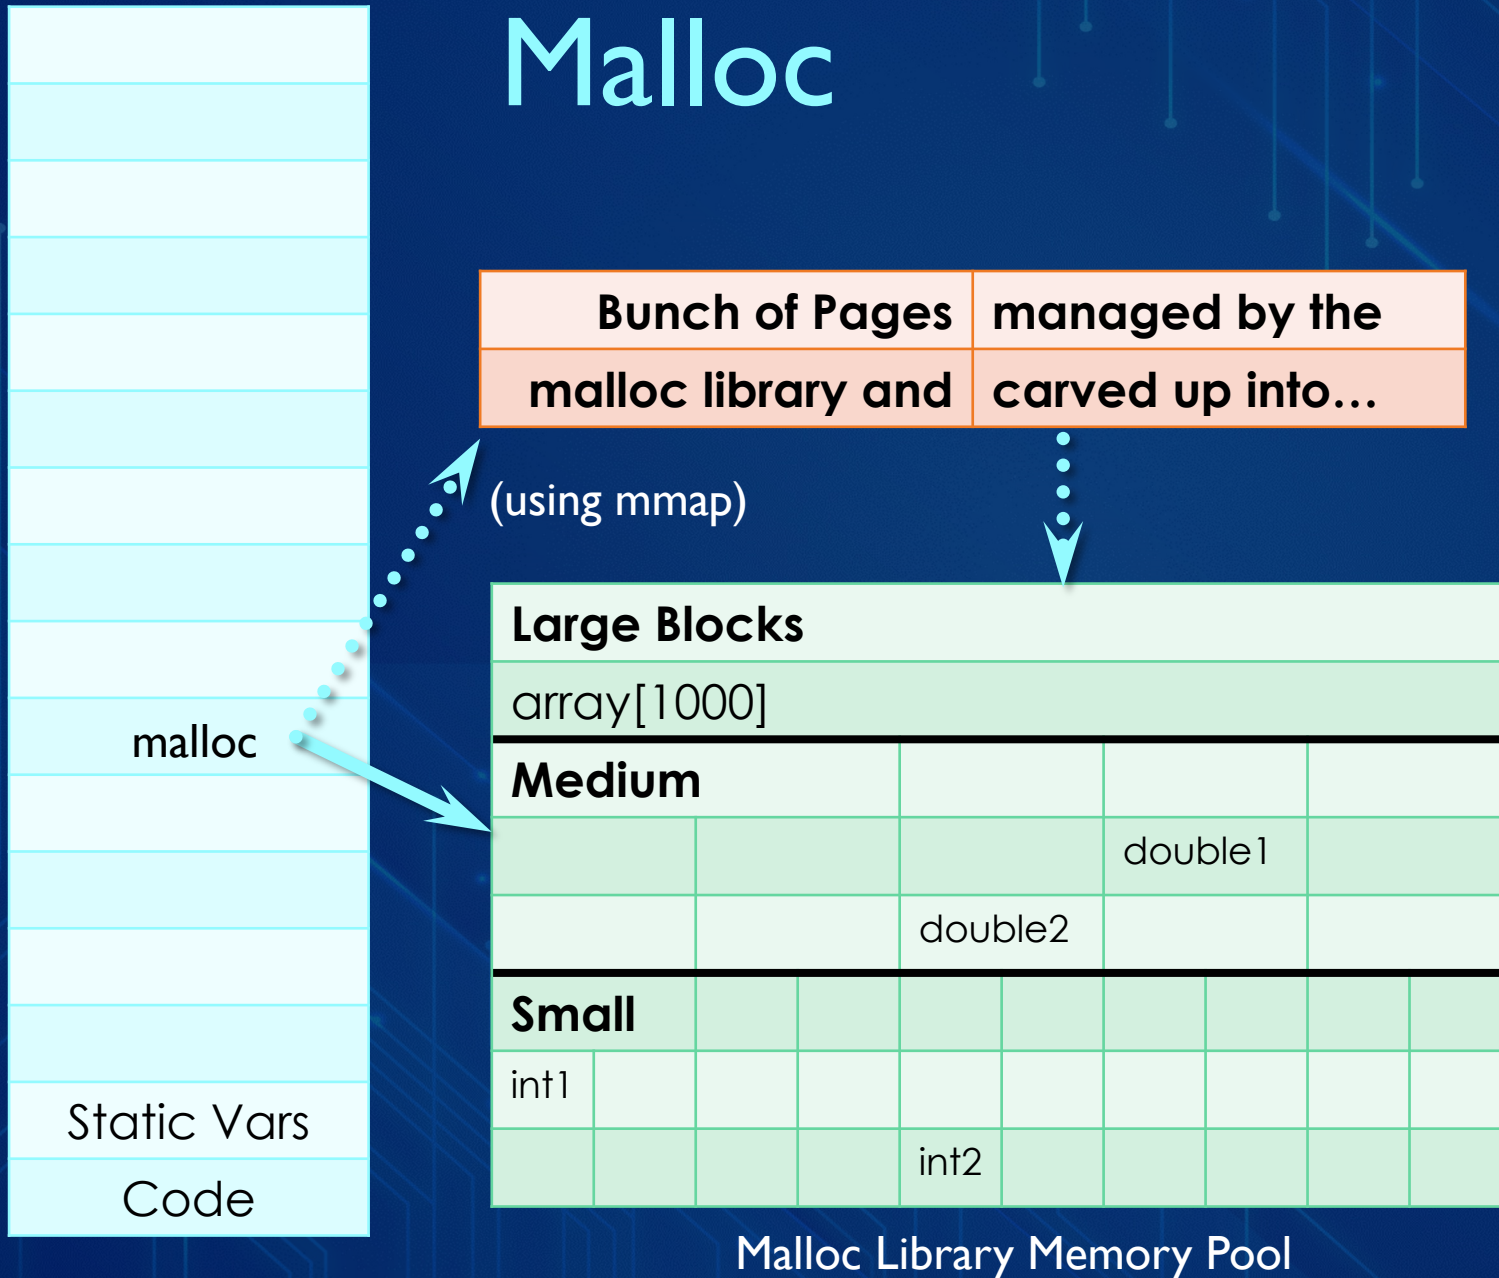
10

- `mov rax, [0]`
 - Copy 64 bits from address 0 to the rax register
- `mov rax, [4]`
 - Copy 64 bits of data from address 4 to the rax register
- Hardware dependent alignment requirements and effect on speed.



Malloc


II



- How does it work?
 - Malloc (library) grabs a bunch of memory and then handles it for you
 - ie, a bunch of pages that it gets via mmap
 - [Mmap gives you exactly what you asked for* and you manage it]
- So on first call to malloc (or more likely as the program initially begins to run)...
- Which the malloc library breaks up into...

Malloc

12



Large Blocks		✗	500 Bytes	
array[1000]				
Medium		✗		
			double1	
		double2		
Small				
int1	4 ✗			
		int2		

Malloc Library Memory Pool

```
int * p1 = malloc( 4 )
*p1 = 468;
double * p2 = malloc( 8 )
char * p3 = malloc( 500 )
```

- The malloc library remembers each of these addresses (and the size associated with them) so that when you call `free()`, it knows what to do.

- And as good memory citizens...

```
free( p1 )
free( p2 )
free( p3 )
```

- Malloc library has reclaimed that memory to provide to the program in the future.

mmap ()

- System call to manipulate address space
- Map anonymous pages to add heap space
 - Not really “adding” space to the heap because the virtual space already exists... mmap is more allocating / mapping the (currently empty) virtual address space to actual (physical) memory pages.

mmap()

14

- `void* mmap(void *addr, size_t len, int prot, int flags, int fd, off_t offset);`
- `addr` is the pointer you'd like to get back. If this is 0 (null), then OS will just give you a pointer to some chunk of memory. Alternatively, can use as a hint to OS about where the page boundary starts.
- `len` is the number of bytes you want to allocate (**beware: it will rounded up to a multiple of the page size!**)
- `prot` describes protections on the chunk of memory, a combination of `PROT_READ`, `PROT_WRITE`, and `PROT_EXECUTE`
- `flags` describe properties of the memory (backed by a file, shared with your child on `fork()`, etc) private, anonymous
- `fd` is the file descriptor for the file that we're using to back this memory (-1 for none)
- If we want to map the middle of the file, we can specify a nonzero `offset`

mmap() flags...

- Need to request read / write pages...
 - Why?
 - Because anything that calls malloc() (and thus your allocator) will want to write data to that memory and read the value from that memory.
 - Use PROT_READ and PROT_WRITE...
 - How do you provide both of these flags in a single parameter (prot) of mmap?

PROT_READ PROT_WRITE are actually (something like):

0010

0001


- Combine them using the bitwise-or operation:

PROT_READ | PROT_WRITE // Which actually means “and” in this case...

0011

Mmap

16



Pages managed by you...

77

Extra Credit

Large Blocks array[500]

Medium

dbl1

dbl2

Small

77

int2

int1

- `mmap()` will return the address of one (or more) pages in memory (depending on how much memory you asked for).

```
void * a = mmap( nullptr, size, ...)
```

- Note, size should be equal to $X * \text{Page_Size}$
 - In other words, you should only allocate in multiples of a page. Why?
 - Because this is what `mmap()` is going to return to you anyway.

```
int * i = malloc( sizeof( int ) )
```

- Becomes (in our assignment)?

```
void* d = mmap( nullptr, 4096... )
return d
```

- `*i = 77`
- Lots of unused space – but we aren't going to worry about that. (See Extra Credit)

Opposite of mmap?

17

- What is the opposite of `ptr = malloc(size)`?
 - `free(ptr)`
- What is the opposite of `ptr = mmap(... size ...)`?
 - `munmap(...)`
 - What parameters does it take?
 - `munmap(ptr)` // ?
 - Almost but not quite...
 - `munmap(ptr, size)` // Yes, both ptr and size!!!
- But `free()` doesn't take size as a parameter?!
How does this work?
 - The malloc library remembers it for us...
- Which brings us to your hash table...

77

Large Blocks		500, ...					
array[1000]							
Medium		double1					
				double8			
Small							
int9							
		int4					

Hash Table and Memory Deallocation

18

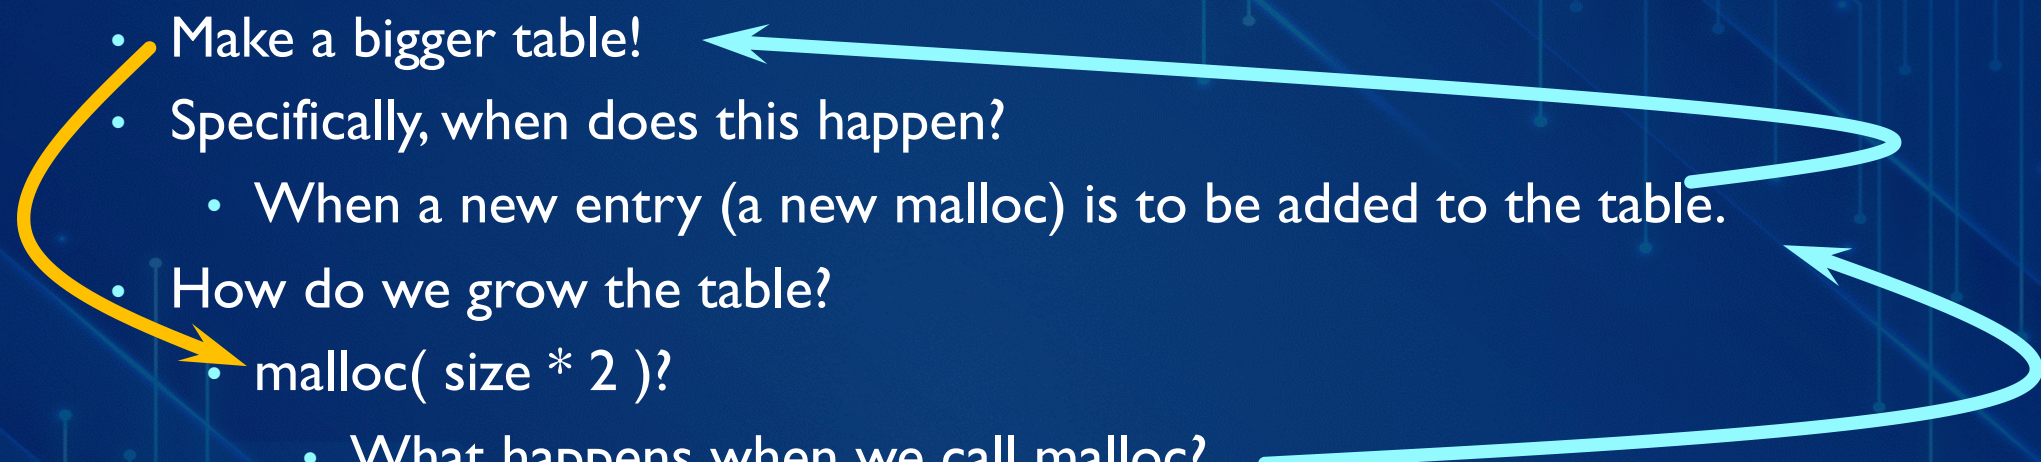
- Our allocator declares its deallocate method as:
`myAllocator.deallocate(void * ptr)`
- To give back the memory, why do we need to know both the size of the data and the pointer (address) to the memory?
 - Because `munmap()` needs both of those pieces of information!
- How do we know the size of data that `myPtr` represents (when we go to deallocate it)?
 - `delete(myPtr);`
 - We have to save it when `myPtr` is originally allocated.
- So we store it in the hash table
 - This is the entire point of having the hash table.
- We are keeping track of all the allocations we do (via the hash table) so that when we deallocate memory, we can look up the size and use the size as a parameter to `munmap`!

The Hash Table's Table

- What data are you storing in your hash table?
 - The size of memory that was allocated (and the key is the address of that memory).
- What does the hash table's table actually look like?
 - Well, we know that it is just an array.
- 1) What is the importance of a hash function? 2) What hash function are you going to use?
 - 1's answer) To minimize the number of collisions
 - 2's answer) What about:
 - `int hash(void * ptr) { return ptr % table_size; }`
 - Good or bad?
 - Probably bad...
 - Try different ways of generating the hash... do some timings and let us know what you find.

The Hash Table's Table


20

- What happens when the table becomes (too) full?
 - Make a bigger table!
 - Specifically, when does this happen?
 - When a new entry (a new malloc) is to be added to the table.
 - How do we grow the table?
 - `malloc(size * 2)`?
 - What happens when we call malloc?
 - **Infinite loop! Crash and burn...**
 - So how do we grow the table?
 - mmap a new table
 - 2x the old size
 - copy the old table into the new table (rehash)
 - munmap the old table.
- 

The Hash Table's Table?

21

- The hash table's table is just...
 - An array
- What does it look like?
 - Entry *[100]?
 - How did you create the new Entry?
 - **new Entry()**?
 - What does new do under the covers?
 - malloc...
 - What happens when malloc is called?
 - **Infinite loop – crash and burn!!!**
- So, what must it look like?



The Hash Table's Table

- `Entry*[100]`
- `Entry[100]`
- What values are in the table to start with?
 - Address?
 - Something that means this entry is not currently in use...
 - `nullptr` seems reasonable
 - Size?
 - Hash table is keying off the address, so if the address is invalid, it really doesn't matter what size is... however, perhaps 0.

Hash Table's Table

[illegible]

Inserting...

- `insert(ptr1, 8)`
 - `ptr1` hashes to location 5.
 - What does the table look like now?

Hash Table's
Table

0	nullptr 0
1	nullptr 0
2	nullptr 0
3	nullptr 0
4	nullptr 0
5	nullptr 0
6	nullptr 0
7	nullptr 0

Inserting...

- `insert(ptr1, 8)`
 - `ptr1` hashes to location 5.
 - What does the table look like now?
- `insert(ptr2, 4)`
 - `ptr2` hashes to location 5
 - What does the table look like now?

Hash Table's
Table

0	nullptr 0
1	nullptr 0
2	nullptr 0
3	nullptr 0
4	nullptr 0
5	<code>ptr1</code> 8
6	nullptr 0
7	nullptr 0

Inserting...

- `insert(ptr1, 8)`
 - `ptr1` hashes to location 5.
 - What does the table look like now?
- `insert(ptr2, 4)`
 - `ptr2` hashes to location 6
 - What does the table look like now?
- `delete(ptr1)`
 - What does the table look like now?

Hash Table's
Table

0	nullptr 0
1	nullptr 0
2	nullptr 0
3	nullptr 0
4	nullptr 0
5	ptr1 8
6	ptr2 4
7	nullptr 0

Inserting...

- `insert(ptr1, 8)`
 - `ptr1` hashes to location 5.
 - What does the table look like now?
- `insert(ptr2, 4)`
 - `ptr2` hashes to location 5
 - What does the table look like now?
- `delete(ptr1)`
 - What does the table look like now?
- Now look up `ptr2`...
 - What happens?
 - Hash to 5... What is at 5?
 - This is a problem!!! How to fix?
 - Need to know that Entry #5 was previously used...

Hash Table's
Table

0	nullptr 0
1	nullptr 0
2	nullptr 0
3	nullptr 0
4	nullptr 0
5	nullptr 8
6	ptr2 4
7	nullptr 0

nullptr
-1? 0?

Inserting...

- Now look up ptr2...
 - What happens?
 - Hash to 5... What is at 5?
 - This is a problem!!! How to fix?
 - Need to know that Entry #5 was previously used...
 - Mark it some how...
 - In this example, I used a -1 “address”
 - Alternatively could add a “deleted” flag to the Entry
 - Though this “wastes” a lot of space.

Hash Table's
Table

0	nullptr 0
1	nullptr 0
2	nullptr 0
3	nullptr 0
4	nullptr 0
5	-1 0
6	ptr2 4
7	nullptr 0

Finding...

- On this slide I've written "ptr101". What is actually in the table there?
 - The address that ptr101 is pointing to
 - 0x29ef47a0
- Is ptr7 in the table? It hashes to 5.
 - How do you know (determine) if it is in the table?
 - Must "wrap" around to continue looking
 - If you hit a nullptr, it is not
 - What is the other case in which ptr7 is not in the table, but you don't find a nullptr?
 - Hint: There are no nullptrs in the table...
 - Answer: If the table were to be completely full.

Hash Table's
Table

0	ptr101 8
1	ptr999 1000
2	nullptr 0
3	nullptr 0
4	ptr19 50
5	ptr2 4
6	ptr222 4
7	ptr123 120

Malloc Assignment

- Do not use:
 - **System** `malloc()`, `delete()` :-p
 - Except for testing / comparison to your version
 - **c++** vectors, etc
 - `new()` / `free()`

Replacing malloc / free

30

- `malloc()` is just a function.
- We can **overload** (well, **override**) it just by creating our own:

```
void * malloc( size_t size ) {  
    // What do we do in here?  
    void * mem = myAllocator.allocate( size );  
    return mem;  
}
```

- And what about `free()`?

```
void free( void * ptr ) {  
    // What do we do in here?  
    myAllocator.deallocate( ptr );  
}
```

- What is `myAllocator`?
 - An object of type of the class you are writing.
- Where does it come from / how to I create it?
- Use a static variable:
 - *// Somewhere near the top of main program...*
`static MyAllocator myAllocator;`
- What is a static variable?
 - One that always exists and is created (constructed) before `main()` begins.
- Note: when our `malloc()` and `free()` are linked in, they will replace the default ones provided by the `malloc` library / compiler.
 - You may want to put a debug print statement in your `malloc()` to verify it is being used. Once you've done this, you can remove the print statement.
- Note: Some of (debug) functions may not work correctly if they use `malloc()` internally.

Timing

31

```
#include <ctime>

// clock() returns the ~amount of CPU time a process
// has used.

clock_t start_time = clock();
// Do some work here..

clock_t end_time = clock();
float time_used = end_time - start_time;
float cpu_time_in_seconds = time_used / CLOCKS_PER_SEC
```


~ Fin ~