

Systems I – CS 60I3

Computer Architecture and Operating Systems

Lecture 29: Locked Data Structures

MASTER OF SOFTWARE DEVELOPMENT (MSD) PROGRAM

J. DAVISON DE ST. GERMAIN

SPRING 2024

(Adapted from Slides by Prof. Andrea Arpaci-Dusseau & Nima Honarmand,
and previous MSD presentations)

Lecture 29 – Topics

2

- Locked Data Structures

Announcements

- Lisp – lots of insidious parentheses

Review

- Concurrency leads to non-deterministic results
 - Different results even with same inputs
 - We call this a **race condition**
 - **Data races** are race conditions involving *write contention*.
- We saw mutexes (locks) for protecting critical sections.
- We saw condition variables for more complicated thread operations.
 - One thread signaling another thread(s)
 - Condition variables + locks = the **monitor pattern**.
 - Monitoring a global condition.

Today

- Some practical applications of locks / mutexes
- We'll design (to work with Threads):
 - A simple counter
 - A concurrent (thread-safe) linked list
 - A concurrent queue
 - A concurrent hash table

Case Study I: A Threaded Counter

```
int value;
void init() {
    value = 0;
}
void increment() {
    value++;
    // How many assembly
    // instructions?
    //      3
}
void decrement() {
    value--;
}
```

- Potential Race Conditions?
 - Everything except init has a race condition.
 - How would you fix these?
 - A simple approach: lock every operation that could have a race condition.

Case Study I: A Concurrent Counter

```
int value;
std::mutex m;
void init(){
    value = 0;
}
void increment(){
    m.lock();
    value++;
    m.unlock();
}
void decrement(){
    m.lock();
    value--;
    m.unlock();
}
```

- Evaluation
 - This is indeed a safe counter.
 - However, is it actually faster to use multiple threads in this example?
 - Thread creation and locks have overhead.
 - This actually slows down as we throw more threads at it!
- This does illustrate the pitfalls of naive locking.
- The solution to this is nontrivial, and we won't go over it.
 - See textbook, Chapter 29, Approximate counter.

Case Study 2: A linked list [in C]

```
typedef struct __node_t {
    int data;
    struct __node_t * next;
} node_t;

typedef struct __list_it {
    node_t * head;
} list_t;

void init( list_t * L ){
    L->head = NULL;
}
```

```
int insert( list_t * L, int data ){
    node_t * new_node =
        malloc( sizeof( node_t ) );
    if( new_node == NULL ){
        return -1;
    }
    new_node->data = data;
    new_node->next = L->head;
    L->head = new_node;
    return 0;
}
```

- Difference from C++?
 - Constructor, functional vs OOP, etc. (see next slides)

C Syntax - Explanation

```
struct node_t {  
    int      data;  
    struct node_t * next;  
};
```

- ‘_t’ <- naming convention to say it is a type
- “struct” must be repeated every time you use node_t, so to create the head of the list:
- struct node_t * head;

```
typedef struct __node_t {  
    int data;  
    struct __node_t *next;  
} node_t;
```

- ‘__’ <- naming convention to say “you should never use/see this”
- typedef defines a type named node_t;
- node_t * head;

Case Study 2: A linked list [in C]

```
typedef struct __node_t {  
    int data;  
    struct __node_t *next;  
} node_t;
```

- In C++, you would say:
 - `list.insert(data)`
- In C, you have to pass both the list, and the data into the insert function.

```
int insert( list_t * L, int data ) {  
    node_t * new_node =  
        malloc( sizeof( node_t ) );  
    if( new_node == NULL ) {  
        return -1;  
    }  
    new_node->data = data;  
    new_node->next = L->head;  
    L->head = new_node;  
    return 0;  
}
```


Case Study 2: A linked list [in C]


11

```
struct node_t { int      data;  
                node_t * next; }
```

```
struct list_t { node_t * head; }
```

```
void insert( list_t * L, int data ){  
    node_t * new_node =  
        malloc( sizeof( node_t ) );  
    new_node->data = data;  
    new_node->next = L->head;  
    L->head = new_node;  
}
```

```
...  
insert( list, 99 );
```



Draw the picture!!!

```
void init( list_t * L ){  
    L->head = NULL;  
}
```


Case Study 2: A linked list [in C]

12

```
void insert( list_t * L, int data ){
    node_t * new_node =
        malloc( sizeof( node_t ) );
    new_node->data = data;
    new_node->next = L->head;
    L->head = new_node;
}
```



- Ask: What happens if multiple threads run this code at the same time?
 - Identify the race conditions.
 - Where should the mutex object live?
 - Where should I lock and unlock the mutex?
 - 1st attempt: don't want any of it interrupted – lock everything

Case Study 2: Concurrent List v2

13

```
void insert( list_t * L, int data ) {
```

```
    lock( m );
```

```
    node_t new_node =
```

```
        malloc( sizeof( node_t ) );
```

```
    if( new_node == NULL ) {
```

```
        return;
```

```
    } m.unlock();
```


```
    new_node->data = data;
```

```
    new_node->next = L->head;
```

```
    L->head = new_node;
```

```
    unlock( m );
```

- ```
}
```
- What is `m`? Where is it defined?
    - Mutex. Somewhere in the global scope so all threads can see it.



- What's missing?
  - How do we fix (avoid) needing this 2<sup>nd</sup> unlock?
  - Use a scoped lock!
  - But is there a better way?



# Case Study 2: Concurrent List v3

14

```
void insert(list_t * L, int data){
 lock(m); // scoped lock
 node_t new_node =
 malloc(sizeof(node_t));
 if(new_node == NULL) {
 return; // throw error?
 }
 new_node->data = data;
 new_node->next = L->head;
 L->head = new_node;
}
```

## • Evaluation

- This is thread-safe.
- However, it is inefficient. There is a giant lock around every bit of code.
- This is a coarse-grained lock.
- A little reading reveals that `malloc()` is thread safe.
- Then we don't need to lock `malloc()`.
- This allows a finer-grained lock.
- What about the update of `new_node`? Does it need to be locked?
  - Is there a difference between the update of `data` and `next`?
  - `new_node` is a local (non-shared) variable.
  - However, `next` is referencing `L`, which is a global variable, so must be locked.
  - (See next slide for updated version of code)



# Case Study 2: Concurrent List v3

15

```
void insert(list_t * L, int data) {
```

```
 node_t new_node =
```

```
 malloc(sizeof(node_t));
```

```
 if(new_node == NULL) {
```

```
 return; // throw error?
```

```
 }
```

```
 new_node->data = data;
```

```
 lock(m);
```

```
 new_node->next = L->head;
```

```
 L->head = new_node;
```

```
 unlock(m);
```

```
}
```

- Performance ramifications?

- Better because less code is locked

- Note, could actually use a scoped lock in this code too (in which case we would not need the unlock at the end)



# Expanding List Functionality

- Let's see another linked-list example for locking.
- Next, we will write a function (`lookup`) to search our linked list.
  - We'll return `false` if the item is not found, `true` if it is found.
- We can write a concurrent version right away if we use a coarse-grained lock.



# Case Study 3: Concurrent List Lookup

17

```
bool lookup(list_t * L, int data){
```

```
 m.lock();
```

```
 node_t * curr = L->head;
```

```
 while(curr){
```

```
 if(curr->data == data) {
```

```
 m.unlock();
```

```
 return true; // success
```

```
 }
```

```
 curr = curr -> next;
```

```
 }
```

```
 m.unlock();
```

```
 return false; // failure
```

```
}
```

- Why do we need to lock the list during a lookup?
  - List could be changed out from under us.
- Why the 2 unlocks?
  - 2 paths out of the function – must remember to unlock in both paths.
- Why is this a really bad way to lock this function?
  - Here, we're locking the whole list before we go through it - a coarse-grained locking strategy.
- What would a fine-grained strategy look like?
  - Think about what actually needs a lock...



# Hand-over-Hand Locking

- **Lock** each node as we traverse the list.
- When done with a node, unlock it and lock the next one in turn.
- If locks can be implemented with low overhead, this is great.
- In practice, this is too fine-grained and will be expensive.
- Lock once every few times? Possible.







# Case Study 4: Michael and Scott

## Two-Lock Concurrent Queue

- This is the algorithm that you must use for your assignment.
- Represent the queue as a singly-linked list with head *and* tail pointers (two ends to the list).
  - \*Always want enqueue and dequeue to be  $O(1)$  operations.
- Use a single lock for enqueue, and a separate lock for dequeue.
- Head always points to a dummy node. Dummy node is the first node in the list.
- Nodes are inserted after the last node of the list.
- Nodes are deleted from the beginning of the list.





# Case Study 4: Michael and Scott

## Two-Lock Concurrent Queue

- Initially, both head and tail point to a dummy node.
- If we didn't have the dummy variable, both head and tail could point to the same node.
  - This would cause problems if enqueue and dequeue occurred simultaneously.
  - Two locks + two threads = potential for “**deadlock**” or “**livelock**”.
    - More later.
- As we enqueue stuff, we update tail pointer.
- As we dequeue stuff, we update the head pointer.



# Case 4: Michael/Scott Concurrent Queue

21

```
struct node_t { int data;
 node_t * next; }
```

```
struct queue_t { node_t * head;
 node_t * tail;
 mutex head_m,
 tail_m; }
```

```
void init(queue_t * q) {
 node_t * tmp = malloc(...);
 tmp->next = NULL;
 q->head = tmp;
 q->tail = tmp;
}
```

```
void enqueue(queue_t * q, int value) {
 node_t * tmp = malloc(...);
 tmp->data = value;
 tmp->next = NULL;

 tail_m.lock();
 q->tail->next = tmp;
 q->tail = tmp;
 tail_m.unlock();
}
```

- Why have a separate tail and head lock?
  - Allows enqueue and dequeue to happen simultaneously!
  - Almost always, the head and tail are not next to each other.



# Case 4: Michael/Scott Concurrent Queue

22

```
struct node_t { int data;
 node_t * next; }
```

```
struct queue_t { node_t * head;
 node_t * tail;
 mutex head_m,
 tail_m; }
```

```
void init(queue_t * q) {
 node_t * tmp = malloc(...);
 tmp->next = NULL;
 q->head = tmp;
 q->tail = tmp;
}
```

```
bool dequeue(queue_t * q, int * value) {
 head_m.lock();
 node_t * tmp = q->head; // dummy node
 node_t * new_head = tmp->next;
 if(new_head == NULL) {
 head_m.unlock();
 return false; // Nothing in queue
 }
 *value = new_head->data;
 q->head = new_head;
 free(tmp);
 head_m.unlock();
 return true;
}
```



# Evaluation

- Enqueue is a nice example of fine-grained locking.
- The dummy node is a great example of how serial data structures must be modified for concurrency
  - ...beyond just adding locks.
- Michael and Scott show that two-lock queue scales very well and avoids deadlock / livelock.
- If you use scoped locks, the lock will automatically unlock when it goes out of scope.
  - Use curly braces { } to define scopes.
  - Allows for unlock even after returning from middle of function!



# Case Study 5: Concurrent Hash Table

- Final example: a hash table without resizing.
- We'll use separate chaining instead of probing.
  - Recall that this means an array of linked lists.
  - Each “bucket” in the table is now a linked list.
  - Collisions are resolved by hashing to the same bucket, then growing the corresponding linked list.
- We can use our concurrent list from earlier!
  - `list_t table[ NUM_BUCKETS ]`.
- Hash table insert is just a call to the `list_t` insert function.
- Lookup is a call to the `list_t` lookup.



# Evaluation

- The concurrent hash table scales well.
- Far better than the linked list does.
- Constant time roughly maintained as we increase the number of threads.



# Summary

- Coarse-grained locking (locking large chunks of code) is easy, but inefficient.
- Fine-grained locking
  - trickier to get right without deadlock/livelock
  - may need to redesign data structure or code
  - may still scale poorly
- Enabling concurrency doesn't always improve performance!
  - Sometimes all we can do is hope it doesn't degrade it.
- Next week: concurrency bugs, lock-free data structures



~ Fin ~