

Systems I – CS 6013

Computer Architecture and Operating Systems

Lecture 14: The Unix Shell

MASTER OF SOFTWARE DEVELOPMENT (MSD) PROGRAM

J. DAVISON DE ST. GERMAIN

SPRING 2024

*(adapted from slides by Scott Brandt at UC Santa Cruz and other general sources, including previous MSD slides)

CS 6013 – Spring 2024

Lecture 14 – Topics

- The Unix Shell Assignment
 - How it fits together
 - Processes, Forking, Pipes, Exec

Miscellaneous

- Vocabulary from Industry Seminar, etc
 - Front end / back end / full stack
 - Excel / VBA
 - **Know** the things on your resume!

Unix Shell Assignment

4

Shell

- Loop
 - Show ">", read user input
 - Type of user input?
 - String
 - What is the user input?
 - Command Line
 - What do we do with the CL?
 - Parse it... into what?
 - Tokens
 - Next step?
 - Create commands
 - ...

sshd

emacs

- What processes are out there?
- What is "Shell" doing?

/sbin/init

- PID?
 - 1

bash

ntpd

Network time
Protocol
Daemon

ps -def

Version I – Run a Single Command

5

Parent Process

Shell

- Ok, now we have a list of commands... For this example, we'll assume only one command:
- `ls -l`
- What does your shell do now?
 - Create a new process to run the "`ls -l`" command... How does a program create another new process?
 - `Fork()`
 - Ok, what does our processes bubble diagram look like now?
 - What does this Shell do now?
 - Waits for child to finish.

New (Child) Process

shell

- Wait! Why is this process named "Shell"? Isn't this "`ls -l`"? (Name not bolded so I can distinguish them.)
- What does `fork()` do?
- Creates an exact copy of the parent process...
- Thus here we are.
- What happens now?
 - Let's assume the CPU swaps us out...
 - And we're back... doing?
 - Replace this "shell" with?
 - `ls -l`

Version 1 – Run a Single Command

6

Parent Process

Shell

- Waiting on child...

New (Child) Process

shell

- Replace this “shell” with?
 - ls -l
- How?
 - exec()
 - And now what does our process bubble diagram look like?

Version 1 – Run a Single Command

7

Shell

- Waiting on child...

ls

- “Ah! I feel like a brand-new process!”
- Where’d the “-l” go?
 - Not part of the process name.
 - Stored in?
 - `main(argv[])`
 - Next, what does `ls` do?
 - **ls** does its thing (lists files in the current directory)
- And then?
 - returns / exits
- And our bubble diagram?

Version 1 – Run a Single Command

8

Shell

- Waiting on child...
 - Child is gone (finished / exited)
- And we're back...
- Shell continues with?
 - Read command line.
 - `w | grep dav > out.txt`

Version 2 – Single Command w/ output

9


Shell

- `ls -l > out.txt`
- What is ">"
 - Redirection of output to a file.
 - What does the `struct Command` look like after parsing the above command line?
 - `execName?`
 - `ls`
 - `argv?`
 - `[ls, -l]` (Notice, no "> out.txt")
 - `background?`
 - `False`
 - `inputFd?`
 - `0` (standard in)
 - `outputFd?`
 - `5` (why 5?)

Kernel Memory

File descriptor table (for Shell)

0	1	2	3	4	5	6	7	8	...
---	---	---	---	---	---	---	---	---	-----



- Note: Kernel memory diagram overlaps Shell process on purpose, why...
 - Because Shell has some of its data in the kernel!
- The FDT is indexed by a bunch of integers, but what does it actually contain*?
 - Each bucket contains a “pointer” to a file!!!
 - Which includes information (for the OS) on how to send/receive data to/from that thing (monitor, keyboard, actual file, pipe, etc).
- Back to our question, why 5?

Version 2 – Single Command w/ output

10


Shell

- `ls -l > out.txt`
- So, what code does the shell run to create this file?
 - `int fd = open("out.txt", ...)`
 - And then does what with this FD?
 - Remember the shell is currently in the process of parsing the command line...
 - What functions do this?
 - `tokenize()`
 - Breaks it into tokens but doesn't know anything about them...
 - `getCommands()`
 - Yes, the open happens here...

Kernel Memory

File descriptor table (for Shell)

0	1	2	3	4	5	6	7	8	...
---	---	---	---	---	---	---	---	---	-----



- Taking a step back... who creates (opens) `out.txt`?
 - The Shell. Why?
 - No easy / generic way to pass “`> out.txt`” to the child process.
 - Shell is responsible for setting up all the “initial” file descriptors for its children to use.

Version 2 – Single Command w/ output

11


Shell

- `ls -l > out.txt`
- `getCommands()`
 - `int fd = open("out.txt"...`
 - What does `getCommands` return?
 - A vector of commands!
 - In this case?
 - A single `struct Command`.
 - At this point, that struct looks like?
 - `execName: ls`
 - `argv: [ls, -l]`
 - `background: False`
 - `inputFd: 0` (standard in)
 - `outputFd?`
 - 5 (why 5?)
 - fd above has a value of 5!

Kernel Memory

File descriptor table...

0	1	2	3	4	5	6	7	8	...
---	---	---	---	---	---	---	---	---	-----



Version 2 – Single Command w/ output

12

Shell

- `ls -l > out.txt`
- Done with `getCommands()` ...
 - And thus we have?
 - One `struct Command` for **ls**
 - Now what?
 - What is the purpose of the Shell right now? What is it (about to) doing for us?
 - Create a new process (for `ls`)
 - How?
 - `fork()`

shell*

- Back to two Shells again!!! What?
- What to do to become an “ls”?
 - `exec()`!
- But not yet... why?
 - `Exec` (wipes) replaces this processes memory...
- What (very important) information do we share with our parent?
 - Info that tells us to become “ls”
 - Remember, we are an exact* copy of Shell.
 - `struct Command!!!`
- We'll get back to this in a minute, but first let's make sure we know...

Version 2 – Single Command w/ output

13

Shell


- `ls -l > out.txt`
- Done with `getCommands()` ...
 - And thus we have?
 - One `struct Command` for `ls`
 - Now what?
 - What is the purpose of the Shell right now? What is it (about to) doing for us?
 - Create a new process (for `ls`)
 - How?
 - `fork()`

shell*

- Where is this process getting input from and writing its output to?
 - Well, just like any new process, `stdin` and `stdout`?
 - Above is not exactly true... this process actually inherits its parent's input and output file descriptors!
 - What is this process's parent?
 - Shell
 - What are Shell's input and output?
 - `stdin`, `stdout`
 - So, is `stdin` and `stdout` what we want?
 - No... What do we want?
 - `stdout` reassigned to 5

Version 2 – Single Command w/ output

14




- Remember Kernel Memory and the FDT?
- **Shell** (seemingly a long time ago) opened “out.txt” and it was associated with FD 5. (before `fork()`)
- That still exists in Kernel memory and we have a copy of it.
- So `dup2` updates this picture to look like...
- Note, there is a FDT for each process. The FDT we see here belongs to shell* (the copy), but began as an exact copy of which (whose) FDT?
 - **Shell** (original / main shell process)

shell*

- How do we change our `stdout` to be file descriptor 5?
- How do we “rename” a file descriptor?
- How do we duplicate a file descriptor and replace it with a different FD?
 - `dup2()`
- So now what are we (finally) ready to do?
 - Replace Shell with ls... how?
 - `exec()`
- But `exec()` will blow away all of our “memory” including our file descriptors...
 - Or does it? Where is the/our In/Out FD info stored?
 - In the Kernel FDT

Version 2 – Single Command w/ output

15



ls

- `exec()` is done, and this process is now ls
- While this process is busy doing its job...

ie: making system calls to ask the OS about the files in the current directory and then making more system calls to send that information to its "stdout",

```
cout << fileInfoString << "\n";
```

...let's jump back to Shell...

Back to Parent (Original/Real Shell)

16

Shell

- The main **Shell**'s FDT still looks like this...
- What does it need to do?
- Does **Shell** do anything with out.txt?
 - So what should it do?
 - `close(fd)`
// Remember fd == 5

Kernel Memory

File descriptor table (**Shell**)

0	1	2	3	4	5	6	7	8	...
---	---	---	---	---	---	---	---	---	-----



ls

Remember,
I'm still out
here. 😊

Version 3 – Multiple | Commands

17

Shell

- `ls -l | nl | tail`
- Step one?
 - `tokenize()`
- Step two?
 - `getCommands()`
- What happened when `getCommands()` parsed the first `|` (pipe) in the token list?
 - What does the `|` mean sitting there between **`ls -l`** and **`nl`**?
 - Send the output from **`ls -l`** to the input of **`nl`**!
- What will **`ls`** and **`nl`** become (shortly)?
 - Processes
 - How to send data between processes?
 - Pipes!

Kernel Memory

File descriptor table (Shell)

0	1	2	3	4	5	6	7	8	...
---	---	---	---	---	---	---	---	---	-----

StdOut

StdIn

ls

Remember,
I'm still out
here. 😊

Version 3 – Multiple | Commands

18


Shell

- `ls -l | nl | tail`
- `getCommands()`
 - See a `|`, create a pipe:
 - `pipe(fds)`
- What are the values in `fds`?
 - Shrug, but let's assume 7 and 8.
 - `fds[0] == 7 // read end of pipe`
 - `fds[1] == 8 // write end of pipe`
- `pipe()` created a pipe object (data structure)... where is it?
- In the Kernel memory!
- Where are the 7 and 8 stored (not the pipe itself, just the FD numbers) with respect to Shell?
 - What function are we in?
 - In `getCommands()`

Kernel Memory

File descriptor table (Shell)

0	1	2	3	4	5	6	7	8	...
---	---	---	---	---	---	---	---	---	-----



Version 3 – Multiple | Commands

19

Shell

- ls -l | nl | tail
- fork -> **ls -l**
- fork -> **nl**
- fork -> **tail** (pipe (b) for **nl** to **tail** was also created).
- **Bubbles now?**

Shell-a

- What does my FTD look like?
 - Copy of Shell


Shell-b

Shell-c

Kernel Memory

File descriptor table (Shell)

0	1	2	3	4	5	6	7	8	...
---	---	---	---	---	---	---	---	---	-----



File descriptor table (Shell-a)


Version 3 – Multiple | Commands

20

Shell-b

- Soon I'll be **nl**
- But before I become **nl**, what do I need to do?
- What is my output supposed to be?
 - Sent to **tail**
- What is my input supposed to be?
 - Received from **ls -l**
- How do we do this?
 - How did we handle **ls > out.txt**?
- dup2, and dup2 again...

Kernel Memory




Version 3 – Multiple | Commands

21

Shell-b

- Soon I'll be **nl**
- But before I become **nl**, what do I need to do?
- What is my output supposed to be?
 - Sent to **tail**
- What is my input supposed to be?
 - Received from **ls -l**
- How do we do this?
 - How did we handle **ls > out.txt**?
- dup2, and dup2 again...
 - Where are the FDs I'm replacing stored?
 - struct command
 - Specifically?
 - .inputFd
 - .outputFd

Kernel Memory



ls -l | nl | tail

Assignments

- Code Review?
 - Anyone want me to review their lab code?
- Unix Shell
 - More Questions?
- Named Pipes (FIFOs) – Coding / Using

~ Fin ~