

# CS 6016



## Database Systems

*Entity Relationship Model*

# Database Design Steps



## This Week

- Database design: **Entity Relationship (ER)** Model
- ER to Schemas

# Database Design Steps

---

1. Requirements Analysis
  - What does user need? What must it do?

# Database Design Steps

1. Requirements Analysis
  - What does user need? What must it do?
2. Conceptual Design
  - High level formal description

# Database Design Steps

1. Requirements Analysis
  - What does user need? What must it do?
2. Conceptual Design
  - High level formal description
3. Schema Refinement
  - Consistency and “normalization”

# Database Design Steps

1. Requirements Analysis
  - What does user need? What must it do?
2. Conceptual Design
  - High level formal description
3. Schema Refinement
  - Consistency and “normalization”
4. Physical Design
  - Indexes, disk layout

# Database Design Steps

1. Requirements Analysis
  - What does user need? What must it do?
2. Conceptual Design
  - High level formal description
3. Schema Refinement
  - Consistency and “normalization”
4. Physical Design
  - Indexes, disk layout
5. Security Design
  - Who accesses it, and how?

# Database Design Steps

## 1. Requirements Analysis

- What does user need? What must it do?

## 2. Conceptual Design

- High level formal description

- Accomplished using “Entity-Relationship” (ER) model



# Problem → Solution

- Problem specification → C++ → assembly

# Problem → Solution

- Problem specification → C++ → assembly
- Problem specification → ER Model → Schema

# Problem → Solution

- Problem specification → C++ → assembly
- Problem specification → ER Model → Schema
- In both of these cases, you *could* skip the middle step

# Problem → Solution

- Problem specification → C++ → assembly
- Problem specification → ER Model → Schema
- In both of these cases, you *could* skip the middle step
- In both of these cases, that would be a bad idea

# Problem → Solution

- Problem specification → C++ → assembly
- Problem specification → ER Model → Schema
- In both of these cases, you *could* skip the middle step
- In both of these cases, that would be a bad idea
- In both of these cases, there is a mechanical (algorithmic) translation to the final result

# ER Model

- What are the **entities**, and their **relationships**?

# ER Model

- What are the **entities**, and their **relationships**?
- For remainder of today, **forget about tables!**



**[shouting] And I don't want any questions  
about the tables!**




# Entity

- A real-world object, distinguishable from other entities
  - {u0123456, "Daniel"}
- An entity is described by a set of **attributes**
  - (uID string, name string)


# Entity Set

- A collection of entities of the same type, e.g.
  - All students
  - All buildings
  - All people
- All entities in the set have the same types of attributes
- An entity set has a key

# ER Diagram



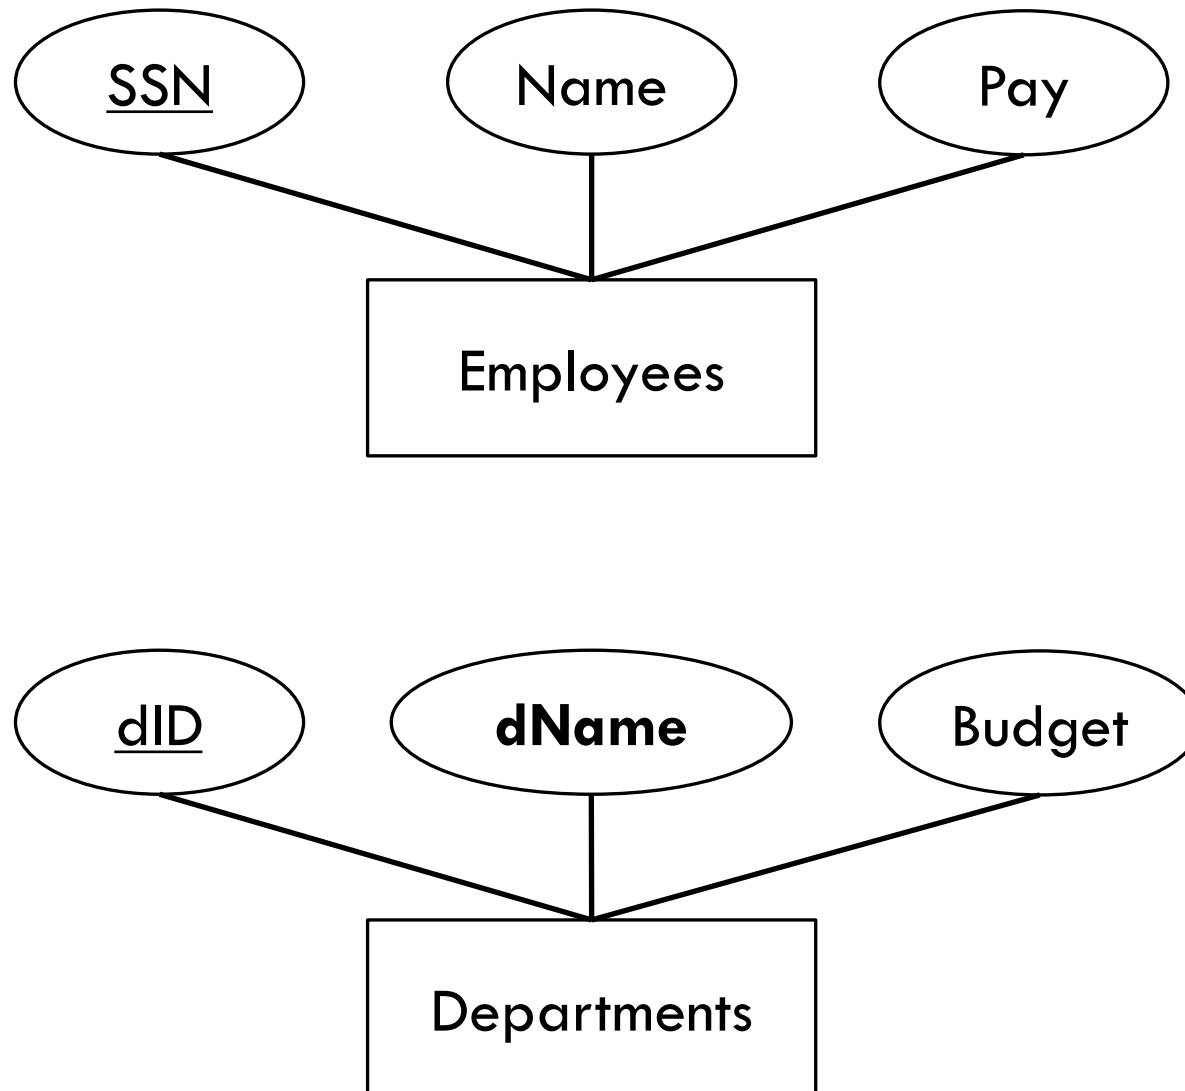
= entity set



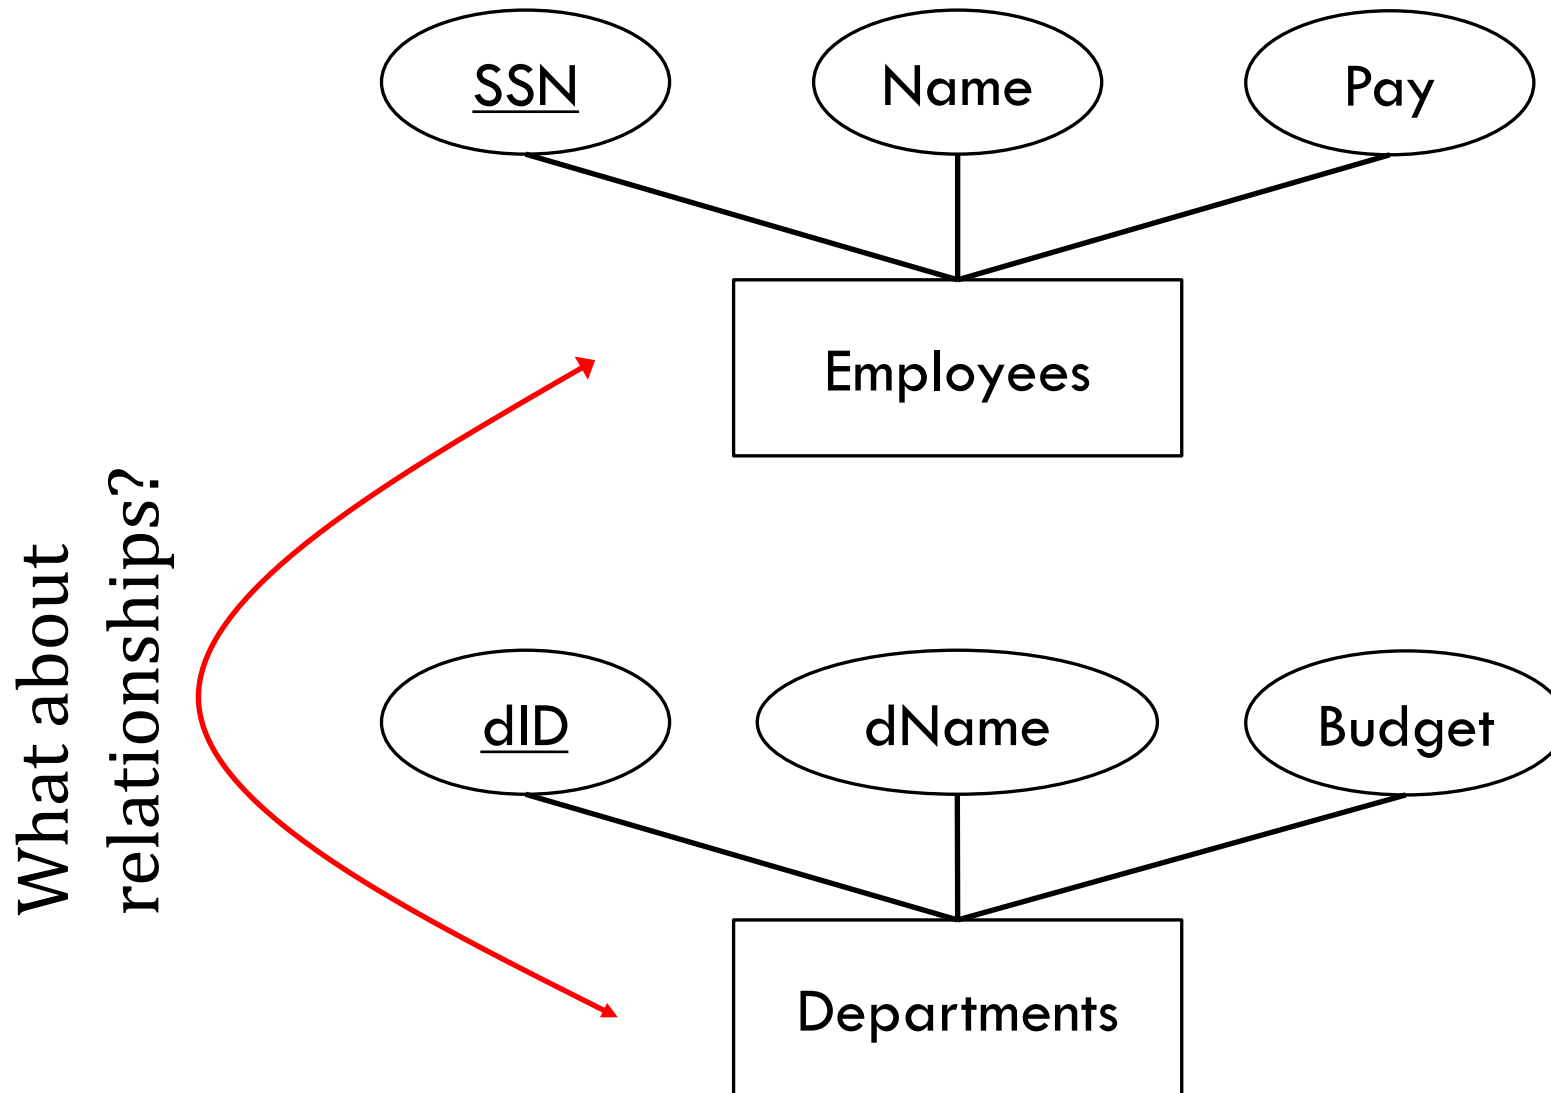
= attribute

underline = key

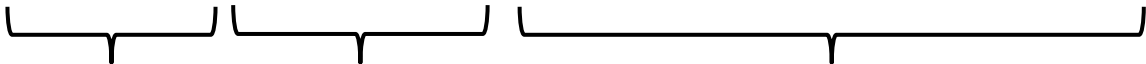
# ER Diagram



# ER Diagram




# Relationship

- Relationship between 2 or more entities:
  - “Daniel works in School of Computing”  
  
entity      relationship      entity

# Relationship

- Relationship between 2 or more entities:

- “Daniel works in School of Computing”



entity

relationship

entity

- Relationship Set:

- Set of relationships between entities of same type
  - e.g. **works in** relates **Employees** to **Departments**

# Relationship Attributes


- Relationships can have attributes as well:

- Daniel works in SoC since 2010  
attribute





# Relationship Attributes

- Relationships can have attributes as well:

- Danny works in SoC since 2010  
A diagram illustrating the concept of a relationship attribute. A horizontal line with an upward-pointing arrow at its right end is positioned below the words 'works in'. The label 'attribute' is placed to the right of this line. A bracket above the line connects the words 'works in' to the label 'attribute'.

- Starting date does not belong to Danny or SoC
  - It belongs to the relationship

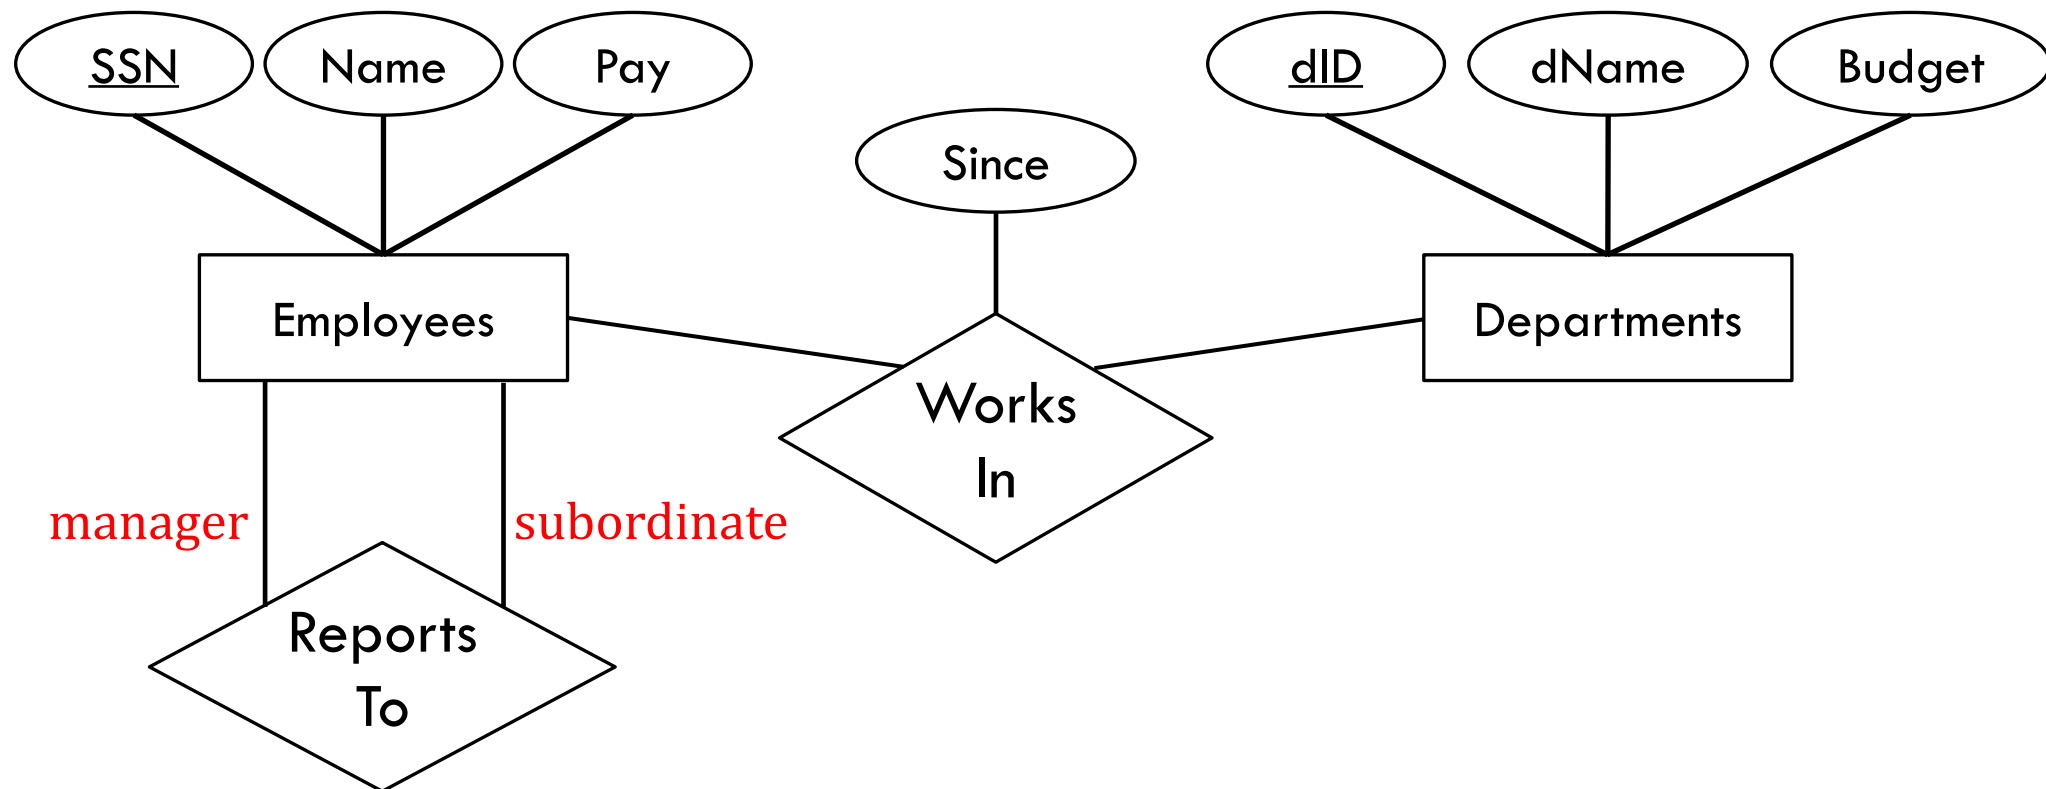
# ER Diagram



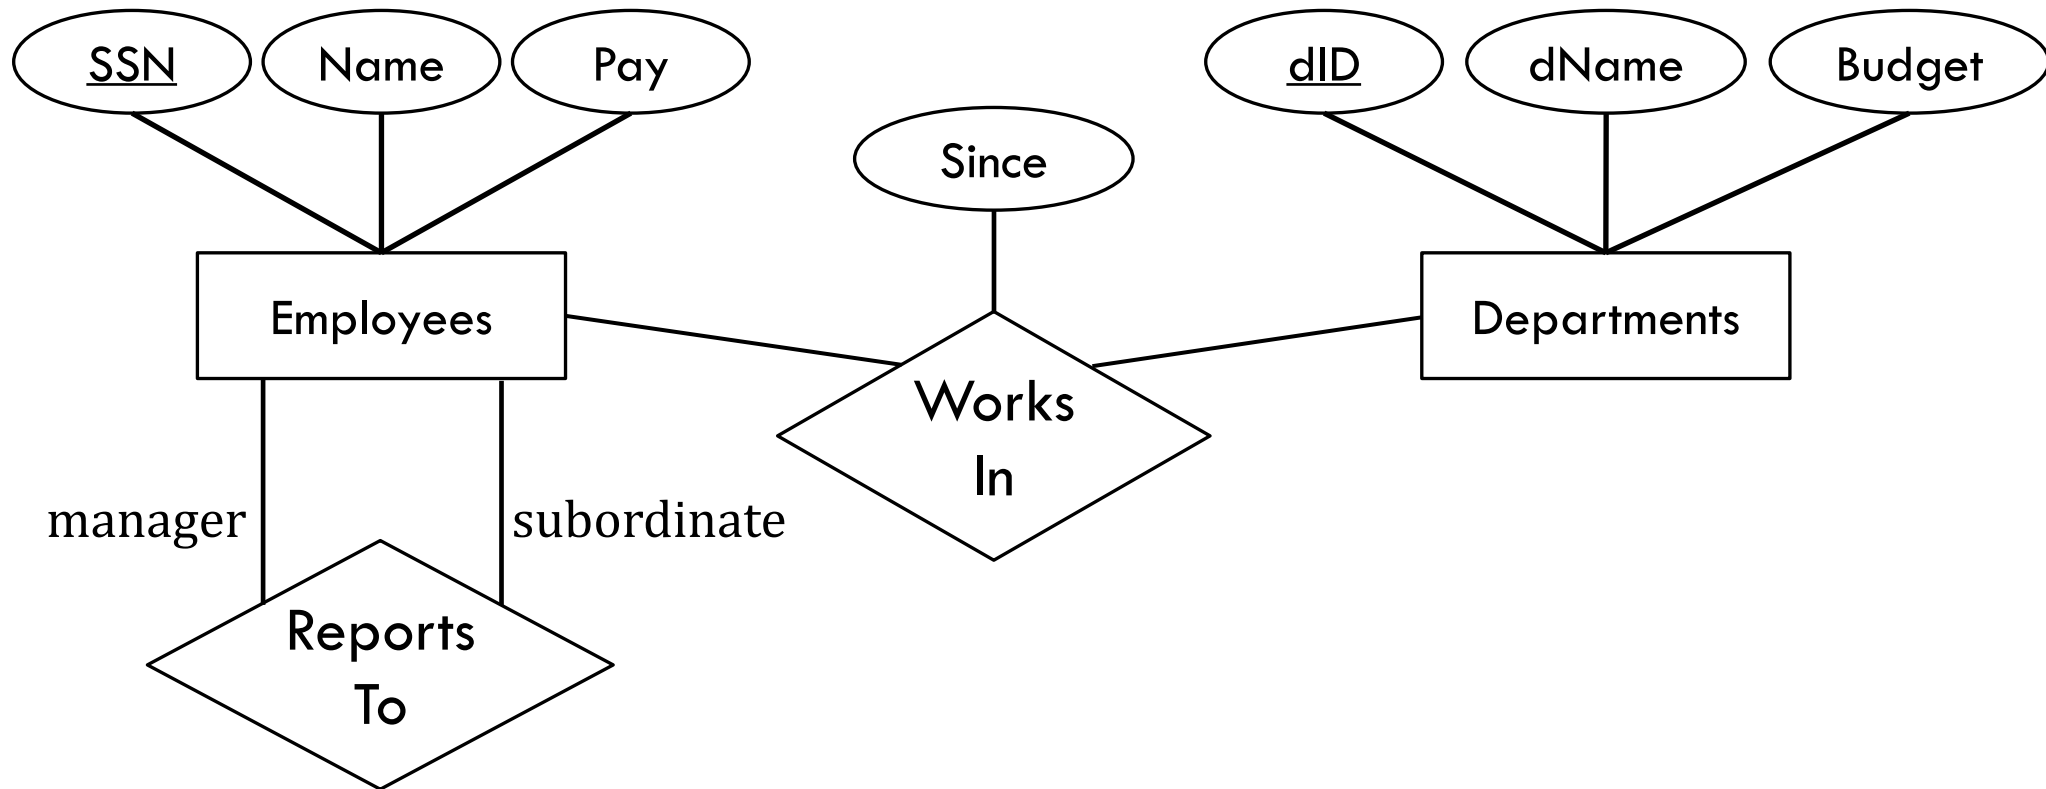
= relationship set

relationship sets do not need a key attribute

# Inter-Relationship

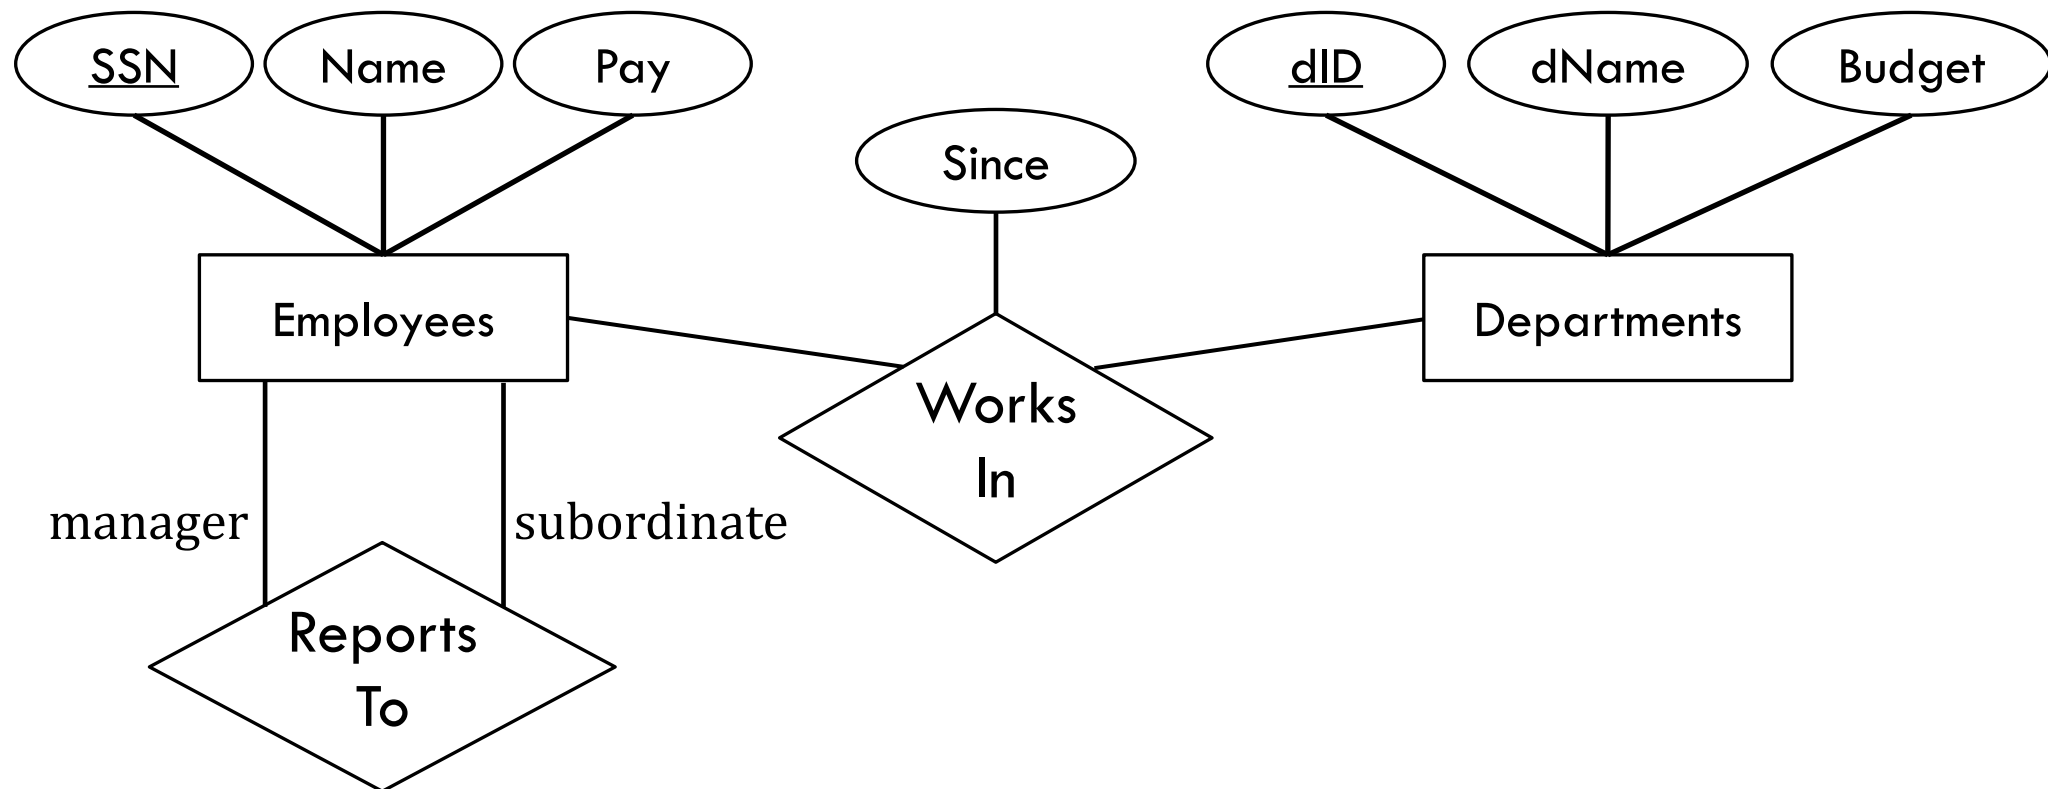


# Cardinality




- Can an employee work in multiple departments?
- Can a department have multiple employees?

# Cardinality




- The diagram does not specify this yet...

# Cardinality




One to one

# Cardinality



Many to many

# Cardinality




One to many




# Cardinality

Three major classifications to restrict relationship sets




Many-to-Many



One-to-Many

VS


~~Many-to-One?~~



One-to-One

# Cardinality


- Annotate **opposite edge** of relationship with cardinality



- “An employee can work for one department, but a department can have many employees”
  - 1-to-Many

# Cardinality


- Annotate opposite edge of relationship with cardinality



- “An employee can work for multiple departments, but a department can only have one employee”
  - 1-to-Many (reversed)

# Cardinality


- Annotate opposite edge of relationship with cardinality



- “A student can take multiple classes, and a class can have multiple students”
  - **Many-to-Many**

# Participation Constraint


- Bold line indicates all entities in the set *must* participate
- Can also use a double line (two parallel lines).



- “An employee must work in 1 department, but a department does not necessarily have any employees”

# Participation Constraint


- Bold line indicates all entities in the set *must* participate
  - At least once



- “An employee must work in one department, a department must have at least one employee”


# Practice (Annotate Diagram)

- An **A** has *at most* one **B**
- A **B** has *at least* one **A**



# Practice (Annotate Diagram)


- An **A** has *at most one* **B**
- A **B** has *at least one* **A**






# Practice

- A department can have multiple employees
- An employee works for exactly one department
- A department has exactly one manager
- An employee can manage up to one department



# Practice


- A department can have multiple employees
- An employee works for exactly one department
- A department has exactly one manager
- An employee can manage up to one department



# Practice

Same as before, but what if

- A department has at least one employee!!!



# Try it Out

- Good options

- draw.io – free!
- lucidchart – better, but not free

- Working options

- Powerpoint
- Illustrator
- MS Visio (online only for macs)
- MS Paint