

# CS 6016



## Database Systems

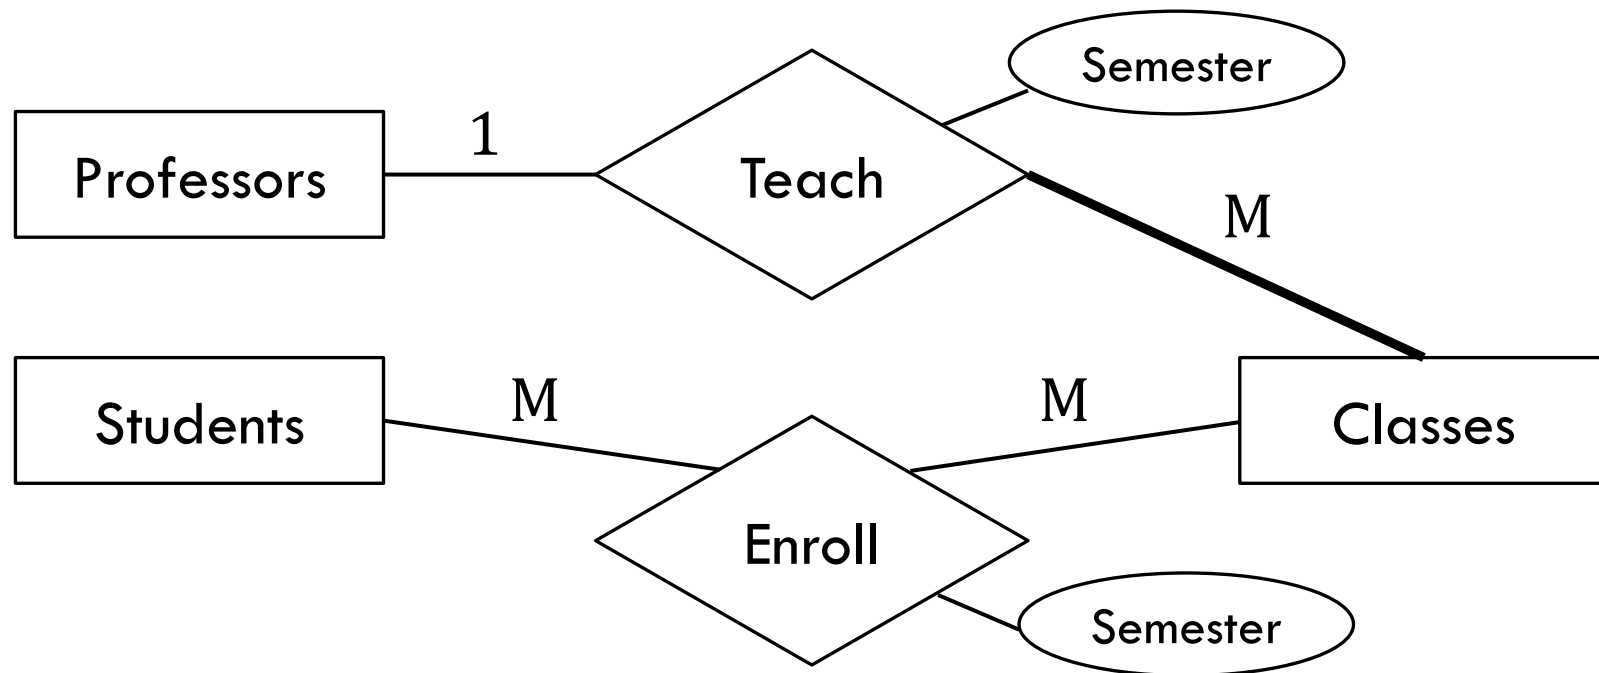
*ER Model Cont.*

*SQL Tables*

*Reducing ER to Schema*

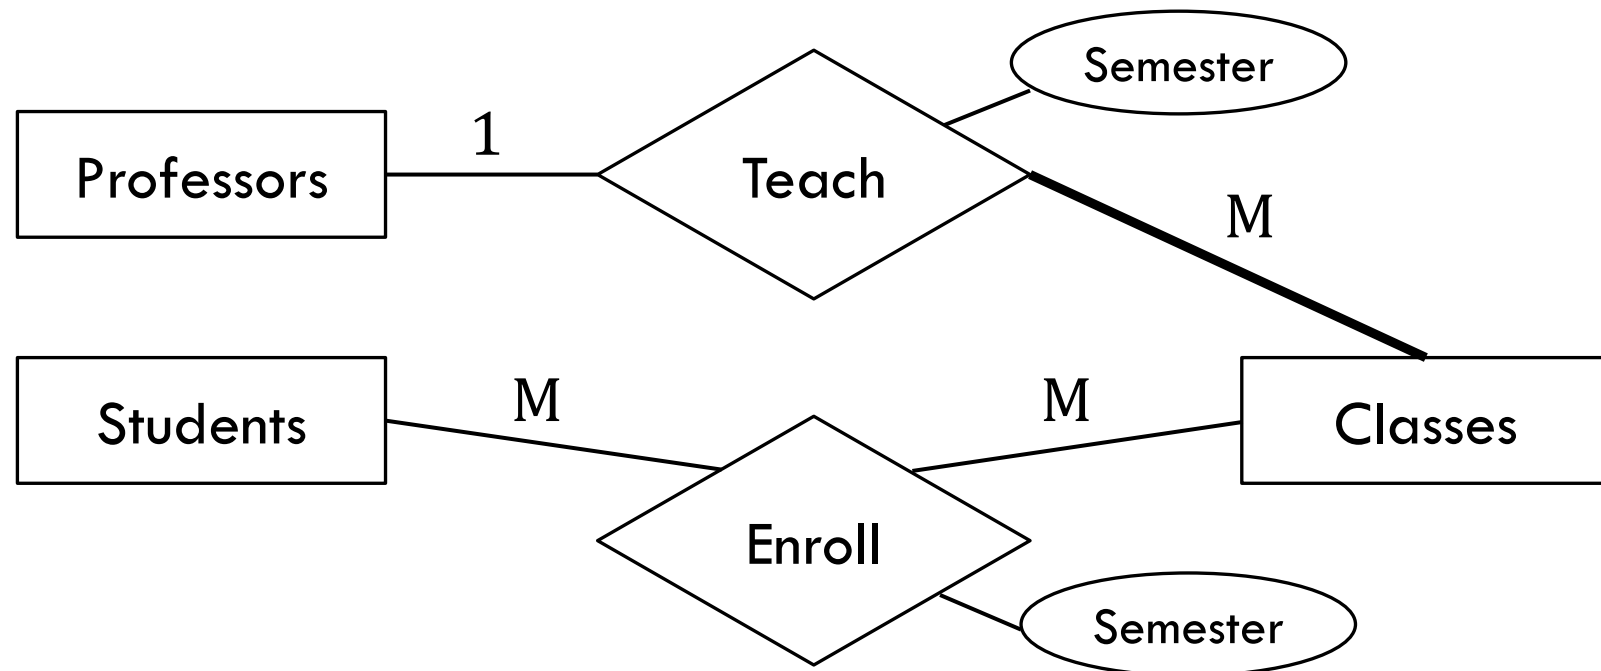
# Practice (Translate Diagram)

- Can a professor take sabbatical?
- Can a class be co-taught?
- Can a class have no teacher?
- Can a class have no students?
- Can a student take multiple classes?
- Can a student take the semester off?



# Practice (Translate Diagram)

- Can a professor take sabbatical? yes
- Can a class be co-taught? no
- Can a class have no teacher? no
- Can a class have no students? yes
- Can a student take multiple classes? yes
- Can a student take the semester off? yes



# Hierarchical Types

```
class Employee { int SSN; string name; }
```

```
class HourlyEmployee extends Employee  
{ float wage; }
```

```
class SalaryEmployee extends Employee  
{ float salary; Benefits b; }
```

# Hierarchical Types

```
class Employee { int SSN; string name; }
```

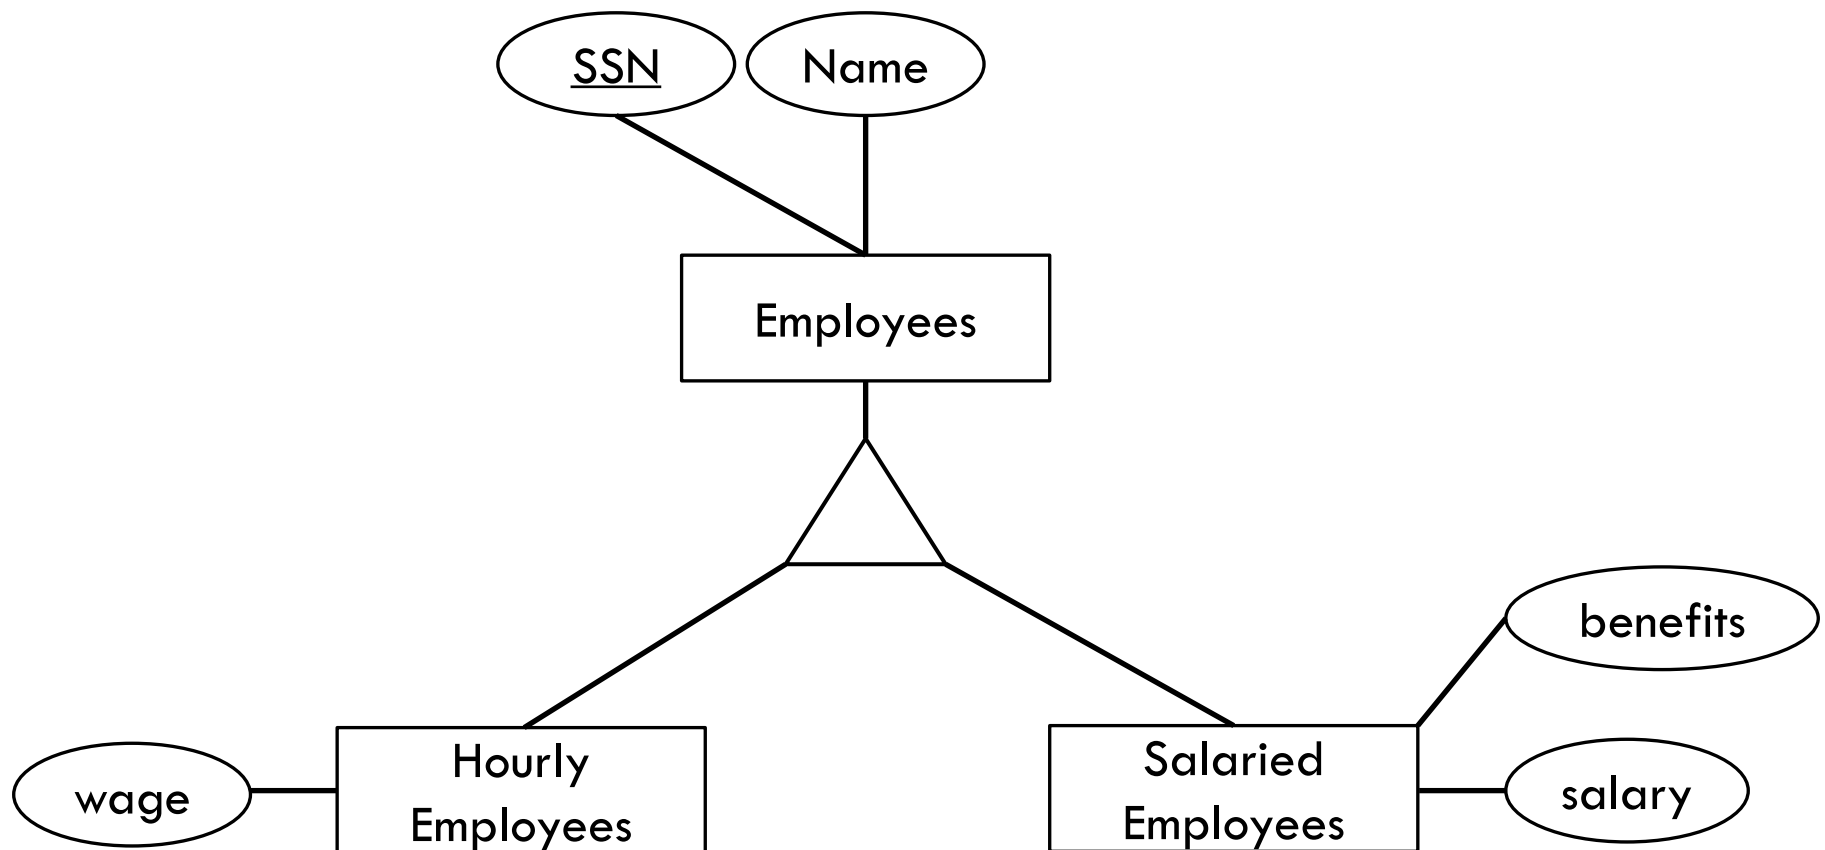
```
class HourlyEmployee extends Employee  
{ float wage; }
```

```
class SalaryEmployee extends Employee  
{ float salary; Benefits b; }
```

- HourlyEmployee “IS-A” Employee



# “IS-A”

△ = IS-A



# Entity or Attribute?

- An employee has an address



# Entity or Attribute?

- It's usually obvious:
  - A student is an entity
  - A student ID is **not** an entity



# Entity or Attribute?

- If it's not obvious, a few questions to ask:
  - Is it complex data that needs its own keys?

# Entity or Attribute?


- If it's not obvious, a few questions to ask:
  - Is it complex data that needs its own keys?
  - Does the data type make sense in another relationship?

# Entity or Attribute?

- If it's not obvious, a few questions to ask:
  - Is it complex data that needs its own keys?
  - Does the data type make sense in another relationship?
  - Can there be more than one?
- “yes” to these usually argues for an **entity**
  - But not always!

# Entity or Attribute?


- Track full residence history



- Now an employee can have multiple addresses

# Ponder

- What if the employee lives at the same residence two different times?




5/1/16 – 5/1/17, 123 Fake Street

5/1/17 – 5/1/18, 555 Creek Rd.

5/1/18 – 5/1/19, 123 Fake Street

# Ponder


- What if the employee lives at the same residence two different times?



Many-to-Many


# Ponder

- What if the employee lives at the same residence two different times?



Many-to-Many


# Ternary Relation






# Entity vs. Relationship

- How to get a new entity out of a relationship?
- An employee can manage multiple departments
- With a different budget for each department



# Entity vs. Relationship


- An employee can manage multiple departments
- With a different budget for each department



- What if the manager has just one budget to split?

# Entity vs. Relationship


- An employee can manage multiple departments
- With a different budget for each department



- What if the manager has just one budget to split?

# Entity vs. Relationship


- An employee can manage multiple departments
- With a different budget for each department



- What if the manager has just one budget to split?
  - **Bad: not all employees are managers**

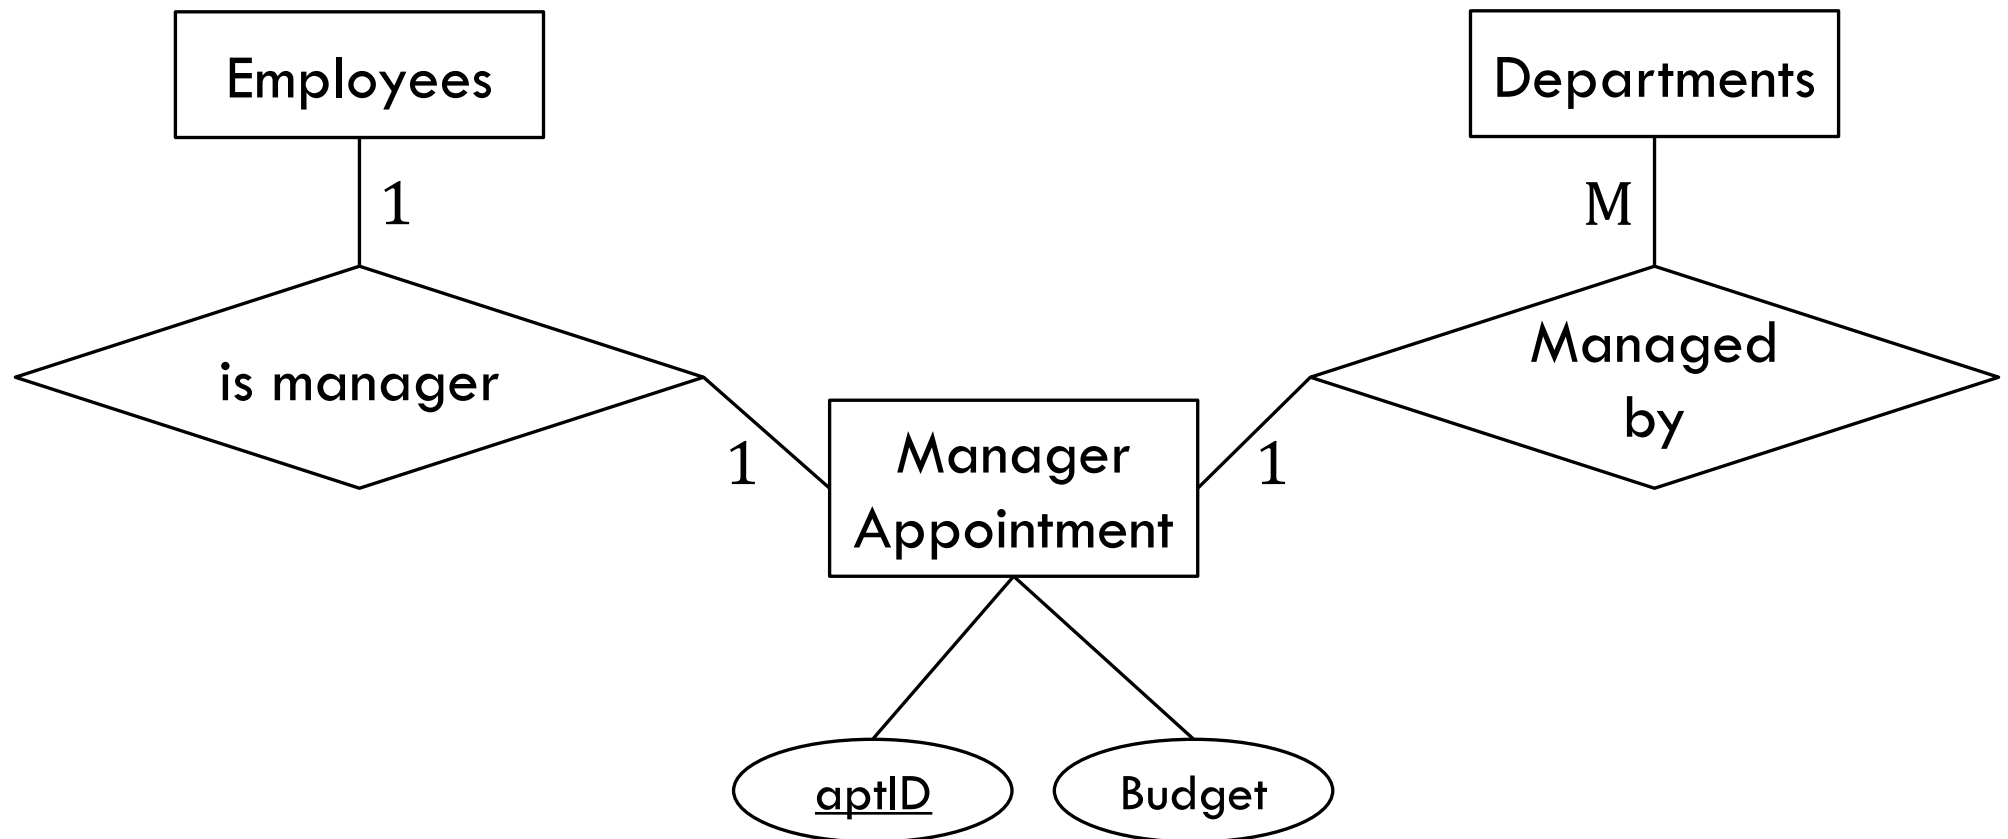
# Entity vs. Relationship

- An employee can manage multiple departments
- With a different budget for each department

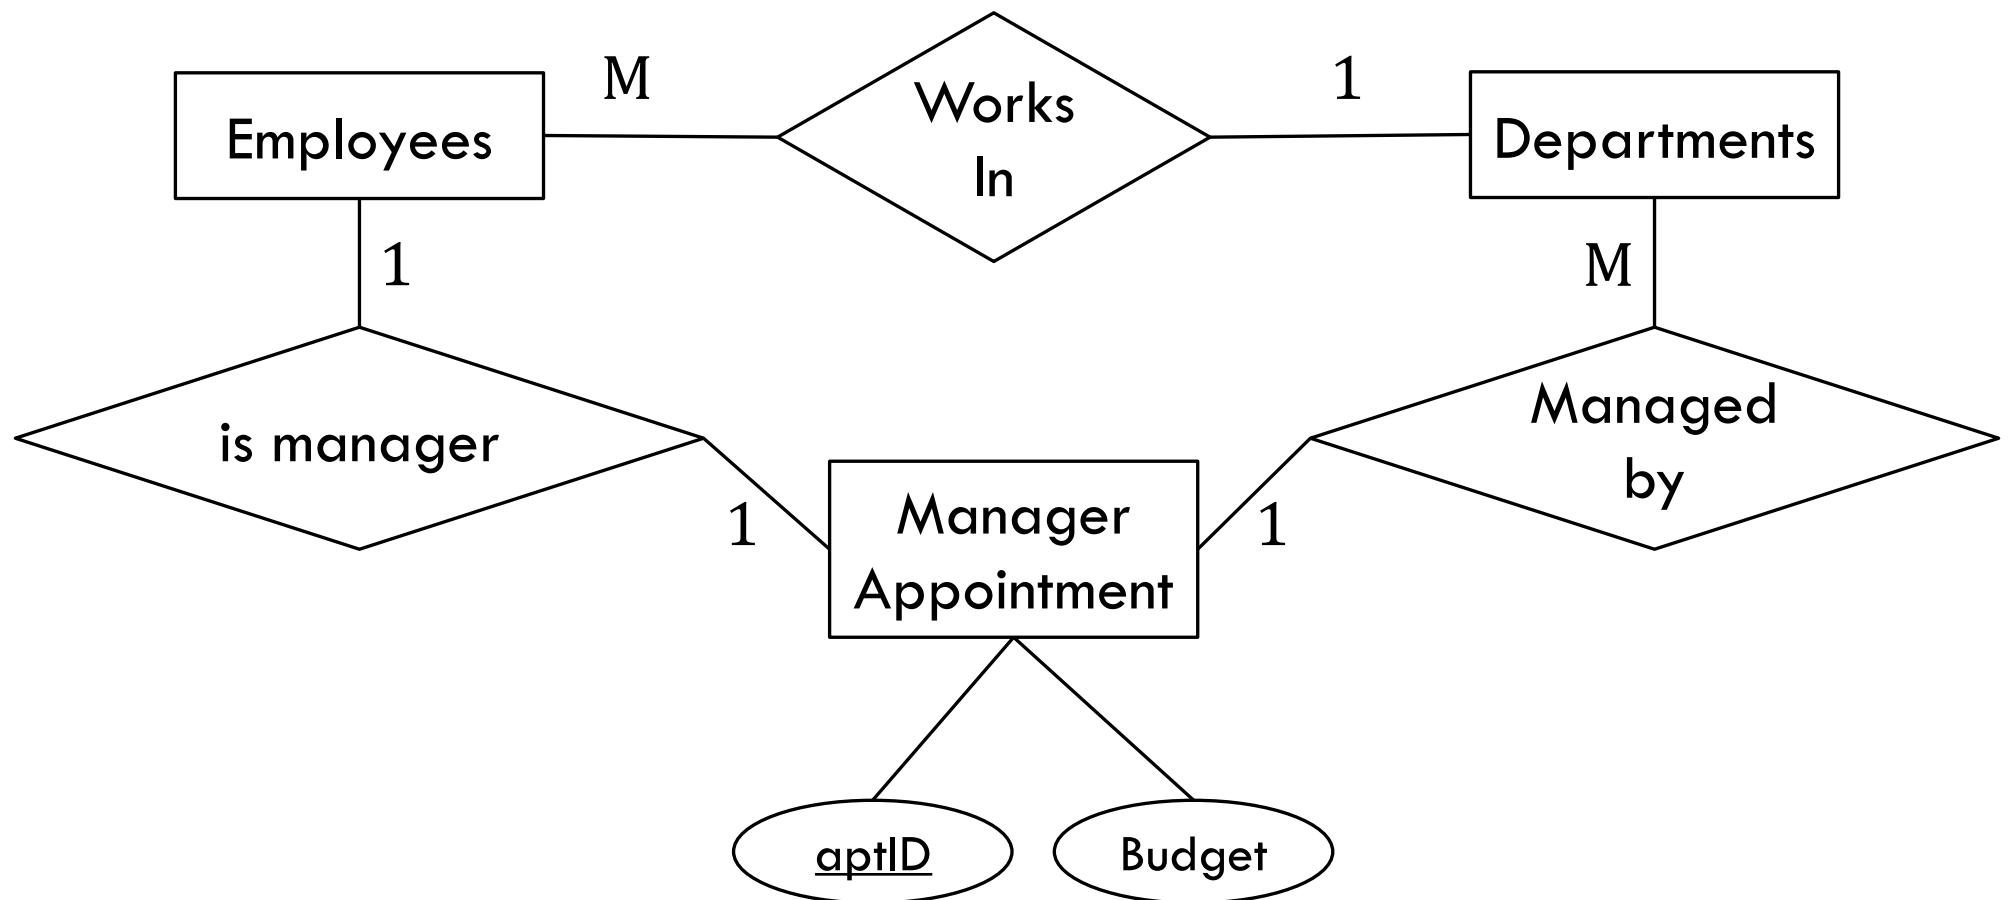


- **What we need:** something that defines a management role

# Entity vs. Relationship



# Almost Complete ER



# Participation Constraints

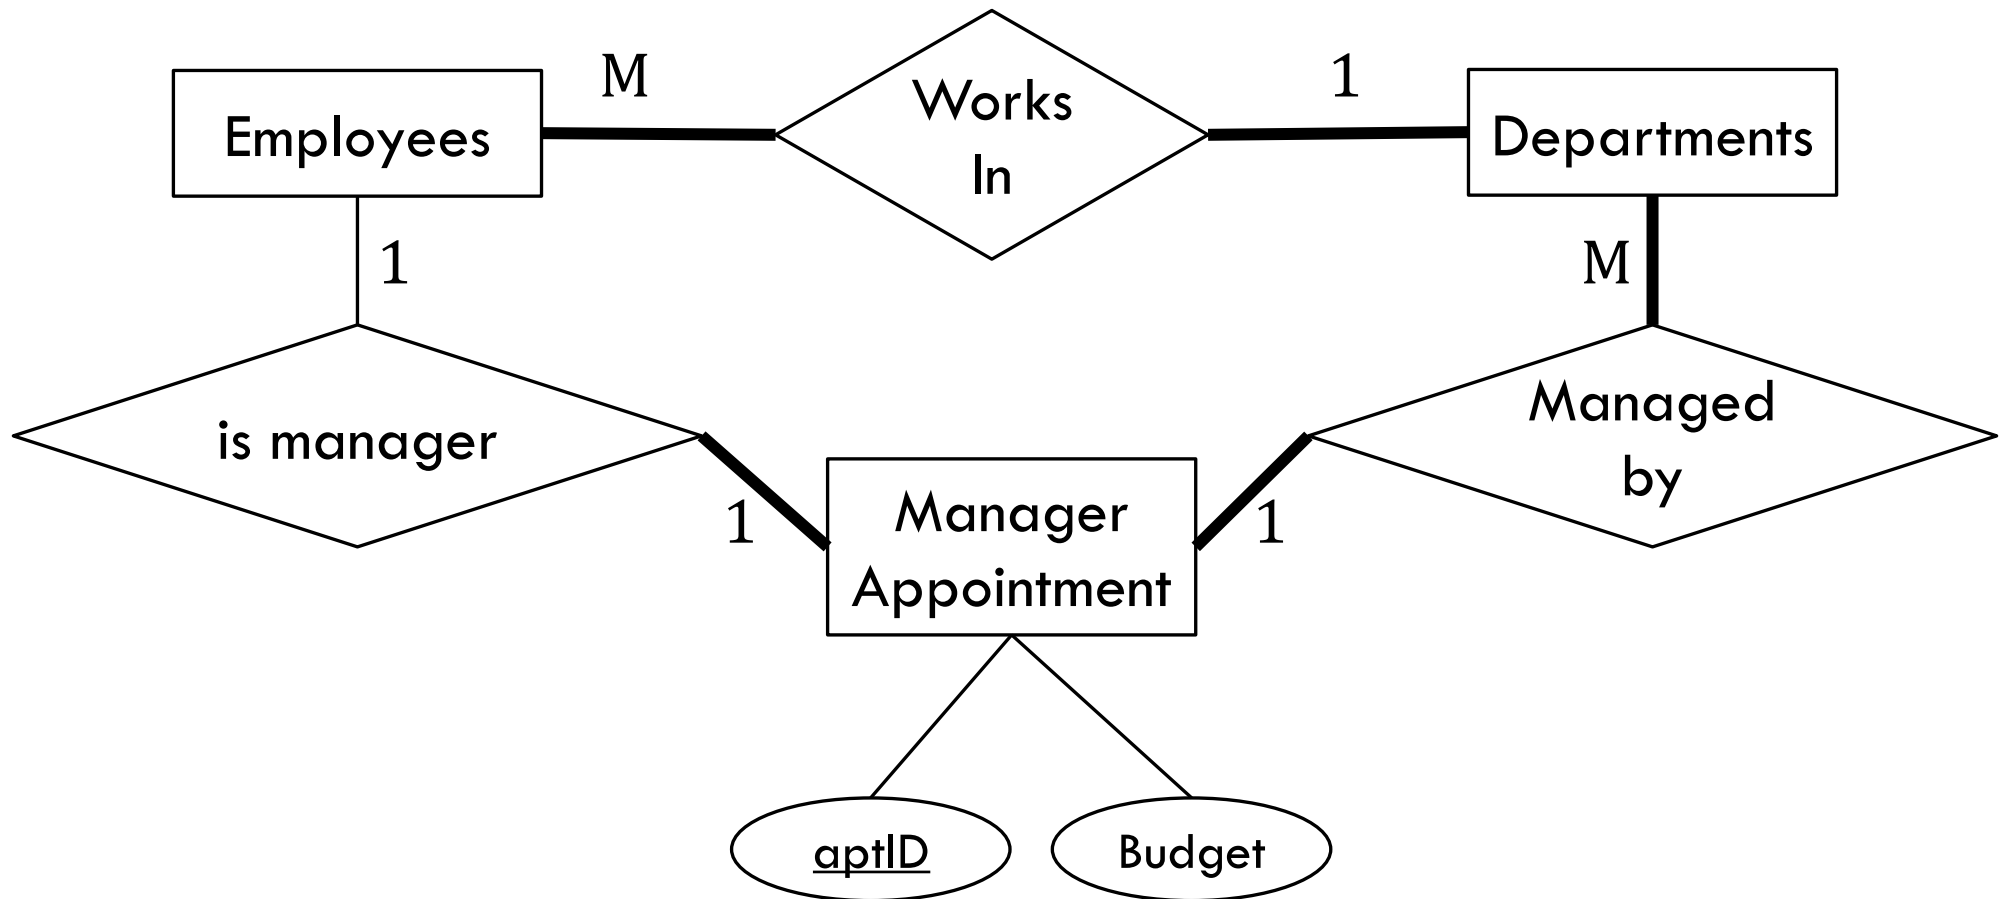
1. Does a department have to have a manager?
2. Does a manager appointment have to be associated with at least one department?
3. Does an employee have to be a manager?
4. Does a manager appointment have to have a person associated with it?
5. Does an employee have to work for a department?
6. Does a department have to have “at least one” employee?



# Participation Constraints

1. Does a department have to have a manager? **Yes, so bold.**
2. Does a manager appointment have to be associated with at least one department? **Yes, so bold.**
3. Does an employee have to be a manager? **No, so no bold.**
4. Does a manager appointment have to have a person associated with it? **Yes, so bold.**
5. Does an employee have to work for a department? **Yes, so bold.**
6. Does a department have to have “at least one” employee? **Yes, so bold.**

# Almost Complete ER



# Weak Entities

- A **weak entity** can't be identified by its own attributes
- It is identified by a combination of:
  - its own attribute(s)
  - and a foreign key (something else's key attribute)
- Can not be uniquely identified by its own attributes

# Weak Entities

- A **weak entity** can't be identified by its own attributes
- It is identified by a combination of:
  - its own attribute(s)
  - and a foreign key (something else's key attribute)
- i.e., its existence only makes sense in the context of another entity

# Example


- Consider the difference between a *course* and a *class*
  - **Course**
    - Subject
    - Number
    - Name
    - Description
- CS 6016 Database Systems, “in this class we study ...”,

# Example

- Consider the difference between a *course* and a *class*
  - **Class**
    - Semester
    - <reference to course>
    - <reference to teacher>
- Spring 21, <reference to CS6016>, <reference to teacher>


# Weak Entity

- Classes is a **weak entity** set
- Instance of is a **supporting relationship** set



# Weak Entity

- Classes is a **weak entity** set
  - Double rectangle for the entity set
  - Double diamond for the **supporting relationship**
  - Supporting relationship must be 1-to-M, and weak entity must participate






# Weak Entities

- A **weak entity** can't be identified by its own attributes
- It is identified by a combination of:
  - its own attribute(s)
  - and a foreign key


# Partial Key

- A weak entity set has a **partial key**
  - Dashed underline



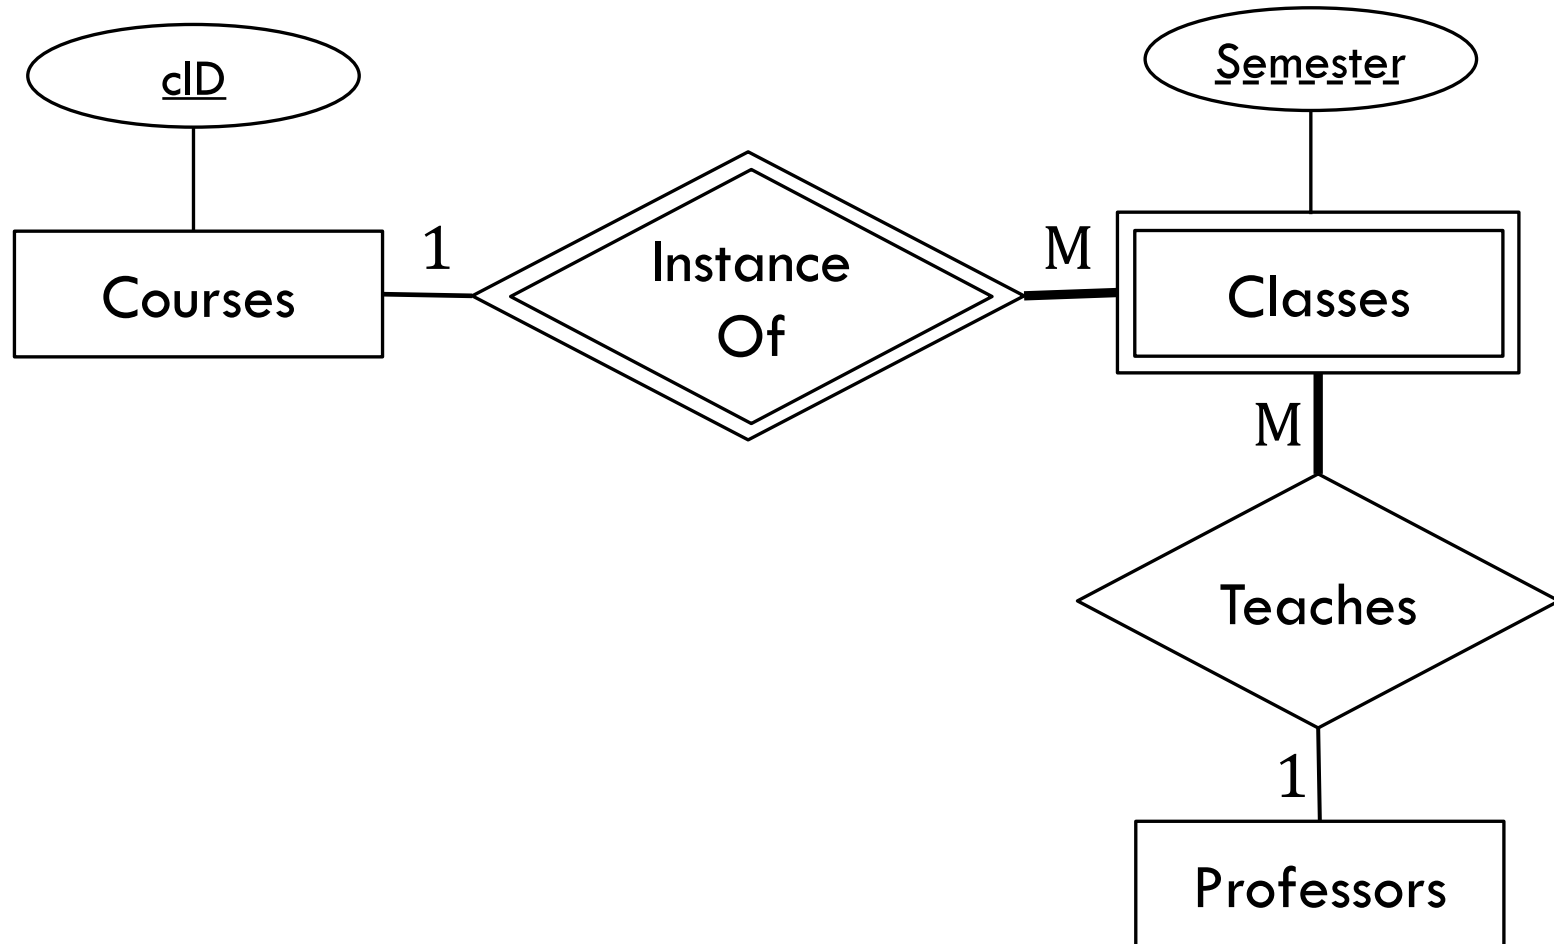
# Partial Key

- A weak entity set has a **partial key**
  - Dashed underline
  - Must be combined with a foreign key
  - {cID, Semester} is the key for Classes



# Partial Key

- Weak entity sets can have non-supporting relationships



# Creating Tables in SQL

```
create table <name> (  
    <column1Name> <type> <properties>,  
    <column2Name> <type> <properties>,  
    ...  
    <table properties>  
) ;
```

- Properties are optional

# MySQL Data Types

- Numeric

- Integers – int, <tiny, small, medium, big>int, <unsigned>
- Reals – float, double, decimal

# MySQL Data Types

- Numeric
  - Integers – int, <tiny, small, medium, big>int, <unsigned>
  - Reals – float, double, decimal
- Dates – very important in DBs
  - Date, datetime, time stamp, time, year

# MySQL Data Types

- Numeric
  - Integers – int, <tiny, small, medium, big>int, <unsigned>
  - Reals – float, double, decimal
- Dates – very important in DBs
  - Date, datetime, time stamp, time, year
- Strings
  - char(m), varchar(m)



# MySQL Data Types

- Numeric
  - Integers – int, <tiny, small, medium, big>int, <unsigned>
  - Reals – float, double, decimal
- Dates – very important in DBs
  - Date, datetime, time stamp, time, year
- Strings
  - char(m), varchar(m)
- Blobs – Binary Large Objects
- Enums

# Strings

- char(N) - exactly N characters
- varchar(N) - up to N characters

# Quiz

- Best type for 'CardNum'?

CheckedOut

CardNum	Serial
1	1001
1	1004
4	1005
4	1006

# Quiz

- Best type for 'CardNum'?

**INT UNSIGNED**

CheckedOut

CardNum	Serial
1	1001
1	1004
4	1005
4	1006

# Quiz

- Best type for ‘Author’?

## Titles

ISBN	Title	Author
978-0590353427	Harry Potter	Rowling
978-0679732242	The Sound and the Fury	Faulkner
978-0394823379	The Lorax	Seuss
978-0062278791	Profiles in Courage	Kennedy
978-0441172719	Dune	Herbert

# Quiz

- Best type for 'Author'?  
`varchar(...)`

## Titles

ISBN	Title	Author
978-0590353427	Harry Potter	Rowling
978-0679732242	The Sound and the Fury	Faulkner
978-0394823379	The Lorax	Seuss
978-0062278791	Profiles in Courage	Kennedy
978-0441172719	Dune	Herbert

# Quiz

- Best type for ‘ISBN’?

## Titles

ISBN	Title	Author
978-0590353427	Harry Potter	Rowling
978-0679732242	The Sound and the Fury	Faulkner
978-0394823379	The Lorax	Seuss
978-0062278791	Profiles in Courage	Kennedy
978-0441172719	Dune	Herbert

# Quiz

- Best type for 'ISBN'?  
`char(14)`  
unless we get rid of the dashes

## Titles

ISBN	Title	Author
978-0590353427	Harry Potter	Rowling
978-0679732242	The Sound and the Fury	Faulkner
978-0394823379	The Lorax	Seuss
978-0062278791	Profiles in Courage	Kennedy
978-0441172719	Dune	Herbert



# Strings

- CHAR(N) - exactly N characters
  - ISBN
  - Phone number
  - We could use BIGINT UNSIGNED if we left out the dash '-'
- VARCHAR(N) - up to N characters
  - Title
  - Author
  - Name

# Strings

- How to pick (N) for VARCHAR(N)?
- Requires N bytes for storage plus 1 or 2 bytes for size
  - 1 size byte if  $N \leq 255$
  - 2 size bytes if  $N > 255$

# Strings

- How to pick (N) for VARCHAR(N)?
- Requires N bytes for storage plus 1 or 2 bytes for size
  - 1 size byte if  $N \leq 255$
  - 2 size bytes if  $N > 255$
- Pick the smallest value that is always large enough

# Properties

- Column properties:

- `NOT NULL`
- `DEFAULT 'hello'`
- `AUTO_INCREMENT`

- Table properties:

- `PRIMARY KEY (column1, ...)`
- `UNIQUE (column1, ...)`

# Not Null

- Improves index optimizations (faster tree-based searches on columns)
- ... it's not just for table design
- Specifying NOT NULL is not required on key values
  - But SQL will automatically set it

# Quiz

- Appropriate properties for 'CardNum'?

Patrons

Name	CardNum
Joe	1
Ann	2
Ben	3
Dan	4

# Quiz

- Appropriate properties for 'CardNum'?

Patrons

Name	CardNum
Joe	1
Ann	2
Ben	3
Dan	4

- not null
- auto\_increment
- primary key

# Quiz

- Appropriate properties for 'CardNum'?

CheckedOut

CardNum	Serial
1	1001
1	1004
4	1005
4	1006



# Quiz

- Appropriate properties for 'CardNum'?

CheckedOut

CardNum	Serial
1	1001
1	1004
4	1005
4	1006

- not null

# Creating Tables

- Let's create this table (without the contents)

Titles

ISBN	Title	Author
978-0590353427	Harry Potter	Rowling
978-0679732242	The Sound and the Fury	Faulkner
978-0394823379	The Lorax	Seuss
978-0062278791	Profiles in Courage	Kennedy
978-0441172719	Dune	Herbert

# Creating Tables

```
create table Titles (  
    ISBN char(14) not null,  
    Title varchar(255) not null,  
    Author varchar(255) not null,  
    primary key(ISBN)  
);
```

Titles

ISBN	Title	Author
978-0590353427	Harry Potter	Rowling
978-0679732242	The Sound and the Fury	Faulkner
978-0394823379	The Lorax	Seuss
978-0062278791	Profiles in Courage	Kennedy
978-0441172719	Dune	Herbert

# Exercise

- Command to create this table? (without the contents)

Inventory

Serial	ISBN
1001	978-0590353427
1002	978-0590353427

```
create table <name> (  
    <column1Name> <type> <properties>,  
    <column2Name> <type> <properties>,  
    <table properties>  
) ;
```

# Solution

```
create table Inventory (  
    Serial int unsigned not null auto_increment,  
    ISBN char(14) not null,  
    primary key(Serial)  
);
```

Inventory

Serial	ISBN
1001	978-0590353427
1002	978-0590353427

# Foreign Keys

```
FOREIGN KEY (<column>) REFERENCES  
<table>(<table's key>)  
ON DELETE <action>  
ON UPDATE <action>
```

# Foreign Keys

```
FOREIGN KEY (<column>) REFERENCES  
<table>(<table's key>)  
ON DELETE <action>  
ON UPDATE <action>
```

- <action> can be:
  - RESTRICT (Default): disallow the change
  - CASCADE: also delete/update in child table
  - SET NULL: nullify key in child table
  - SET DEFAULT: set to column's default value

# Foreign Key Example

```
CREATE TABLE Phones (  
    ...  
    FOREIGN KEY (CardNum) REFERENCES Patrons (CardNum)  
    ON DELETE CASCADE  
)
```

Patrons

Name	CardNum
Joe	1
Ann	2
Ben	3
Dan	4

Phones

CardNum	Phone
1	555-5555
2	666-6666
3	777-7777
4	888-8888
4	999-9999



# ER $\rightarrow$ RM Algorithm

- We've been using a “naïve”, but correct relational model
  - Sometimes resulting in unnecessary tables

# ER $\rightarrow$ RM Algorithm

- If we start with a good ER model, the translation to a good RM is mechanical

# ER $\rightarrow$ RM Algorithm


- If we start with a good ER model, the translation to a good RM is mechanical
- Basic algorithm:
  1. Every entity set becomes a schema
  2. Relationship sets *might* become schema depending on cardinality, otherwise they are *merged*

# Entity Sets to Schema

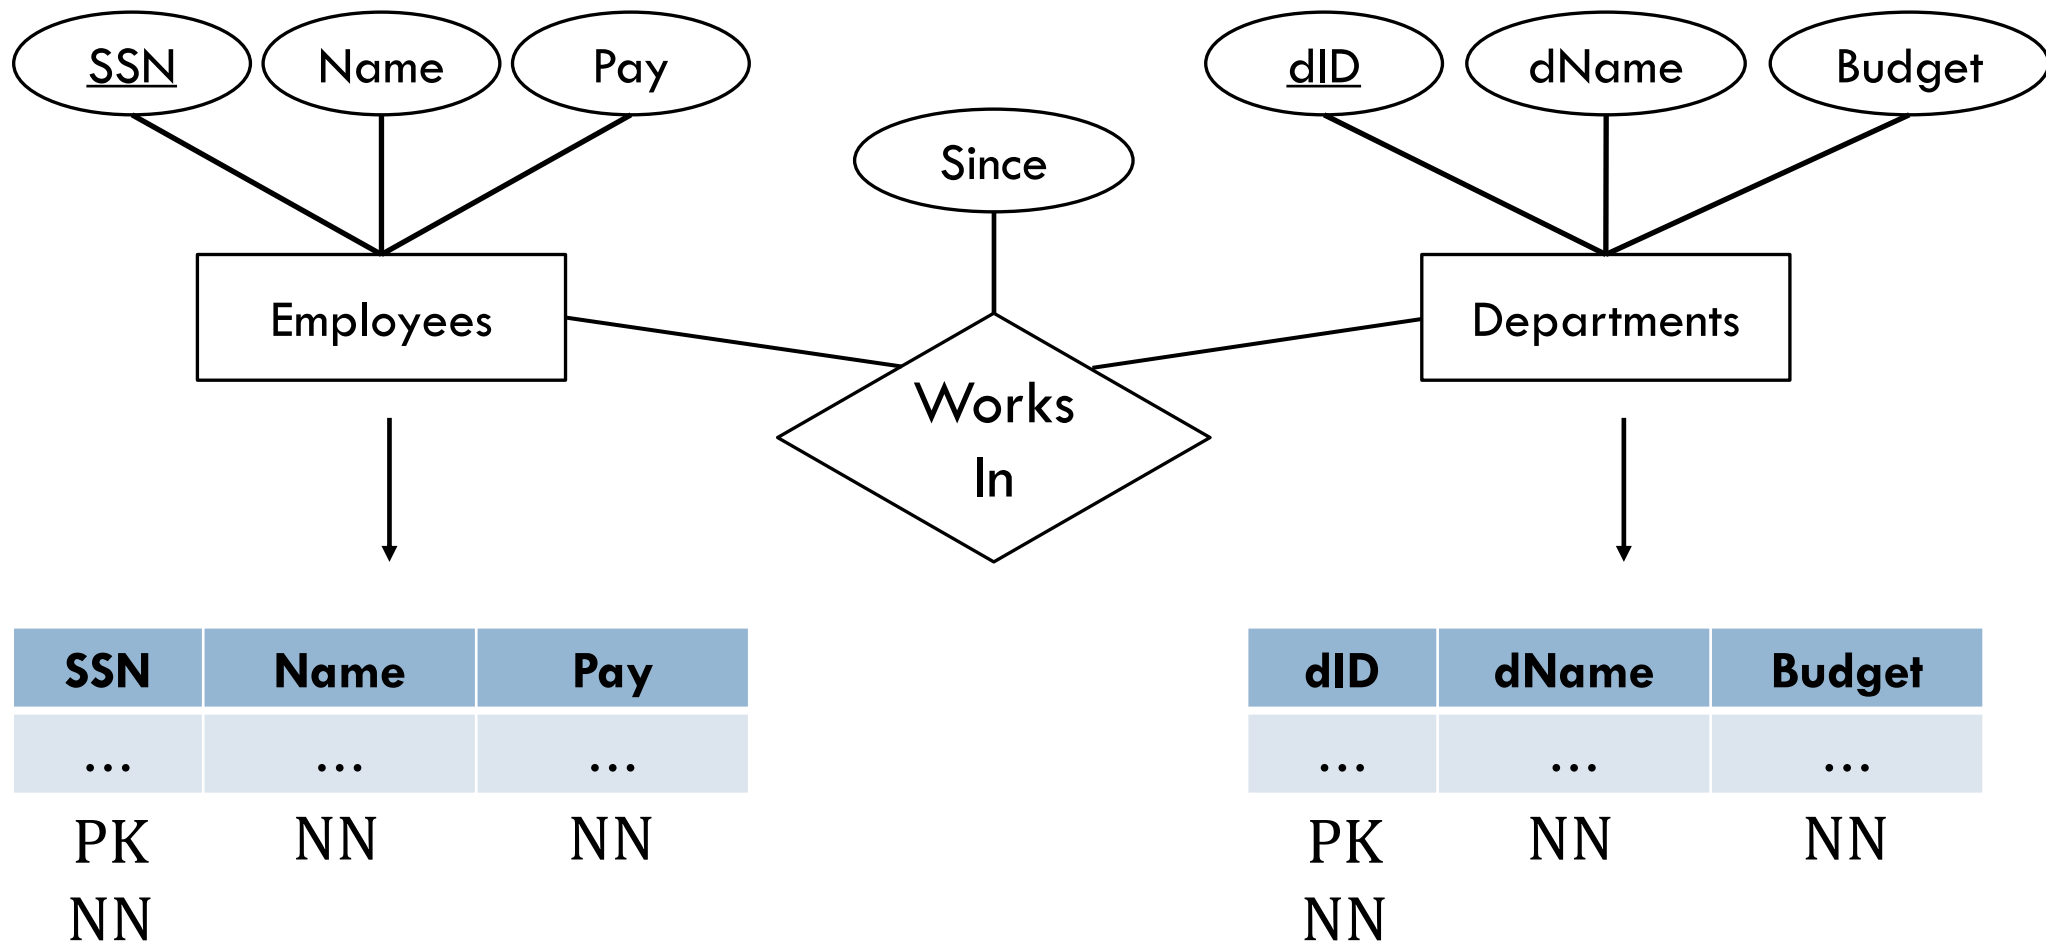
- **Entity Sets**

- Every entity set translates directly to a schema
- Attributes become columns
- Pick one of the key attribute sets as the primary key

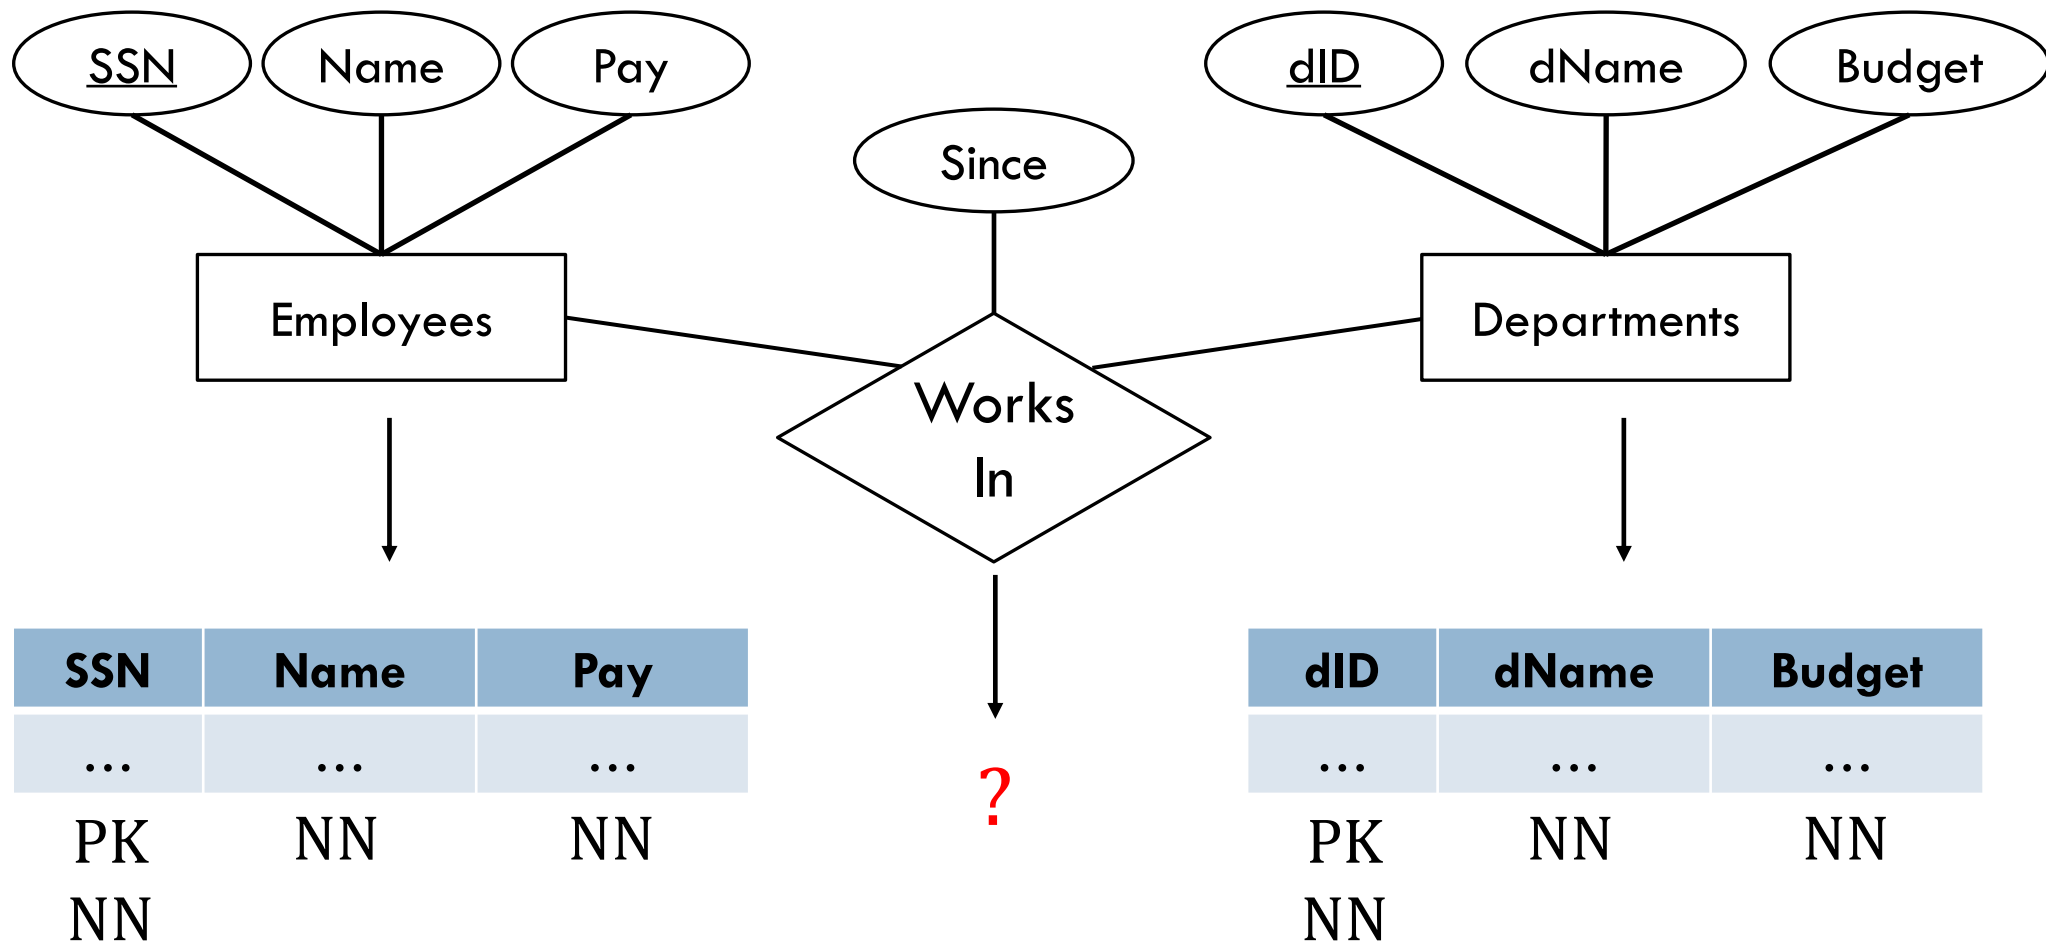
# Entity Sets to Schema



# Entity Sets to Schema




# Entity Sets to Schema



# Relationship Sets to Schema

- **Many-to-Many**

- New schema for the relationship



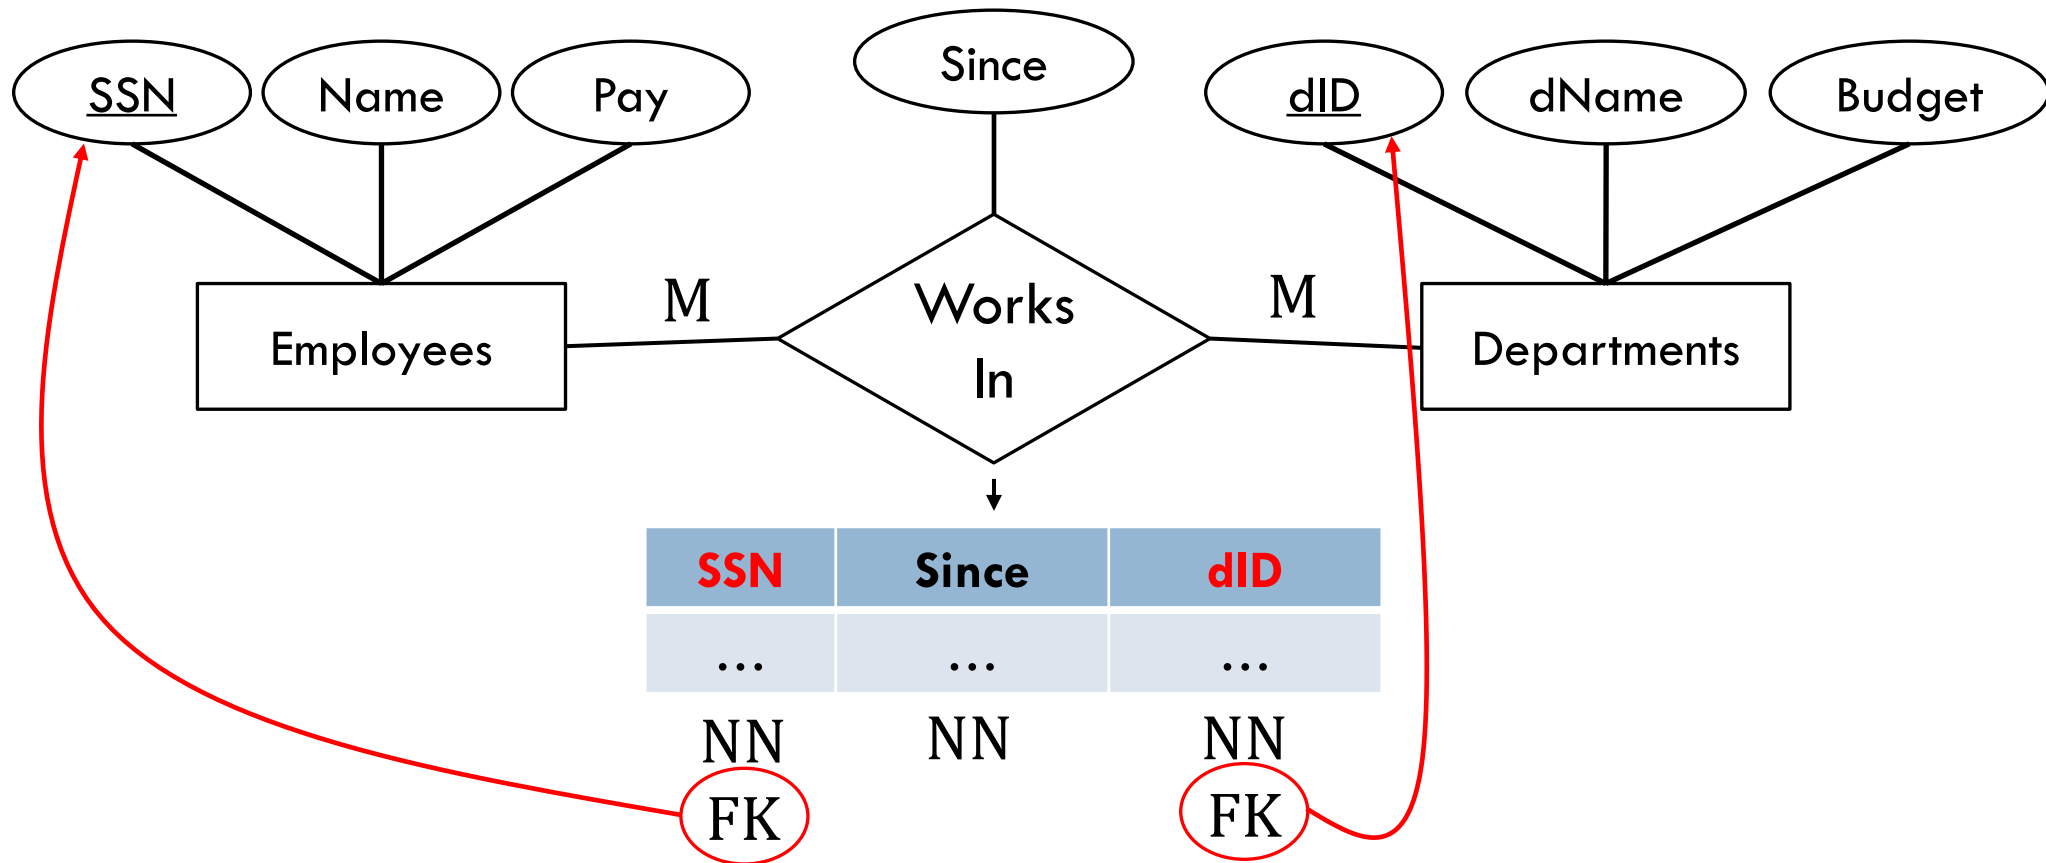
?	Since	?
...	...	...



# Relationship Sets to Schema

- **Many-to-Many**


- Primary keys of relating entities as foreign keys



# Relationship Sets to Schema

- **Many-to-Many**

- What is the key of the WorksIn table?




SSN	Since	dID
...	...	...
NN	NN	NN
FK		FK

PK?

# Relationship Sets to Schema

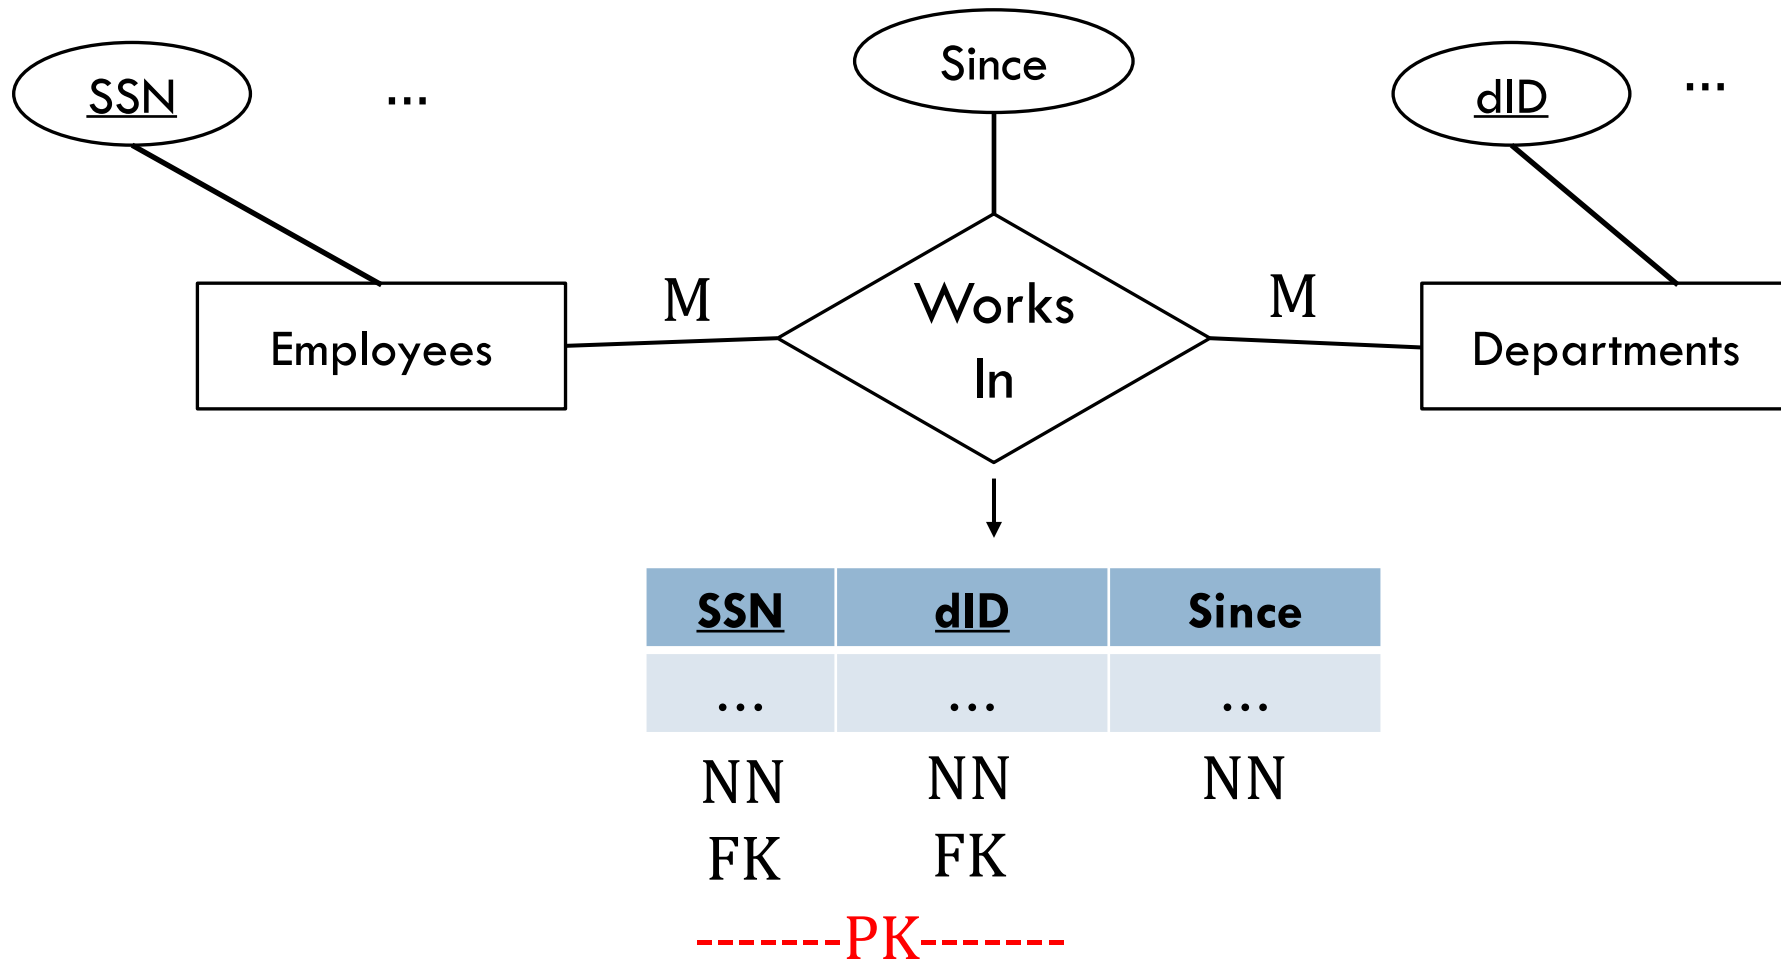
- How do we uniquely identify a line (many-to-many)?



# Relationship Sets to Schema


- **Many-to-Many**

- Key is the combination of the foreign keys



# Relationship Sets to Schema


- **1-to-Many**



# Relationship Sets to Schema

- **1-to-Many**

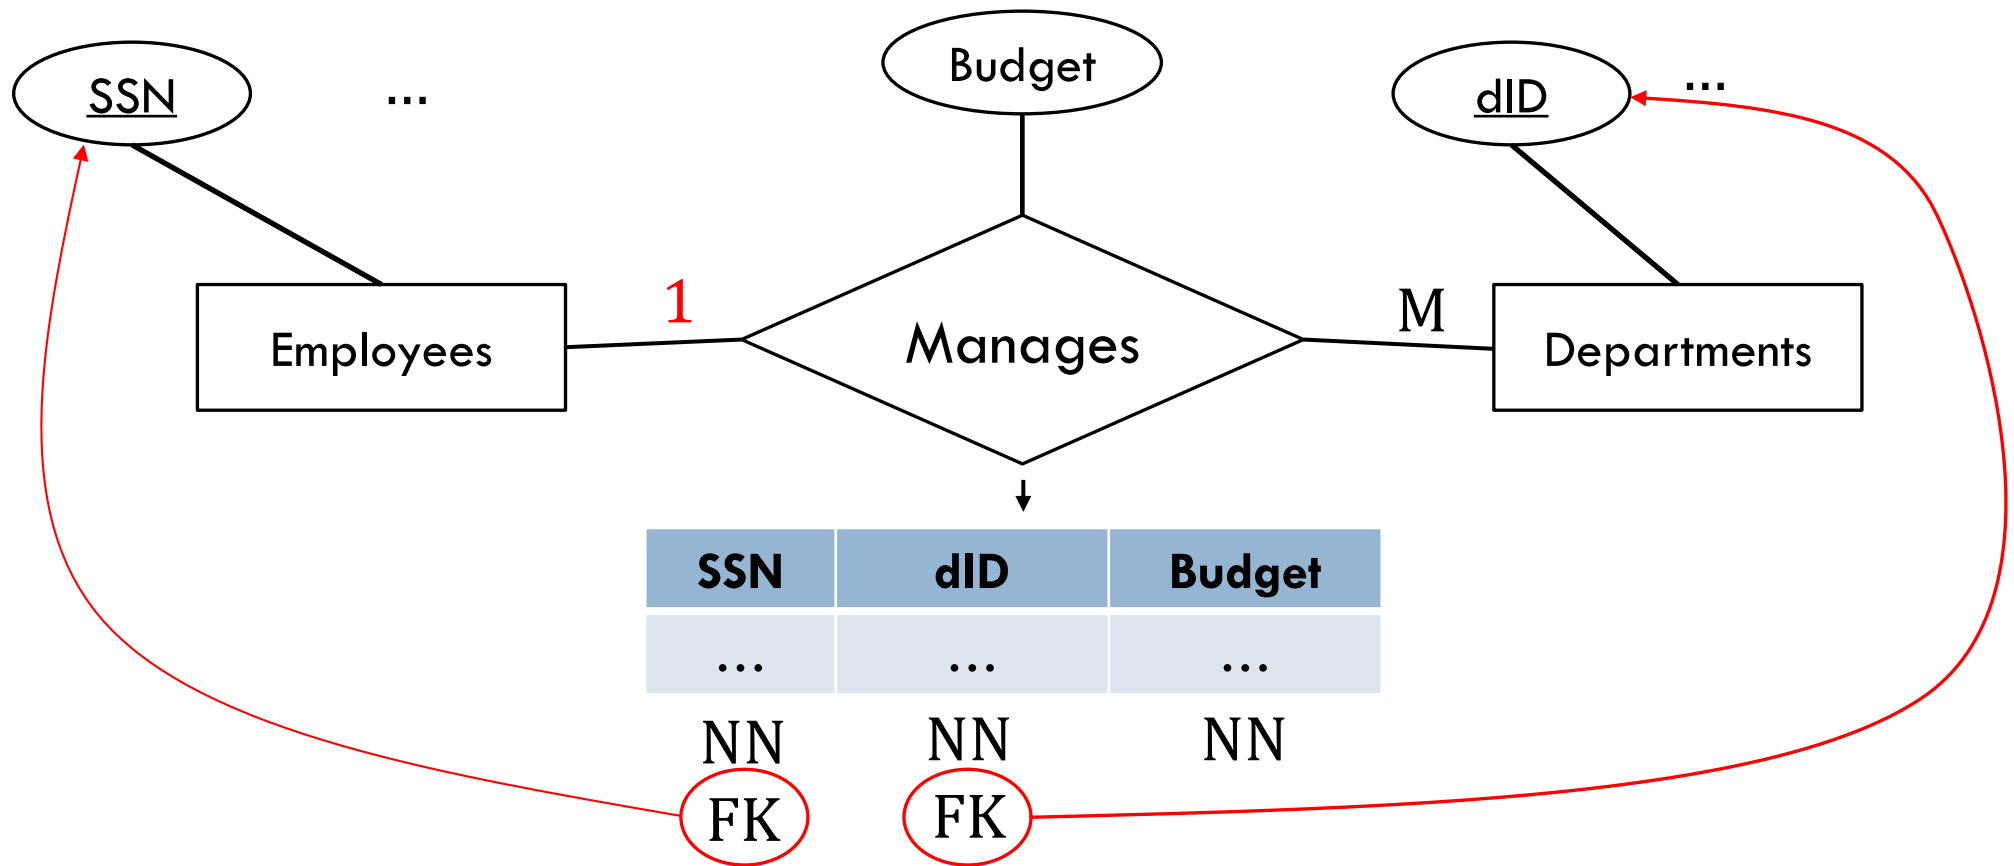
- Department can only have 1 manager



# Relationship Sets to Schema


- **1-to-Many (first attempt)**

- New schema for the relationship
- Primary key of related entities as foreign keys



# Relationship Sets to Schema

- How do we uniquely identify a line (1-to-M)?
  - “a dID can only appear once in Manages table”

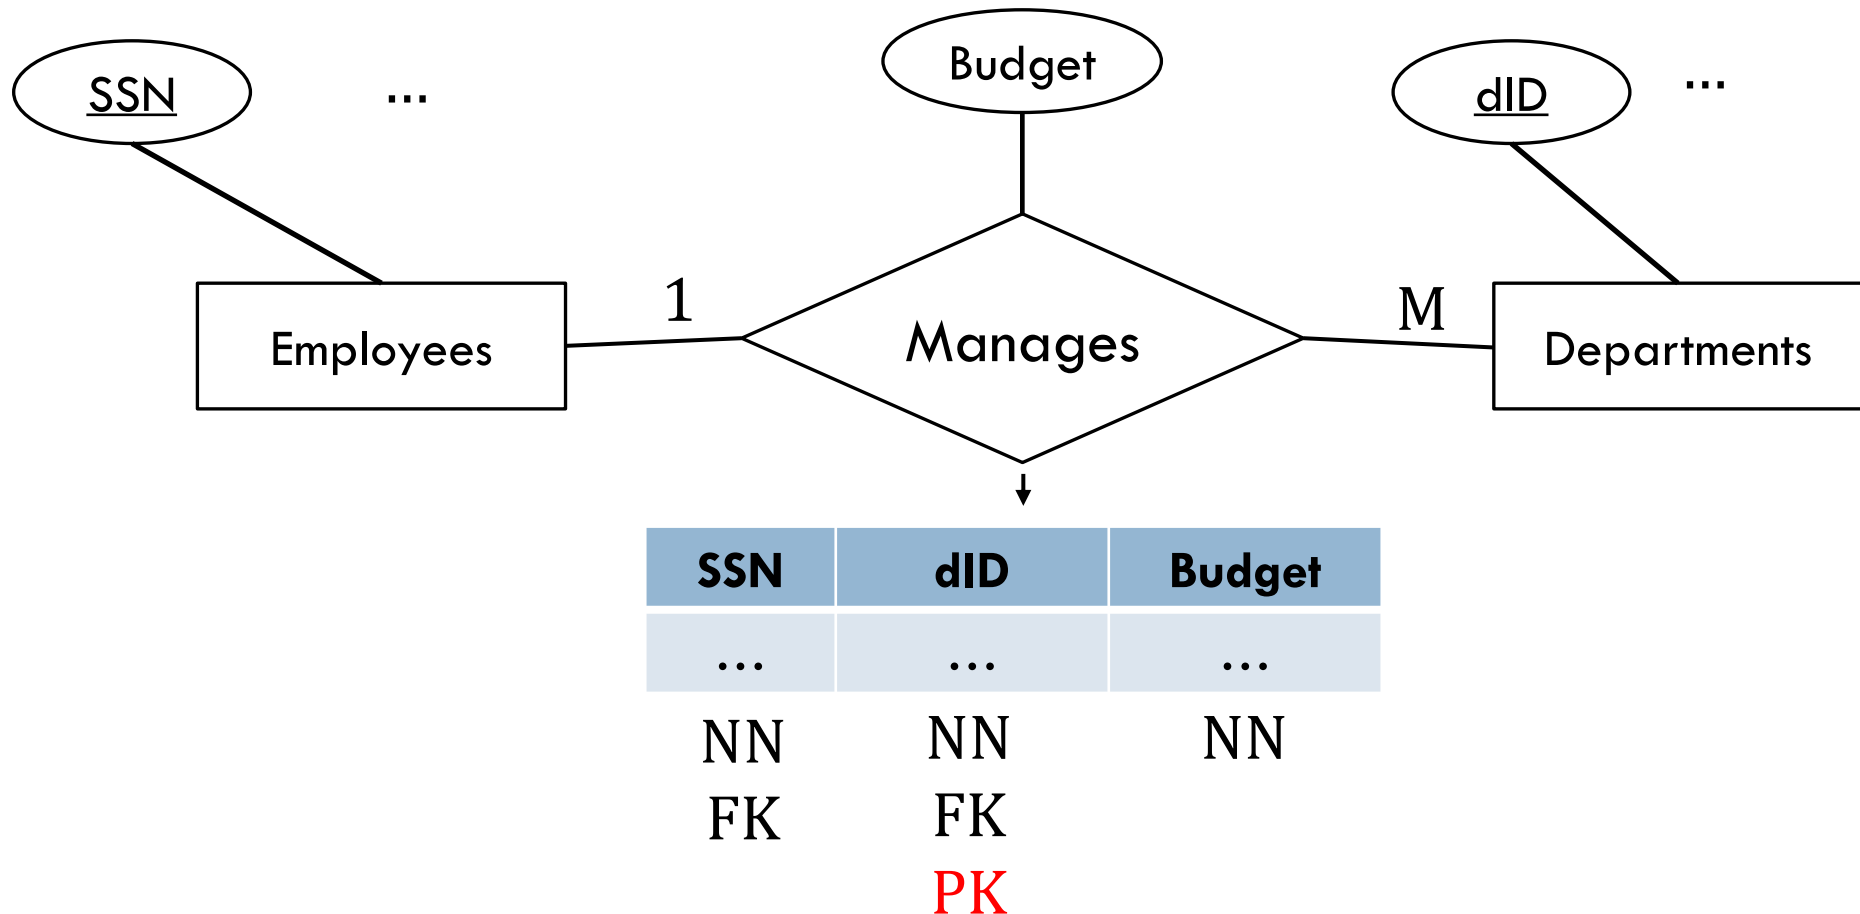




# Relationship Sets to Schema

- **1-to-Many (first attempt)**


- Key is the foreign key from the 'M' side



# Relationship Sets to Schema


- **1-to-Many (better option)**

- Department can only have 1 manager
- Merge the relationship into the Entity



# Relationship Sets to Schema

- **1-to-Many (better option)**
  - Department can only have 1 manager




SSN	Name	Pay
...	...	...
PK	NN	NN
NN		

dID	dName	Mnger	Budget
...	...		
PK	NN	FK	
NN			

# Relationship Sets to Schema

- **1-to-Many (better option)**
  - Can Mnger and Budget be NULL?



SSN	Name	Pay
...	...	...

PK  
NN

dID	dName	Mnger	Budget
...	...		


PK  
NN

FK  
NN

NN

# Participation Constraints

- If participation is required, set NOT NULL
- Else, NULL is allowed




...	Mnger	Budget
...		

FK

NN

NN



...	Mnger	Budget
...		

FK

# Foreign Key + NULL

- Referential integrity only enforced when inserting non-null data
- Not all departments have a manager

Employees

SSN	Name	Pay
1	Joe	...
2	Steve	...
3	Meg	...

Departments

dID	dName	Budget	Mnger
a	Chem	...	1
b	Phys	NULL	NULL
c	CS	...	3

FK

# Naïve Library

## Patrons

Name	CardNum
Joe	1
Ann	2
Ben	3
Dan	4

## Inventory

Serial	ISBN
1001	978-0590353427
1002	978-0590353427
1003	978-0679732242
1004	978-0394823379
1005	978-0394823379
1006	978-0062278791

## CheckedOut

CardNum	Serial
1	1001
1	1004
4	1005
4	1006

## Phones

CardNum	Phone
1	555-5555
2	666-6666
3	777-7777
4	888-8888
4	999-9999

## Titles

ISBN	Title	Author
978-0590353427	Harry Potter	Rowling
978-0679732242	The Sound and the Fury	Faulkner
978-0394823379	The Lorax	Seuss
978-0062278791	Profiles in Courage	Kennedy
978-0441172719	Dune	Herbert

# Reduced

## Inventory

Serial	ISBN	CheckedOutBy
1001	978-0590353427	1
1002	978-0590353427	NULL
1003	978-0679732242	NULL
1004	978-0394823379	1
1005	978-0394823379	4
1006	978-0062278791	4

## Patrons

Name	CardNum
Joe	1
Ann	2
Ben	3
Dan	4

## Phones

CardNum	Phone
1	555-5555
2	666-6666
3	777-7777
4	888-8888
4	999-9999

## Titles

ISBN	Title	Author
978-0590353427	Harry Potter	Rowling
978-0679732242	The Sound and the Fury	Faulkner
978-0394823379	The Lorax	Seuss
978-0062278791	Profiles in Courage	Kennedy
978-0441172719	Dune	Herbert




# Performance

- In general:
  - Fewer joins = better performance
  - Fewer tables = fewer joins
- 1-to-M and M-to-1 don't need their own tables

# Relationship Sets to Schema

- **1-to-1**

- Treat as 1-to-M
- Merge relationship into one of the other tables



SSN	...	Manages	Budget
...	...	...	...

PK  
NN

FK

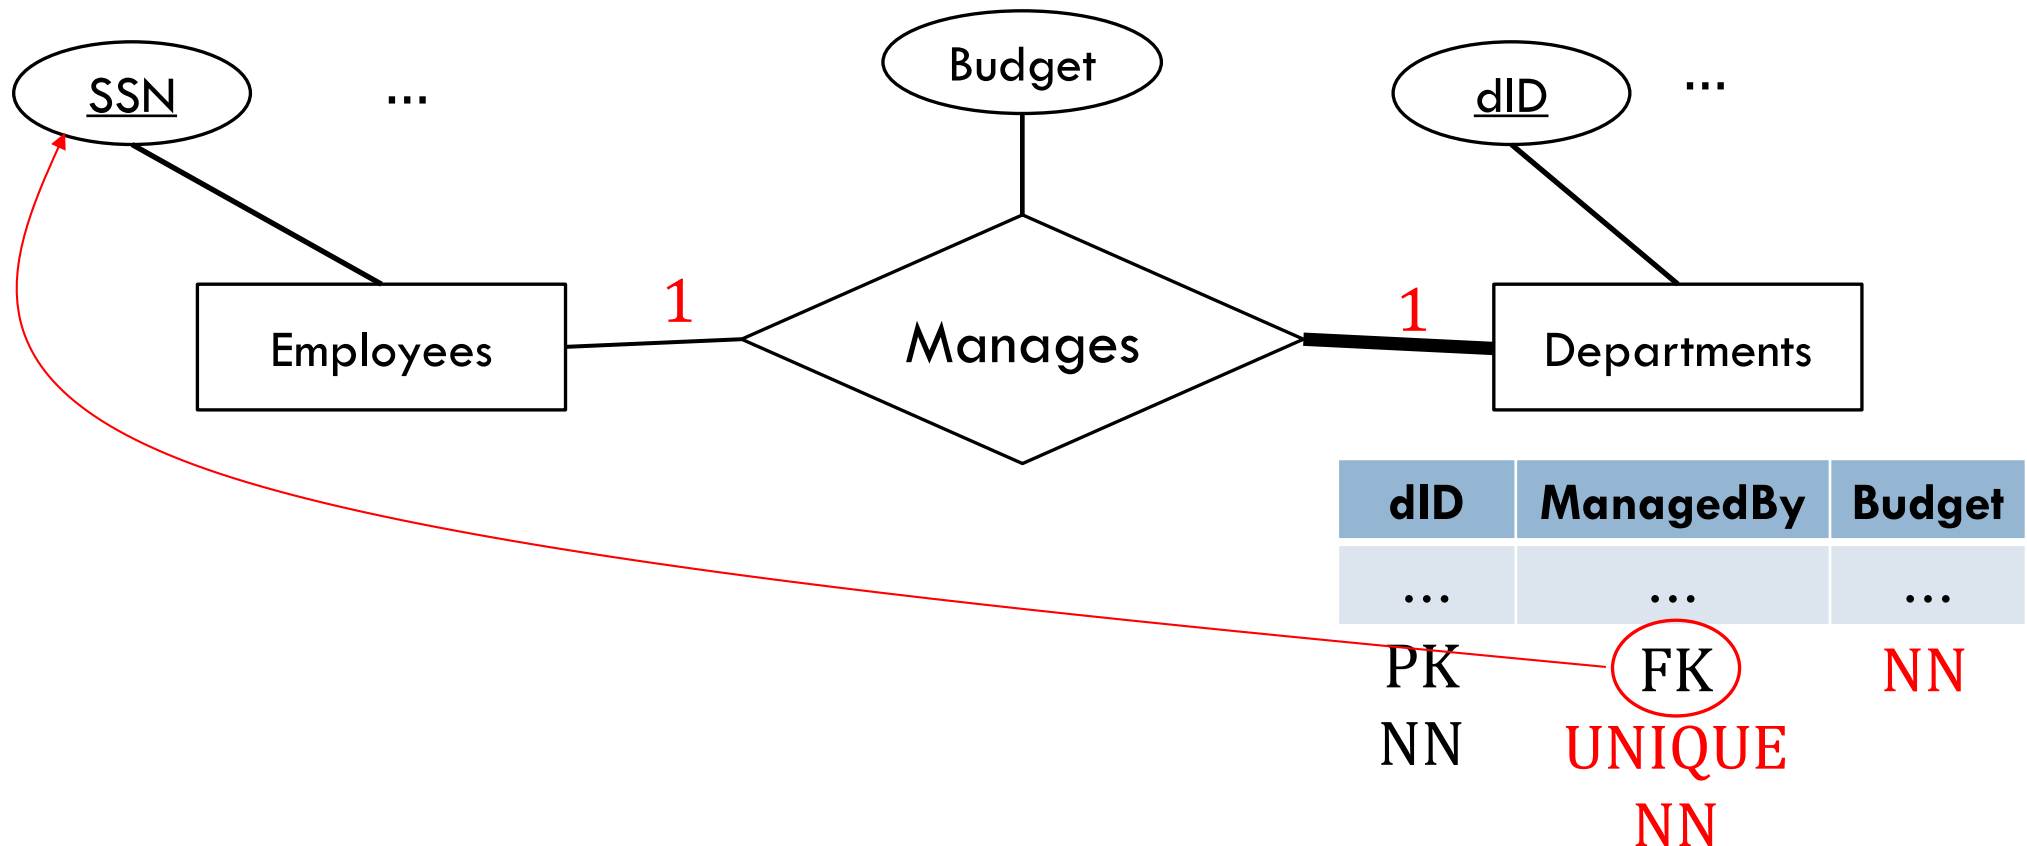
UNIQUE

dID	...
...	...

# Relationship Sets to Schema

## • 1-to-1


- If participation required on one side, use that side
- Add NOT NULL



# Relationship Sets to Schema

- **Supporting Relationship**

- A supporting relationship is **1-to-M**
- Same procedure, except supporting key is combined with partial key



Semester	cID	...
...	...	...
NN	NN	
	FK	

-----PK-----

# Double Participation

- If participation required on **both** sides...
  - Regardless of cardinality




# Double Participation

- If participation required on **both** sides...
  - Regardless of cardinality
  - **Chicken/egg problem**





- Difficult to capture with schema design
  - Instead, enforce with SQL commands

# Algorithm Summary



# Algorithm Summary

■ = FK





$$E1 = \{\underline{a_1}, \dots, a_N\} \quad E2 = \{\underline{b_1}, \dots, b_N\}$$

$$R = \{\underline{a_1}, \underline{b_1}, r_1, \dots, r_k\}$$




# Algorithm Summary

■ = FK



$$E1 = \{\underline{a_1}, \dots, a_N\} \quad E2 = \{\underline{b_1}, \dots, b_N\}$$

$$R = \{\underline{a_1}, \underline{b_1}, r_1, \dots, r_k\}$$





$$E1 = \{\underline{a_1}, \dots, a_N\}$$

$$E2 = \{\underline{b_1}, \dots, b_N, a_1, r_1, \dots, r_k\}$$


# Algorithm Summary

■ = FK




$$E1 = \{\underline{a_1}, \dots, a_N\} \quad E2 = \{\underline{b_1}, \dots, b_N\}$$

$$R = \{\underline{a_1}, \underline{b_1}, r_1, \dots, r_k\}$$



$$E1 = \{\underline{a_1}, \dots, a_N\}$$

$$E2 = \{\underline{b_1}, \dots, b_N, a_1, r_1, \dots, r_k\}$$





$$E1 = \{\underline{a_1}, \dots, a_N, b_1, r_1, \dots, r_k\}$$

$$E2 = \{\underline{b_1}, \dots, b_N\}$$


# Algorithm Summary

■ = FK




$$E1 = \{\underline{a_1}, \dots, a_N\} \quad E2 = \{\underline{b_1}, \dots, b_N\}$$

$$R = \{\underline{a_1}, \underline{b_1}, r_1, \dots, r_k\}$$




$$E1 = \{\underline{a_1}, \dots, a_N\}$$

$$E2 = \{\underline{b_1}, \dots, b_N, a_1, r_1, \dots, r_k\}$$



$$E1 = \{\underline{a_1}, \dots, a_N, b_1, r_1, \dots, r_k\}$$

$$E2 = \{\underline{b_1}, \dots, b_N\}$$




Treat as 1:M or M:1  
Mark foreign key as unique

# Algorithm Summary

- NOT NULL determined by participation constraints
- Total participation on both sides – don't capture with schema

# Relationships with Arity $> 2$



Relation might be “1 to 1 to M” or “1 to M to M”:

If so, can add a foreign key to that entity set's table

Otherwise, create a table for the relationship with foreign keys for each participating entity