

Lecture I I: Spatial Partitioning, Part 2

CS 6017 – Data Analytics and Visualization

MASTER OF SOFTWARE DEVELOPMENT (MSD) PROGRAM

J. DAVISON DE ST. GERMAIN

SUMMER 2023

Lecture 12 – Topics

2


- Spatial Partitioning
 - Uniform Grid / Bucketing
 - Quad Tree
 - KD Tree (brief intro)

Miscellaneous

- Questions?
- Remember, HW3 is due next Tuesday...


Making KNN Fast

- Naïve: Look at every point
- Today we look at better approaches...



Approach One: Bucketing


- Break space into fixed size "buckets"
- Within each bucket, we just store an array (list) of the points that are contained inside that bucket
- Which training points would we look at to find a value for the red point?
 - `KNN(red_pt, 3)`
 - Which bucket(s) are they in?
 - `[1, 1]`
 - How many buckets do we look at?
 - `[0, 0], [0, 1], [0, 2]`
 - `[1, 0], [1, 1], [1, 2]`
 - `[2, 0], [2, 1], [2, 2]`
 - As many as it takes (moving outward) until we find K points.
 - Note: Do we only consider the points in the red circle?
 - No



- How do we calculate these bucket indices? (Red point is: 13, 14)
 - Can be calculated in $O(1)$
 - How did we put data values together in histograms?
- This is a **fixed size (or uniform) grid**


Approach One: Bucketing (cont.)

- To do a range query, we find out which buckets $(x - \text{radius} \text{ and } x + \text{radius}, y - \text{radius}, y + \text{radius})$ are in, then loop through all the buckets "between" them
- In a single bucket, we just check the distance from p_t of each point and keep it if it's less than the radius
- We can quickly figure out which bucket a point is in with:
for each dimension
$$\text{bucketIndex} = (\text{position} - \text{bucketCorner}) / \text{bucketSize}$$
- If we create N buckets along each dimension, we will have a total of $N^{\text{dimensions}}$ buckets.
 - Note: this can be bad for high dimension data!
- To do a KNN query, we do a range query and increase the radius if we don't find enough points
- Note: *Training* is just creating this data structure and storing the appropriate points in the appropriate bucket.




Approach One: Bucketing (cont.)

- What is a *dimension*? Note: I'm currently displaying 2 dimensions...
- Let's think about this in terms of house sell price... What might the axes (ie, the dimension) be for each axis? x? y?...
 - X axis: # of square feet
 - Y axis: Location
 - Z axis: # of Garage Spaces
- More axes... and I've run out of letters!!?
 - X_3 : # of Bathrooms
 - X_4 : Lot Size
 - X_5 : Age of roof
 - ...



Approach One: Bucketing (math)


- How to determine the x and y values of upper left and lower right points on the grid itself?
 - $\min(\text{pts.x}), \max(\text{pts.x})$
 - $\min(\text{pts.y}), \max(\text{pts.y})$
- Bucket index math?
 - $\text{col} = (\text{p.x} - \text{grid_min.x}) / \text{cell_width}$
 - $\text{row} = (\text{p.y} - \text{grid_min.y}) / \text{cell_height}$
- Which buckets are “near” the red point?
 - Determine the “square” around point of interest
 - $\text{p.x} \pm \text{radius}, \text{p.y} \pm \text{radius}$



Uniform Grid / Bucket Approach


9

- Issues
 - 1 (giant) bucket vs 1000s of small buckets?
- Trickiness
 - Choosing the # of buckets...
 - S splits (3 in this example)
 - S^{dim} buckets...
 - As *dim* gets big?
 - “Curse of Dimensionality”
 - Storing the array of buckets...
 - Could “flatten” the array
 - Need a way to covert “row 1, column 2” into an index in a 1D array
 - Note: most languages support 1 and 2 dimensional arrays... but we may have many dimensions...
 - Could use a hash function for this. If a lot of buckets are empty, this works very well! (**Spatial Hashing** – Row,Col as key to hash table)
 - In 2D you could say `index = row*numColumns + column`.
 - Can produce a similar formula for any number of dimensions




Uniform Grid Issues

- What are some issues that a uniform grid of buckets might have?
 - Non-uniform data...
 - Perhaps make more buckets?
 - Empty buckets cause overhead
 - Still issues, so...
- Uniform grid works well if data points are relatively evenly distributed...
 - Note: Non-uniform (stretched) grid is a possibility – though finding buckets can be more difficult.



10




Spatial Partitioning

- These approaches we are discussing are in a family of algorithms for **spatial partitioning**
 - Breaking up “space” into regions
- Uniform grid works well on uniformly distributed data...
 - For Non-uniform data → Want non-uniform buckets
- Need a “data-aware” structure
 - Will look at two different *tree-like* data structures which are designed to split up the space where there are more points.
 - Building these data structures is more complicated, but they are better adapted for a lot of real data sets.

Approach Two: Quad Tree


- Create a grid
 - *Adaptively* add more cells based on...
 - distribution of data
- Fix number of splits to 1...
- Thus always dealing with a 2x2 grid.
- How do we reduce the number of points in a grid cell?
 - Split cells with too many points again
 - Repeat recursively...



Quad Tree

13


- A node is either a "bucket" (array of points)
 - Or an internal node with 4 children (each of the children covers 25% of parent's area)
- We pick a "threshold" number of points to decide whether to split a node or make it a bucket (leaf)
- Each node stores the region it covers (an “axis aligned bounding box” (AABB)) and either a list of points, or pointers to its 4 child nodes
 - Children are sometimes named NW, NE, SW, SE
- Quad Tree because each sub-cell is broken into 4 children... What is the 3D version of this called?
 - Oct Tree
 - Note: # of children per node is $2^{\text{dimensions}}$.
 - Thus usually these are only used up to 3 dimensions.



Example where each node is only allowed to store one point (before being sub-divided).

Quad Tree Construction


- Input – “Point Soup”
 - Points are not organized in any way
- Since a Quad Tree is a Tree:
 - `root = new Node(points, aabb)`
 - `aabb` specifies the area the node is responsible for
- Recursive approach...
- So what is the base case?
 - # of points is less than X , then this is a leaf node (and we are done).
 - Note: X is a *knob* that is used to *tune* the data structure
- Otherwise, create four children...
 - How do we decide which points to send to each child?
 - Find midpoints: $(aabb.max + min) / 2$
 - Loop through the points and place them in the NW, NE, SW, SE list of points
 - “Array Partitioning”
 - `NW = new Node(NW_list, nw_aabb)`



Quad Tree Querying

15

- Function - `QT::rangeQuery(pt, radius)`
- Recursive algorithm, so:
 - `points = Node::rangeQuery(pt, radius)`
- `points = root.rangeQuery(my_pt, r)`
- Only need to “think about” handing this for a single node, and let recursion take care of the rest.
- Two choices:
 - Leaf node:
 - Look through the array of points this leaf contains.
 - Is point within r of `my_pt`?
 - otherwise?
 - determine which child/children `my_pt+radius` is/are in, and call (for example):
 - `ne_points = NE.rangeQuery(my_pt, radius)`
 - Note: Depending on radius, we may have to call the `rangeQuery` on multiple children. (Bounding box checks with AABB)



- `Node::knnQuery(pt, K, result)`
 - **Note:** `result` keeps track of the list of the nearest `K` points
- Similar to range query, but if statements are more complicated...

```

if leaf
    for each point in bucket...
        if len( result ) < K
            add point
        else if distance( point, pt ) < distance( pt, worst in result )
            replace worst with point
else // internal node
    for each child
        if necessary // <- return length < K or closestPointInAABB( pt ) is closer than worst in list
            recurse


```

- Can think of this as having a “search radius”
 - The search radius starts at infinity, but shrinks as we add points to the `result` list.

Quad Tree KNN Query

Quad Tree Query Complexity



- Asymptotic Complexity (Big O)?
- Height of tree? (Number of levels down?)
 - $\sim \log_4 N$
- Query time: $\sim \log_4 N \diamond K \diamond r$
- If tree is not balanced... And depending on K and radius the number of cells we are looking at can change (think the “width” of the search)...
- How could we split better?
 - See next slide...



Approach Three: KD Tree

18

- Generalization of a BST to many "keys"
- Each node stores the median of it's subtree according to one of the dimensions (x, y, z, etc)
- Each time we go down one level, we move to the next dimension (if I split by x, my children split by y)
- Each node stores a point and the dimension it splits by
- More tomorrow



~ Fin ~